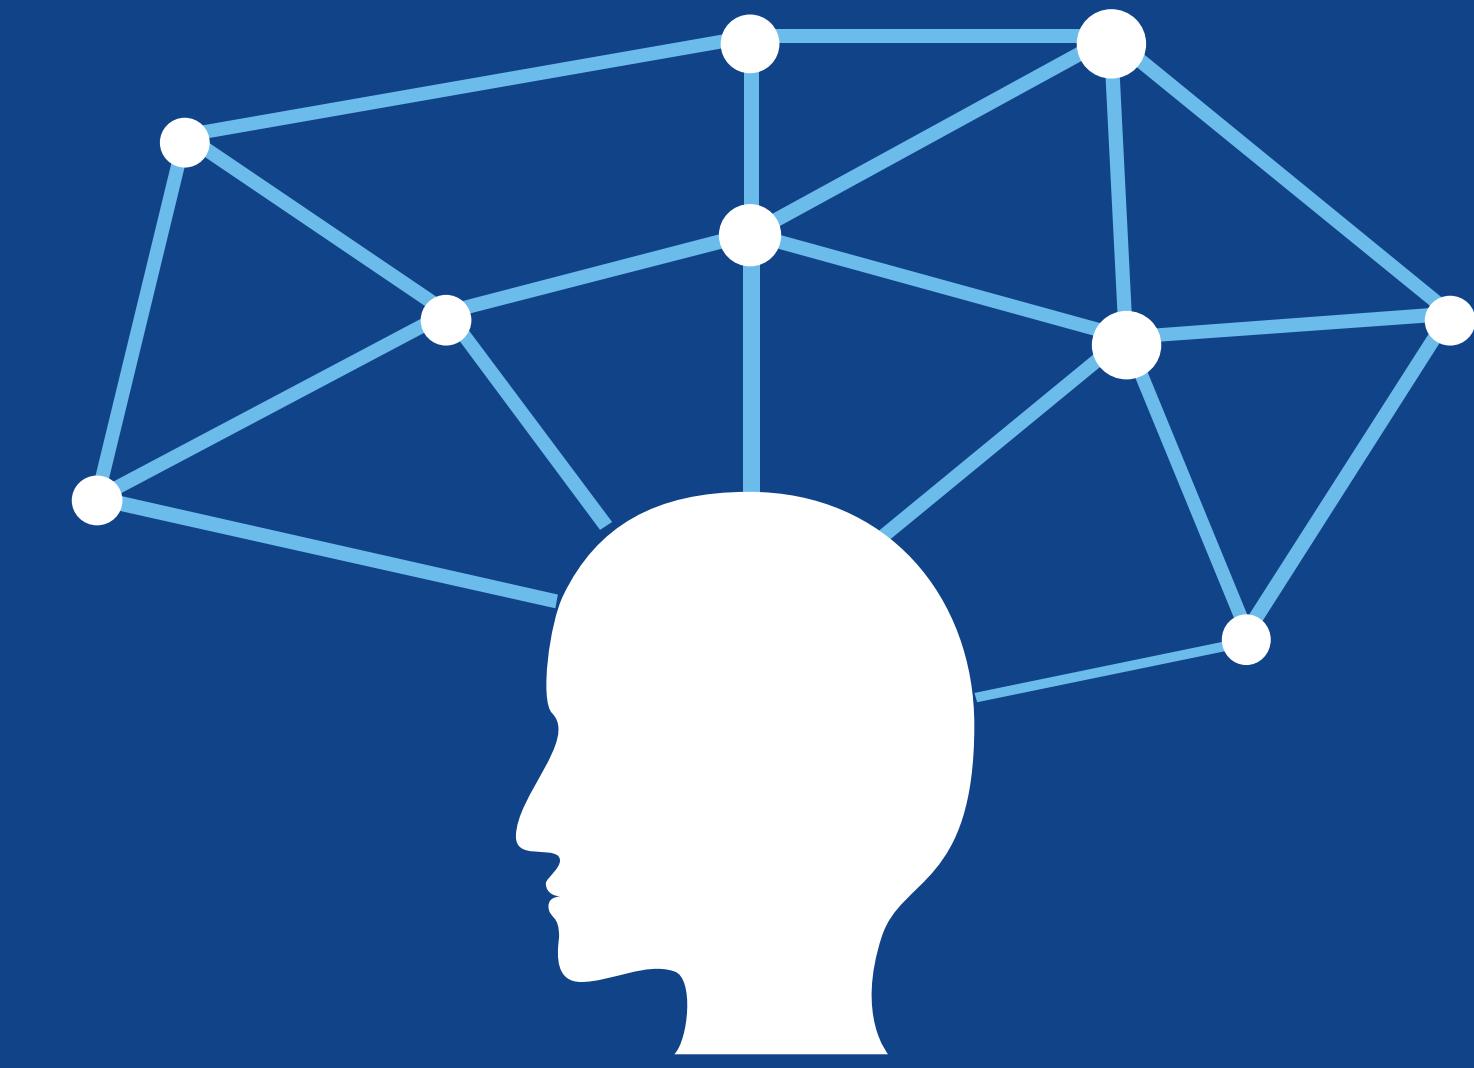
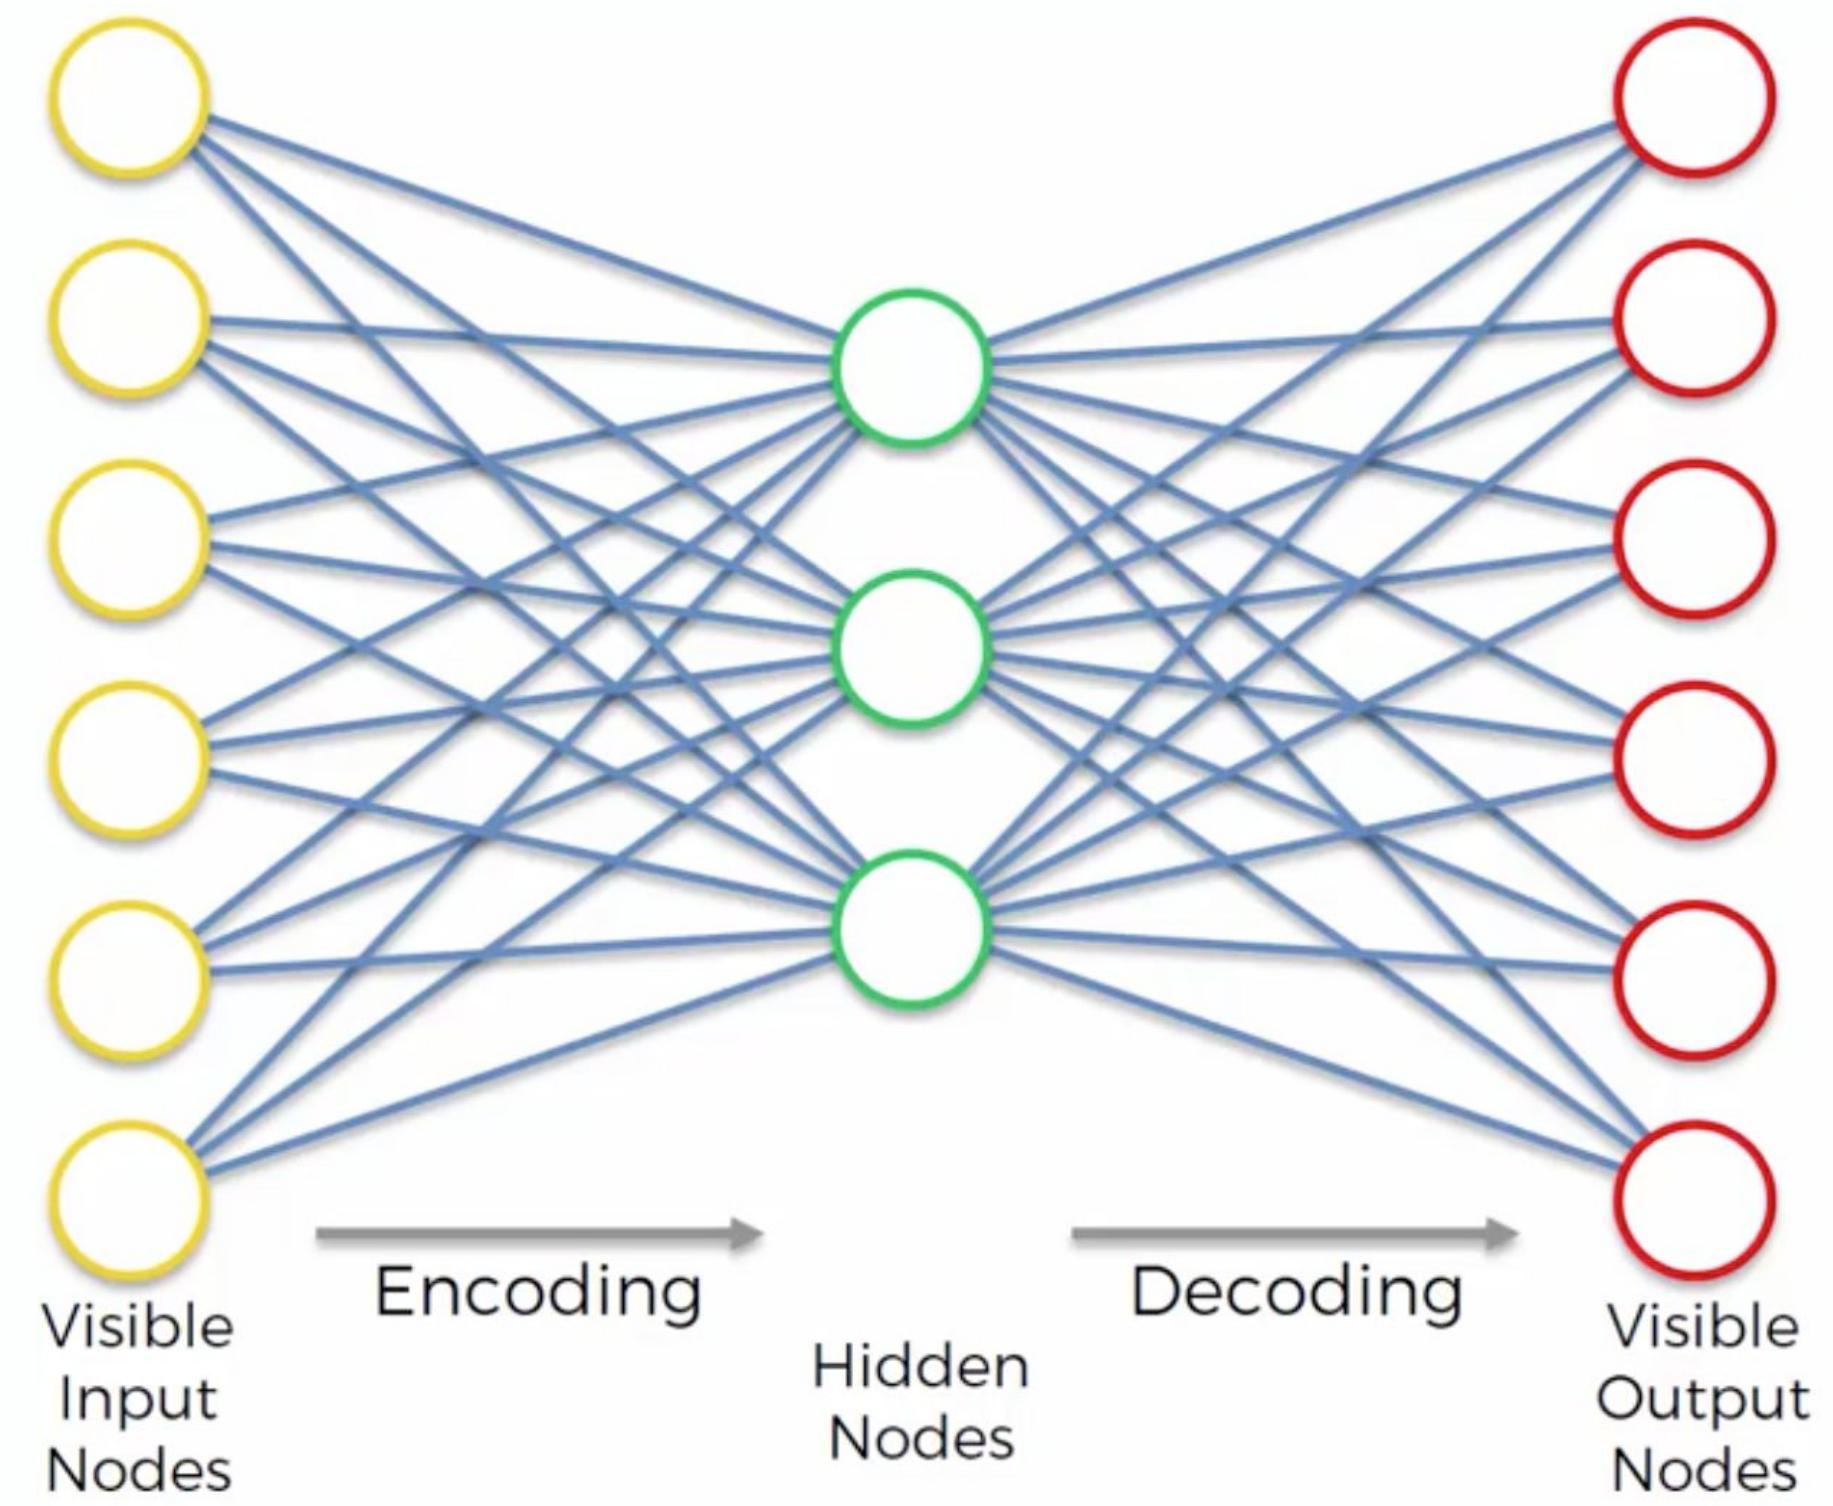


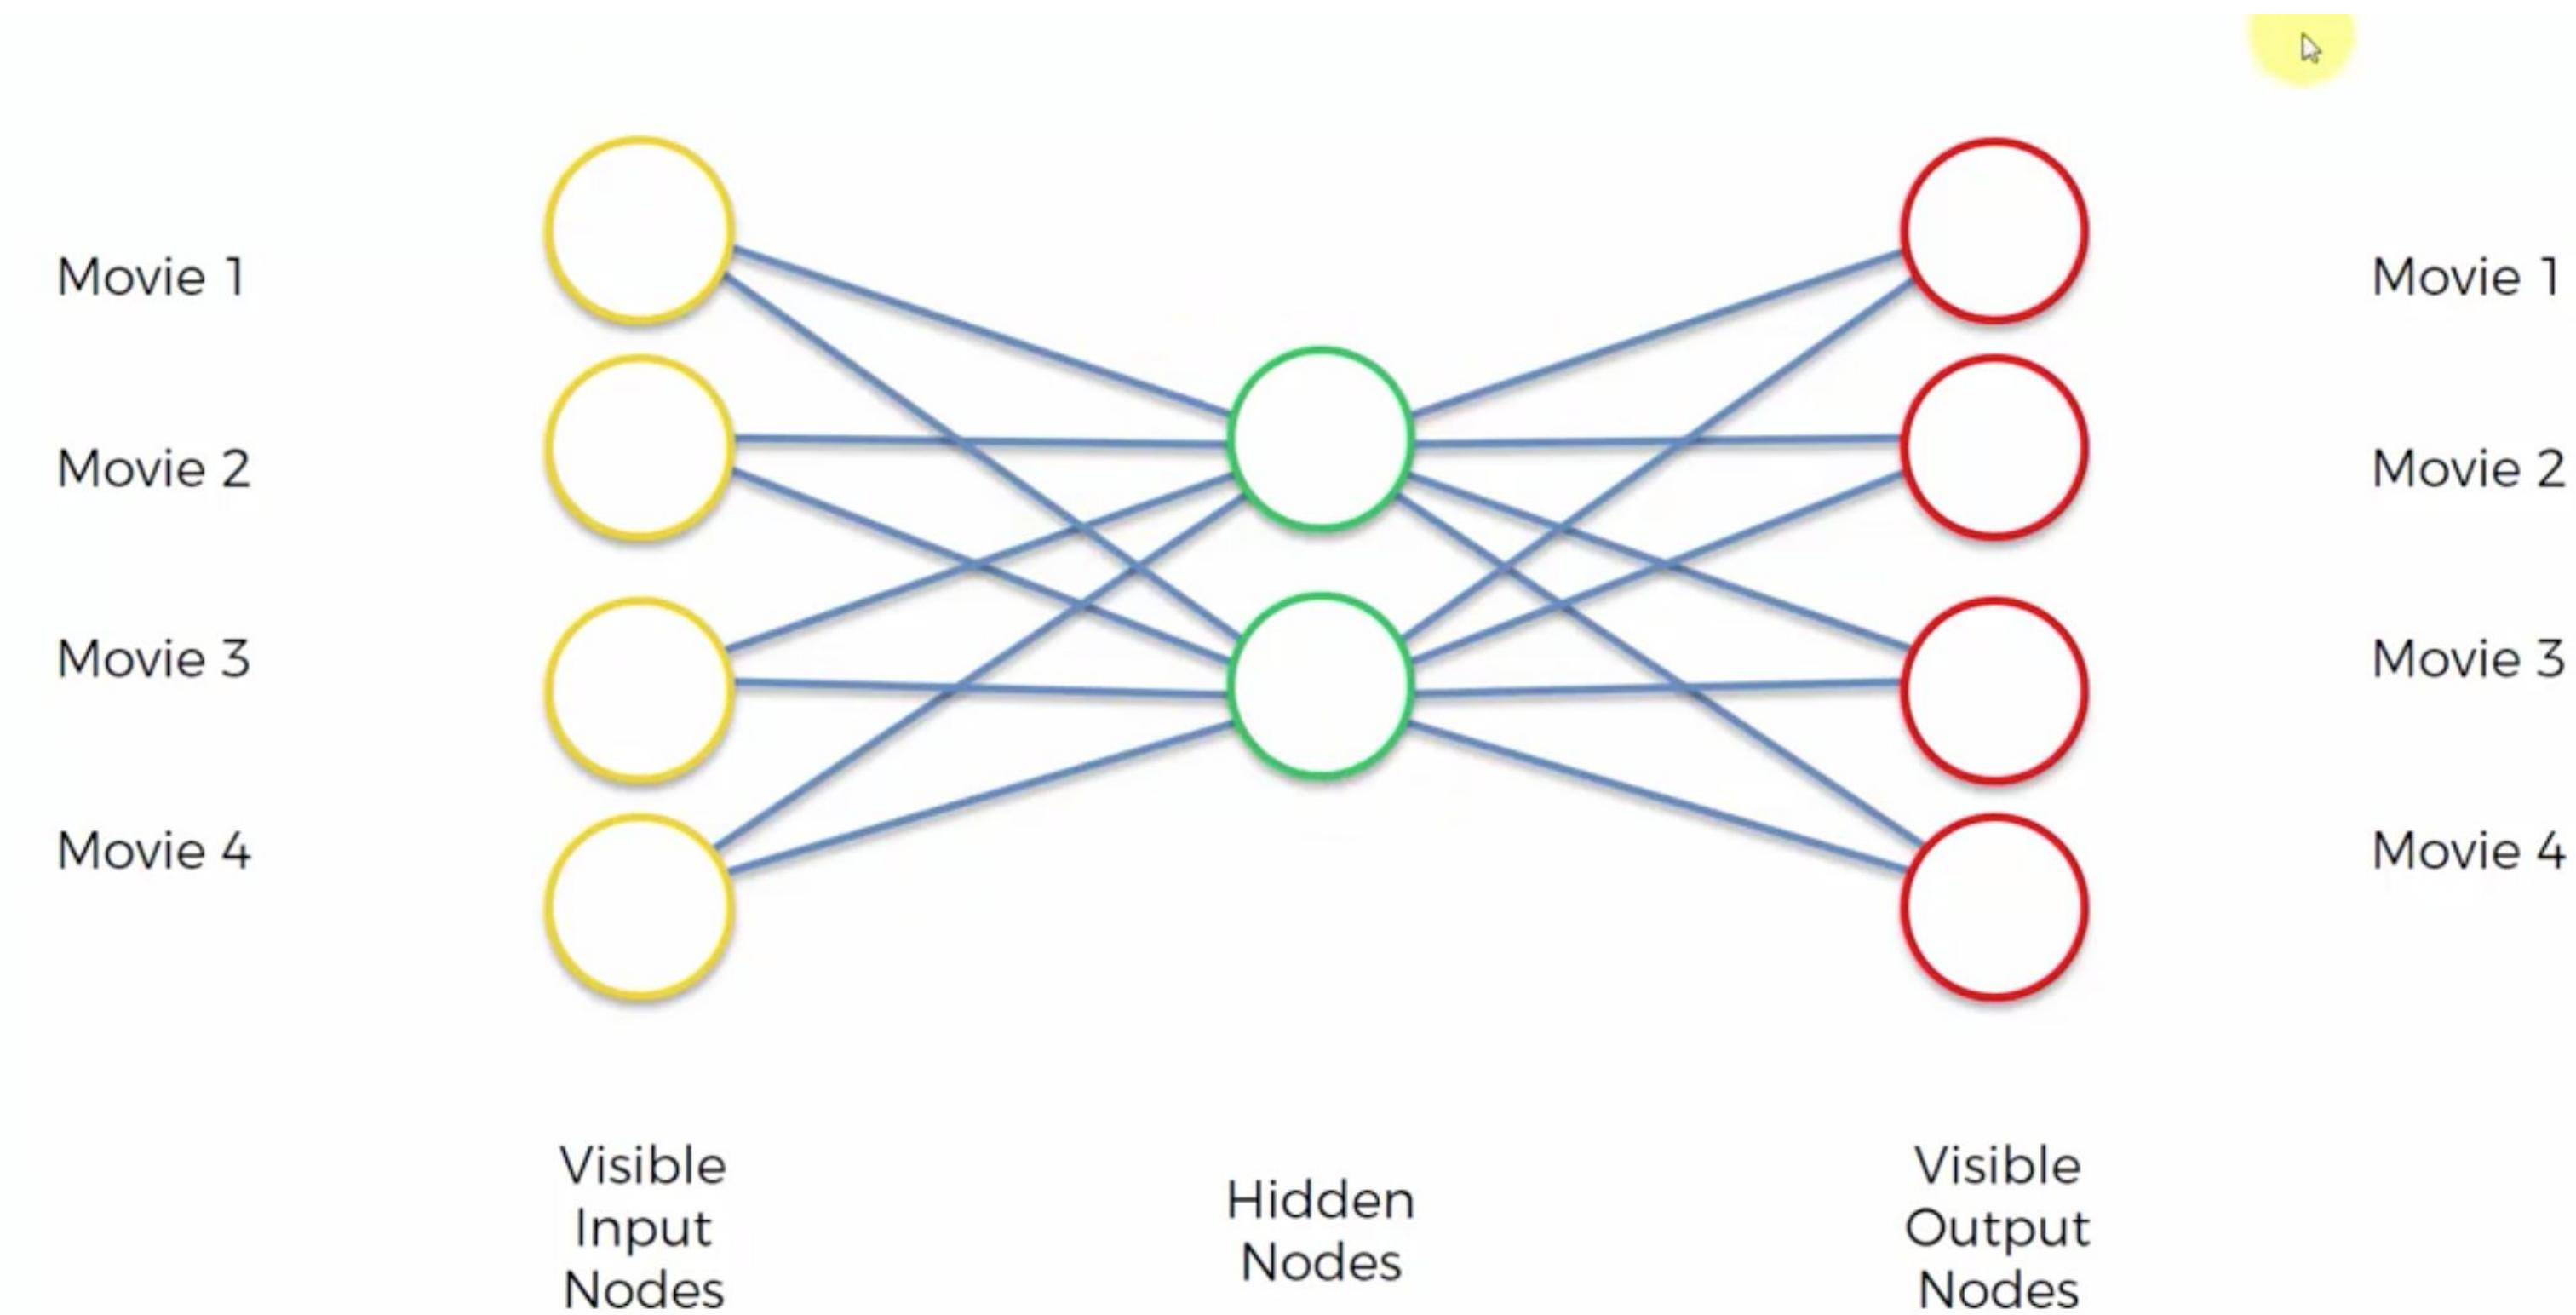
# Auto Encoders



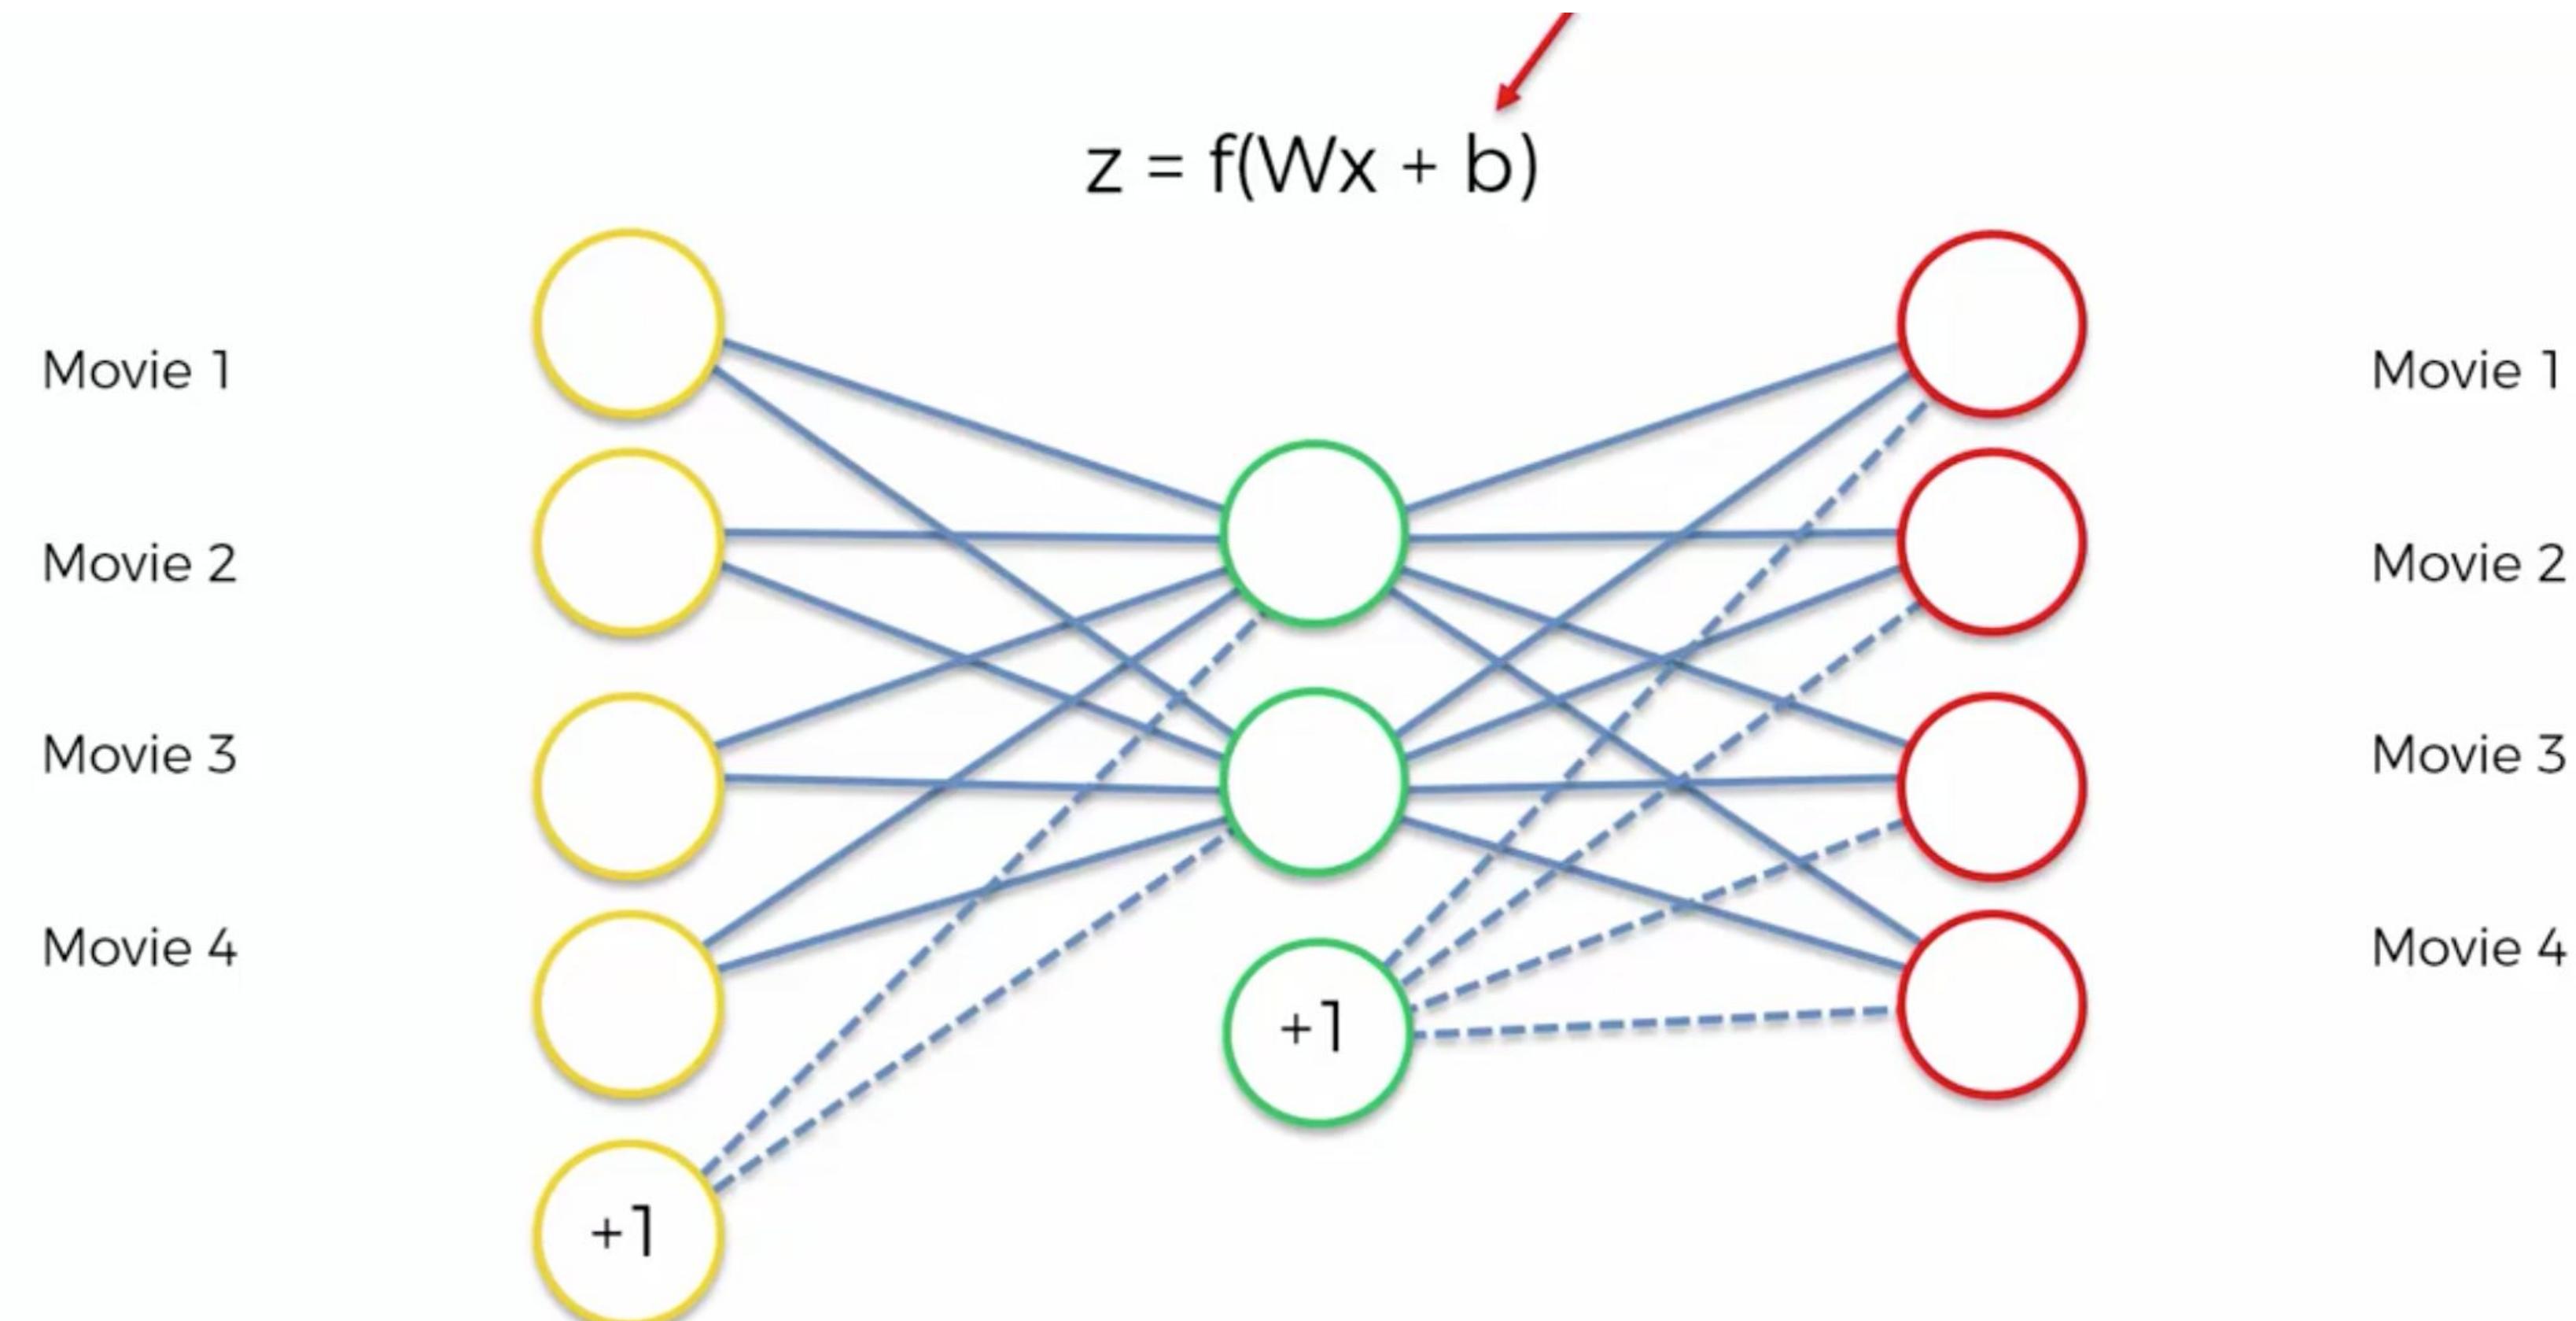
# Auto Encoders



# Auto Encoders



# Auto Encoders - Bias



# Auto Encoders - Bias

**STEP 1:** We start with an array where the lines (the observations) correspond to the users and the columns (the features) correspond to the movies. Each cell ( $u, i$ ) contains the rating (from 1 to 5, 0 if no rating) of the movie  $i$  by the user  $u$ .



**STEP 2:** The first user goes into the network. The input vector  $x = (r_1, r_2, \dots, r_m)$  contains all its ratings for all the movies.



**STEP 3:** The input vector  $x$  is encoded into a vector  $z$  of lower dimensions by a mapping function  $f$  (e.g: sigmoid function):

$$z = f(Wx + b) \text{ where } W \text{ is the vector of input weights and } b \text{ the bias}$$



**STEP 4:**  $z$  is then decoded into the output vector  $y$  of same dimensions as  $x$ , aiming to replicate the input vector  $x$ .



**STEP 5:** The reconstruction error  $d(x, y) = \|x-y\|$  is computed. The goal is to minimize it.



**STEP 6:** Back-Propagation: from right to left, the error is back-propagated. The weights are updated according to how much they are responsible for the error. The learning rate decides by how much we update the weights.

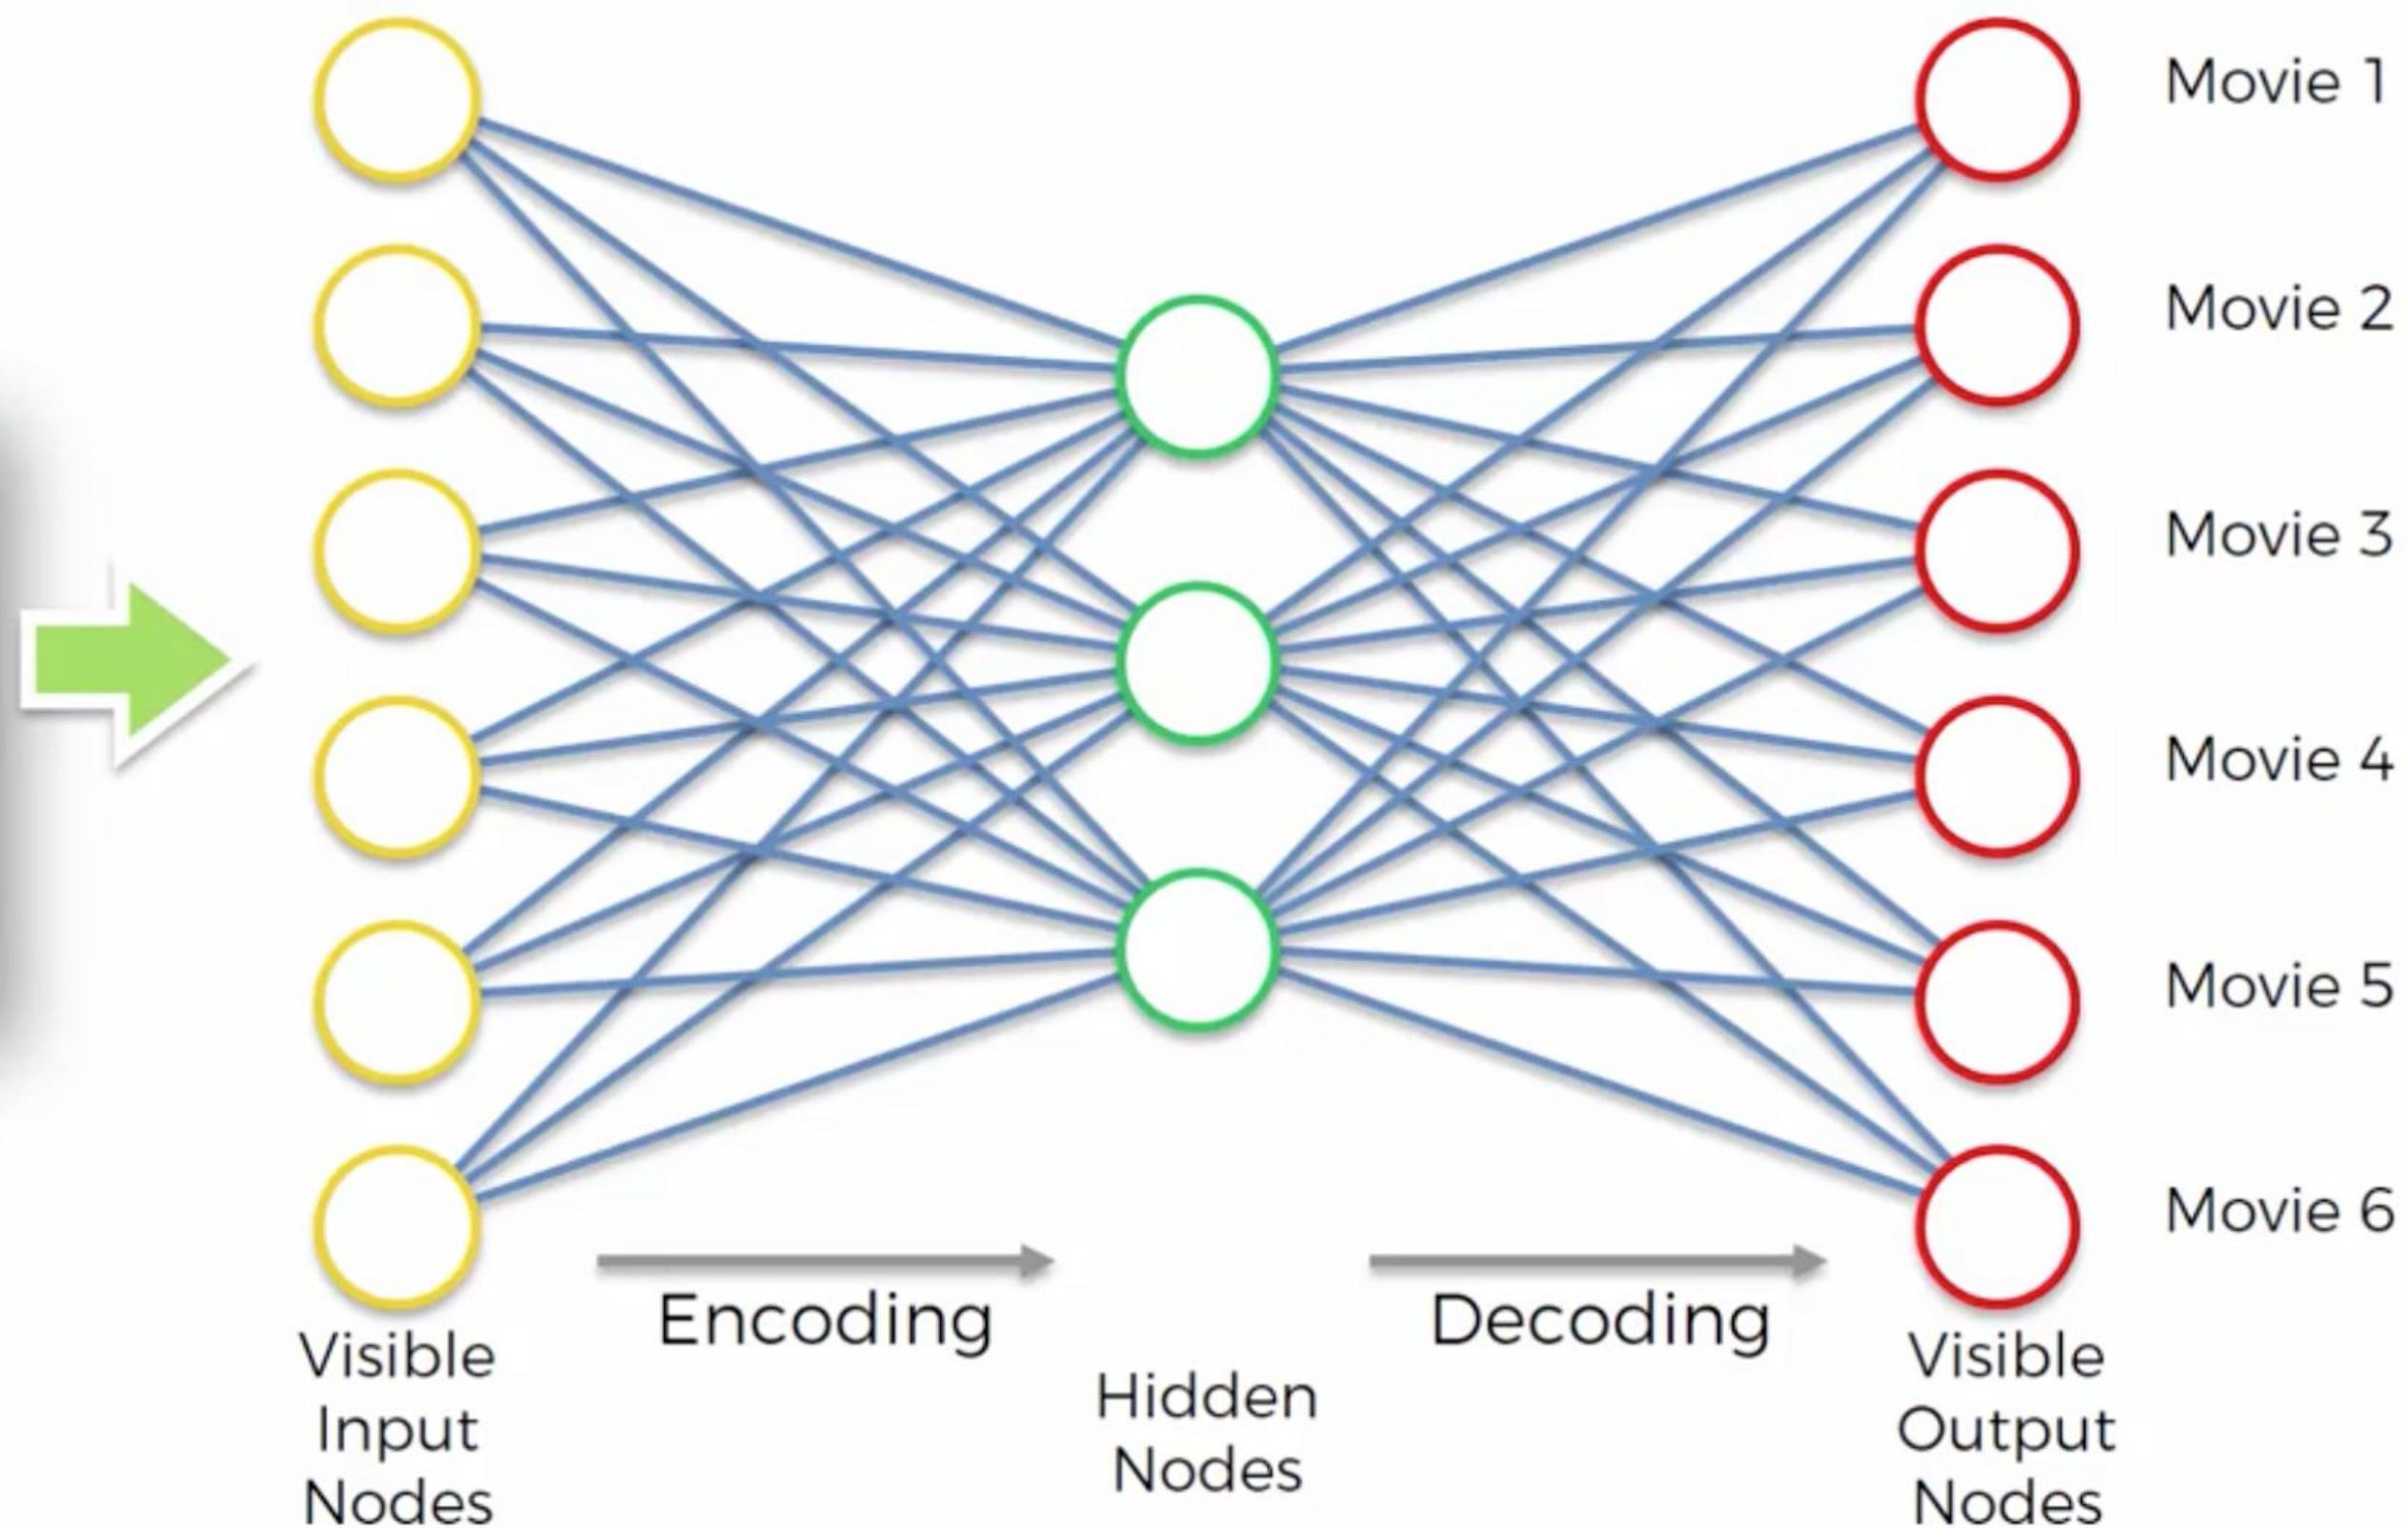


**STEP 7:** Repeat Steps 1 to 6 and update the weights after each observation (Reinforcement Learning). Or:  
Repeat Steps 1 to 6 but update the weights only after a batch of observations (Batch Learning).

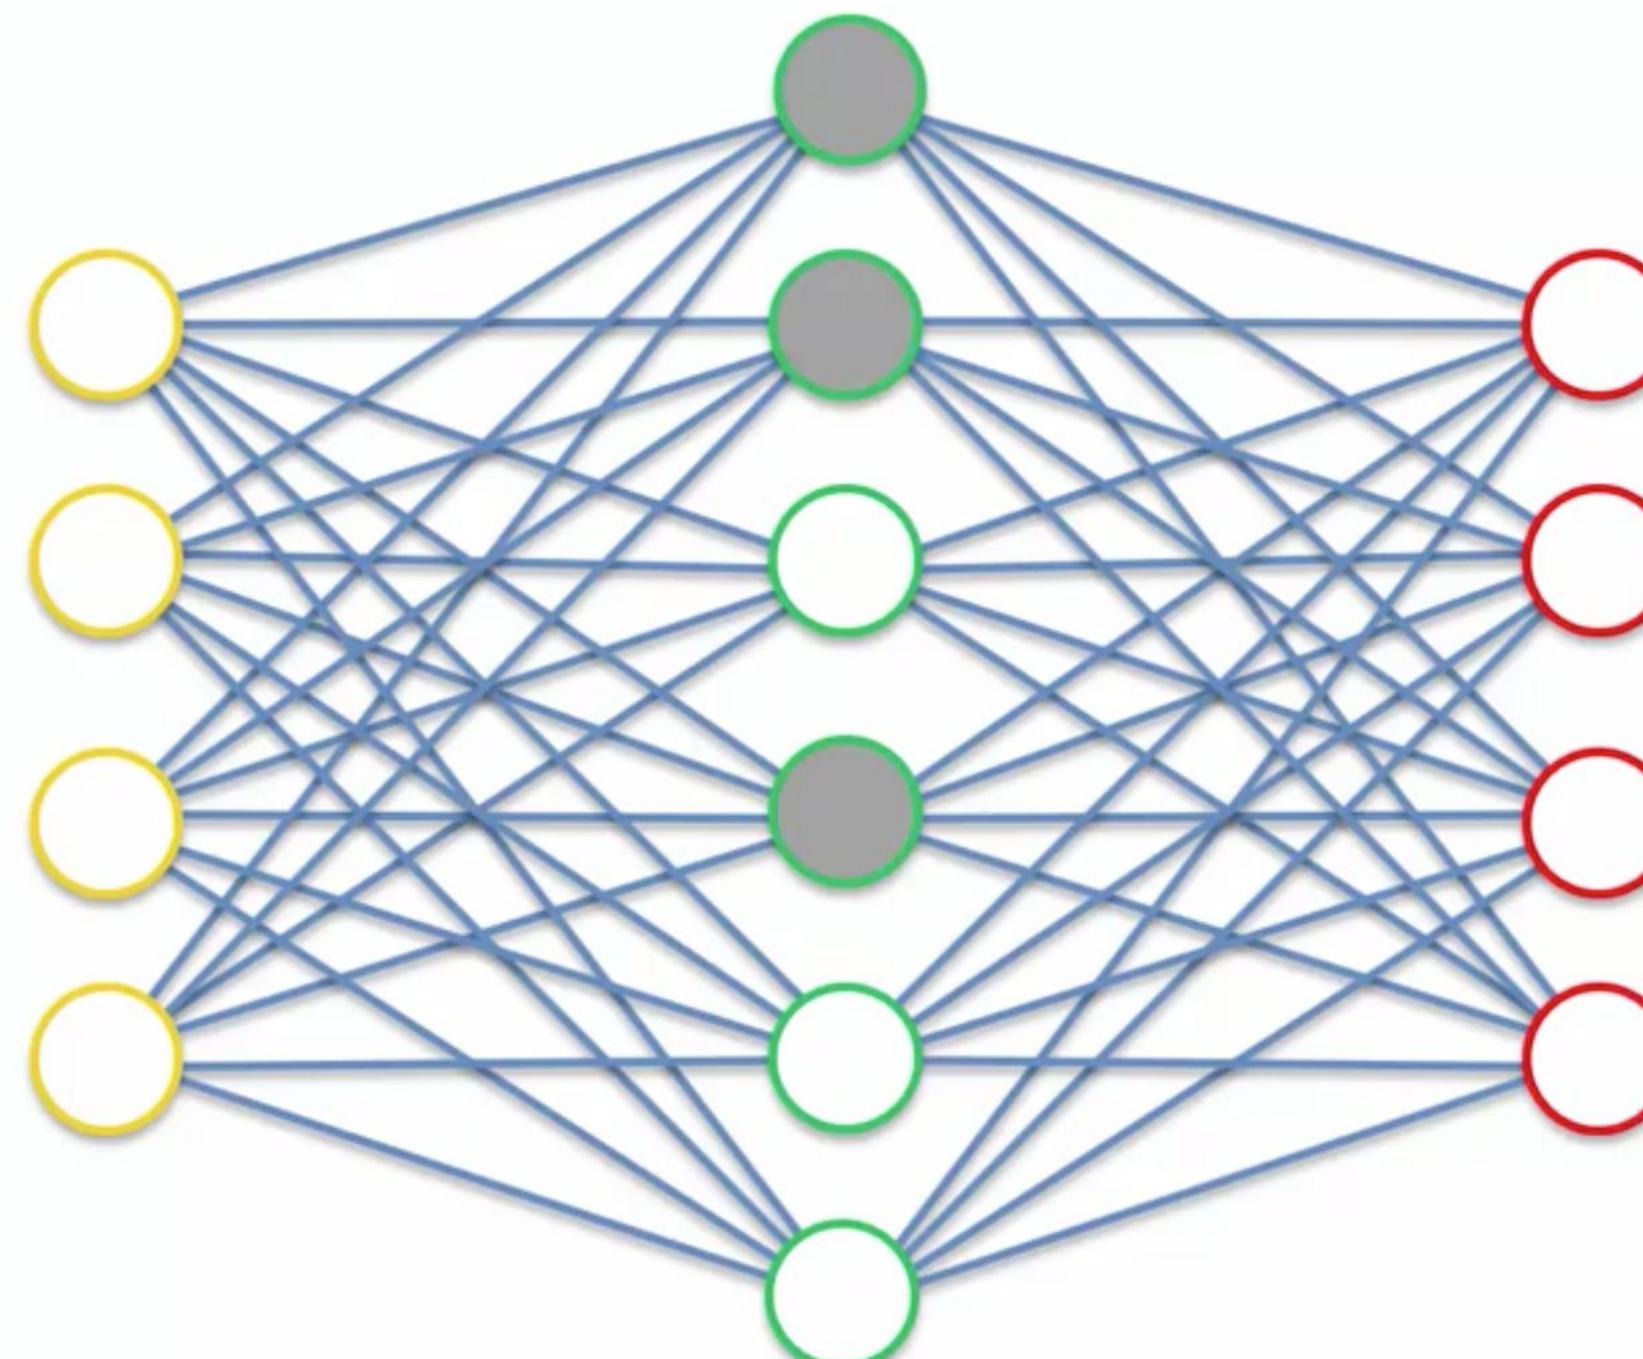
# Auto Encoders - Movies Example

STEP 7

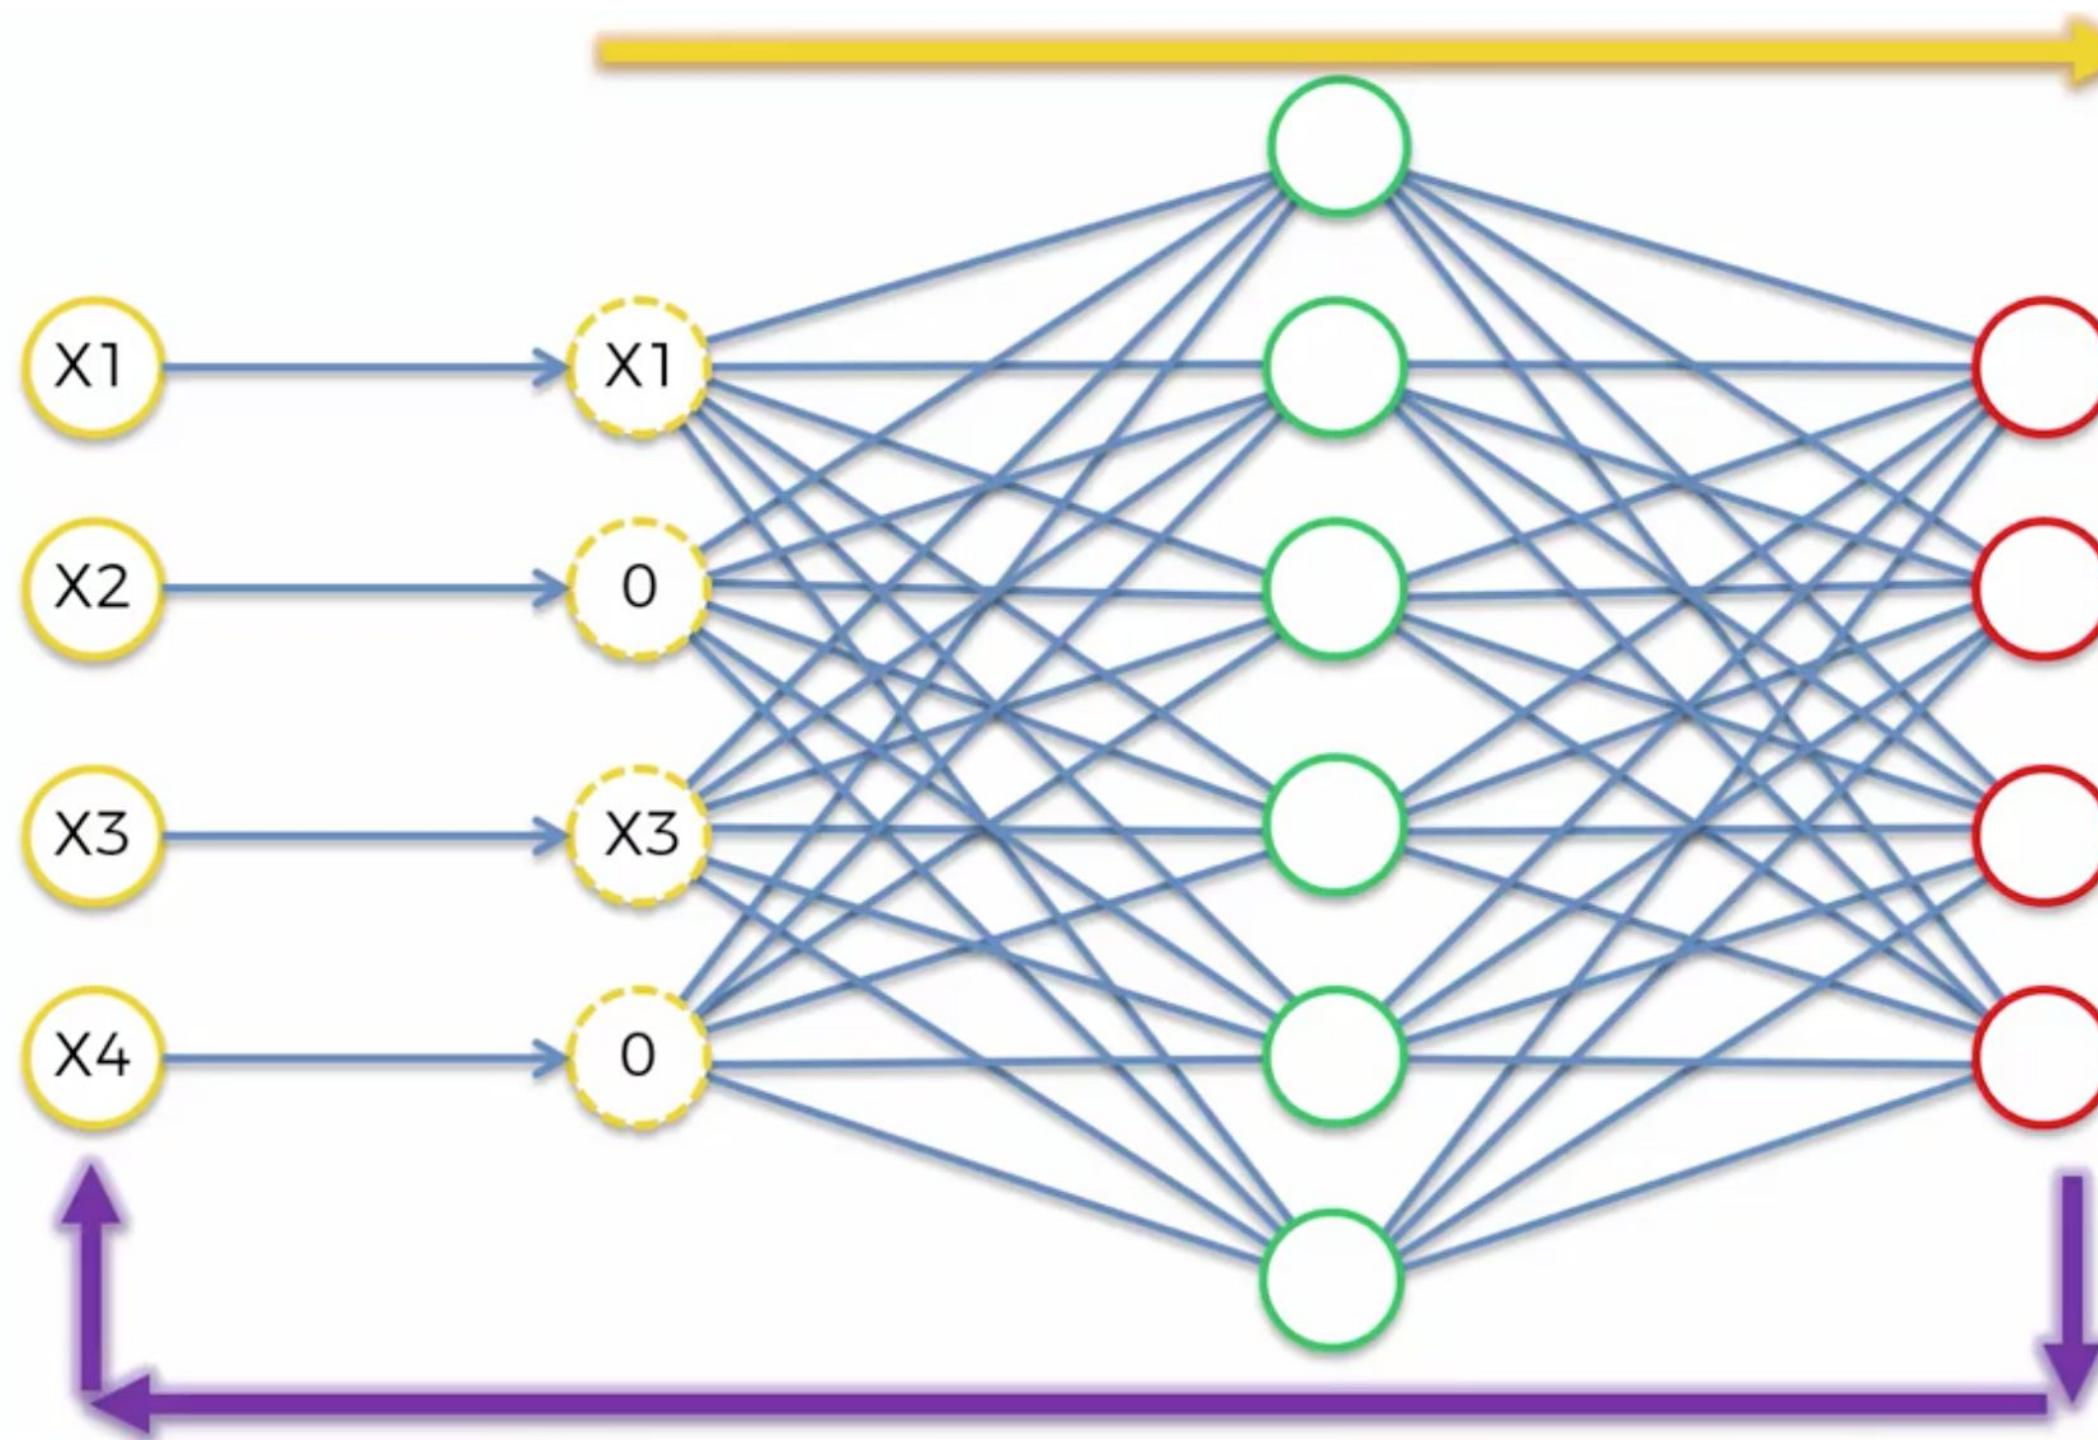
	Movie 1	Movie 2	Movie 3	Movie 4	Movie 5	Movie 6
User 1	1	0		1	1	1
User 2	0	1	0	0	1	0
User 3		1	1	0	0	
User 4	1	0	1	1	0	1
User 5	0		1	1		1
User 6	0	0	0	0	1	
User 7	1	0	1	1	0	1
User 8	0	1	1		0	1
User 9	0	1	1	1	1	
User 10	1		0	0		0
User 11	0	1	1	1	0	1



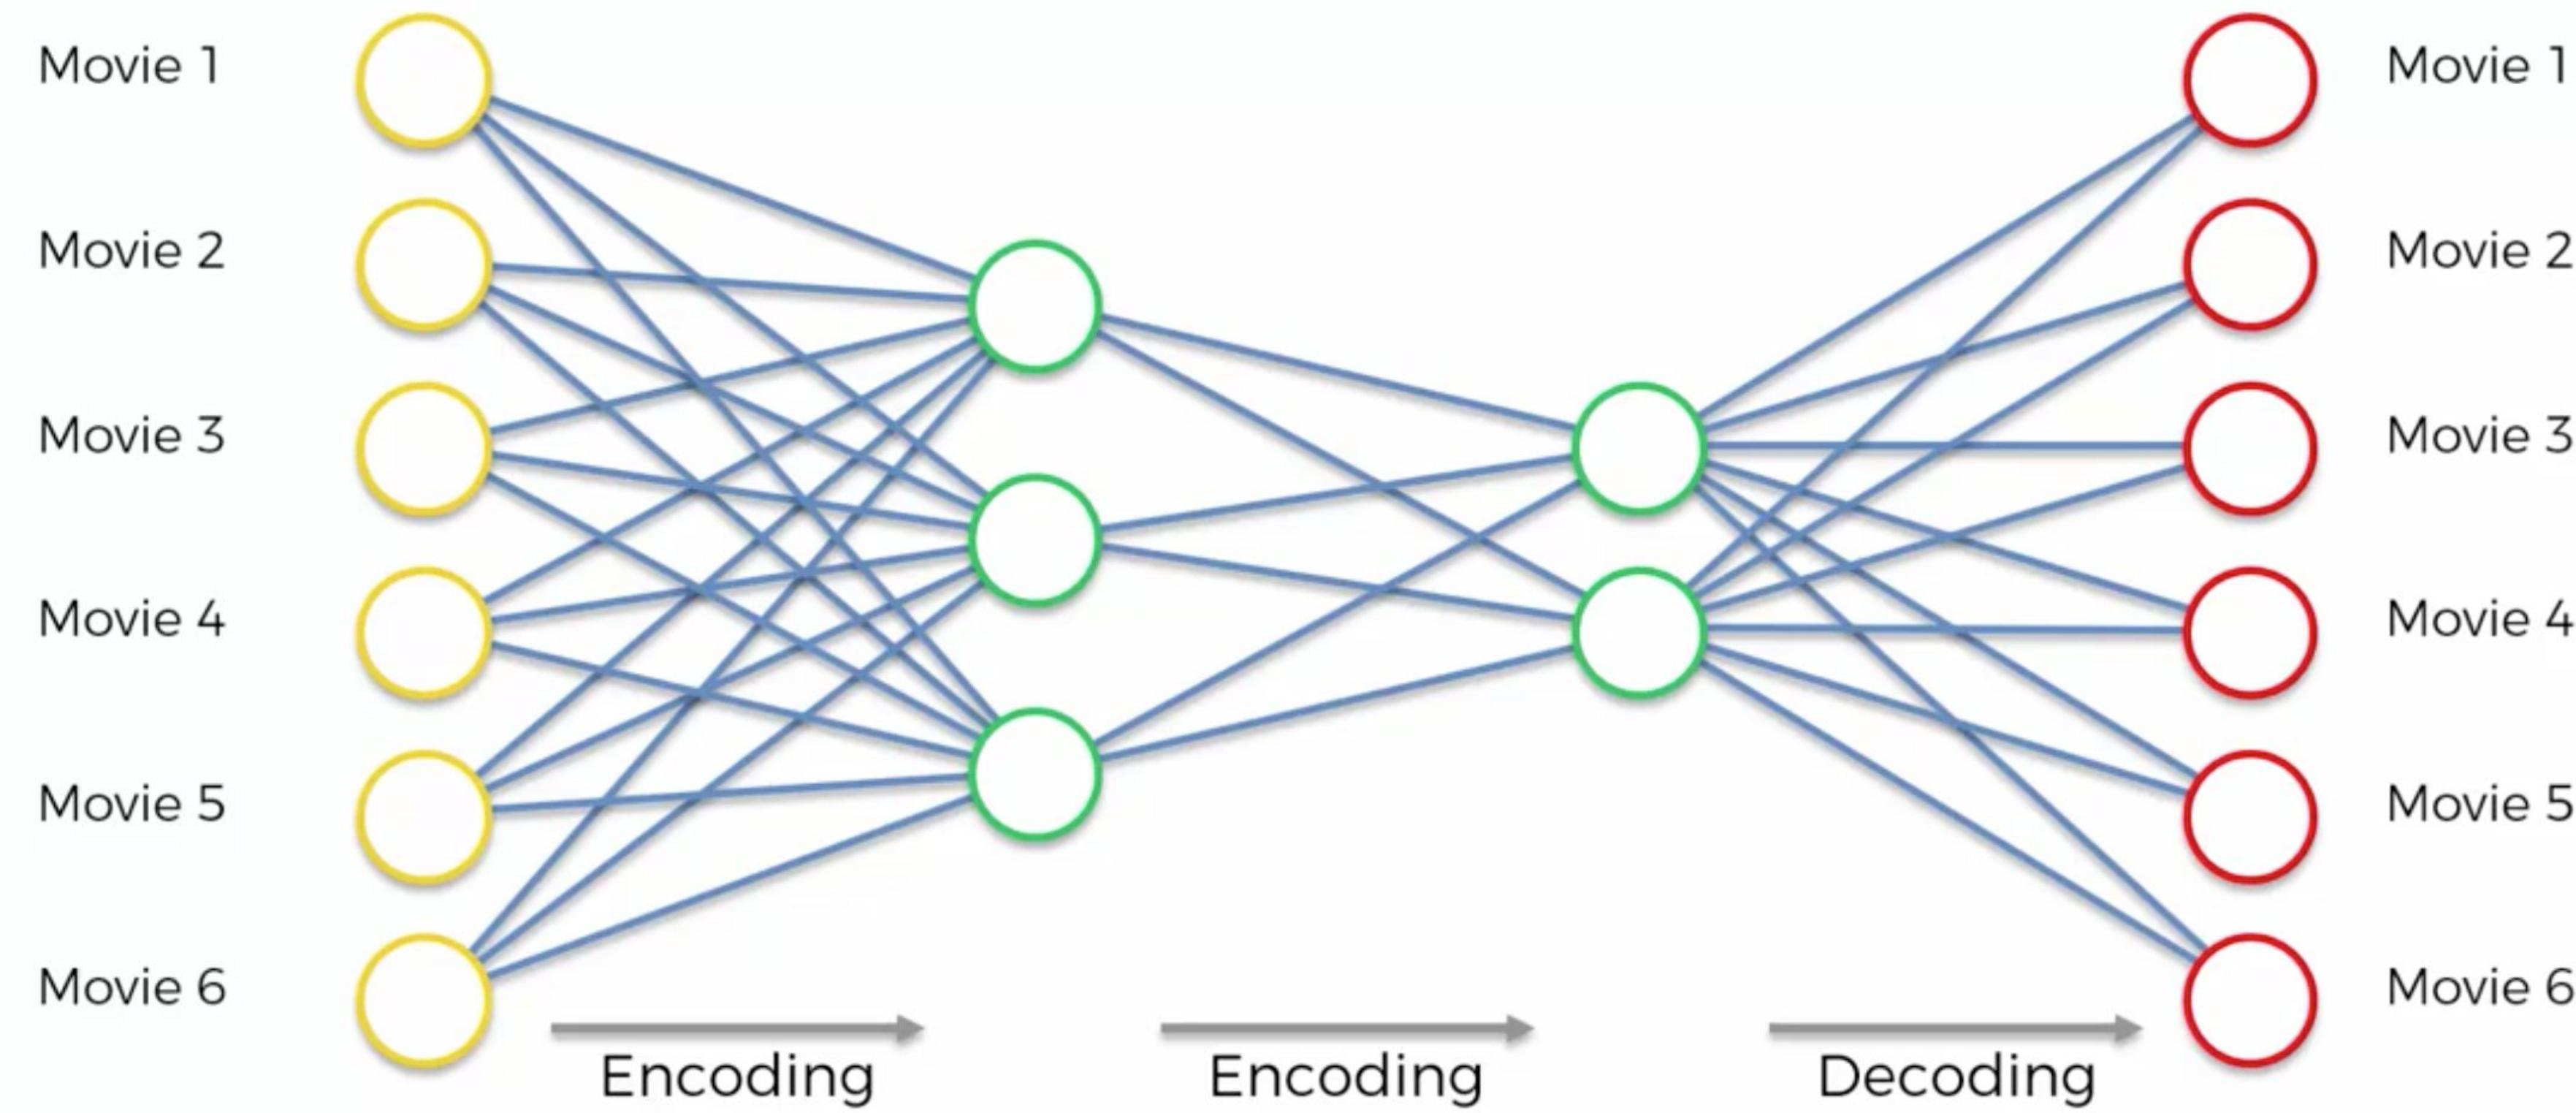
# Over Complete Hidden Layers Sparse AutoEncoders



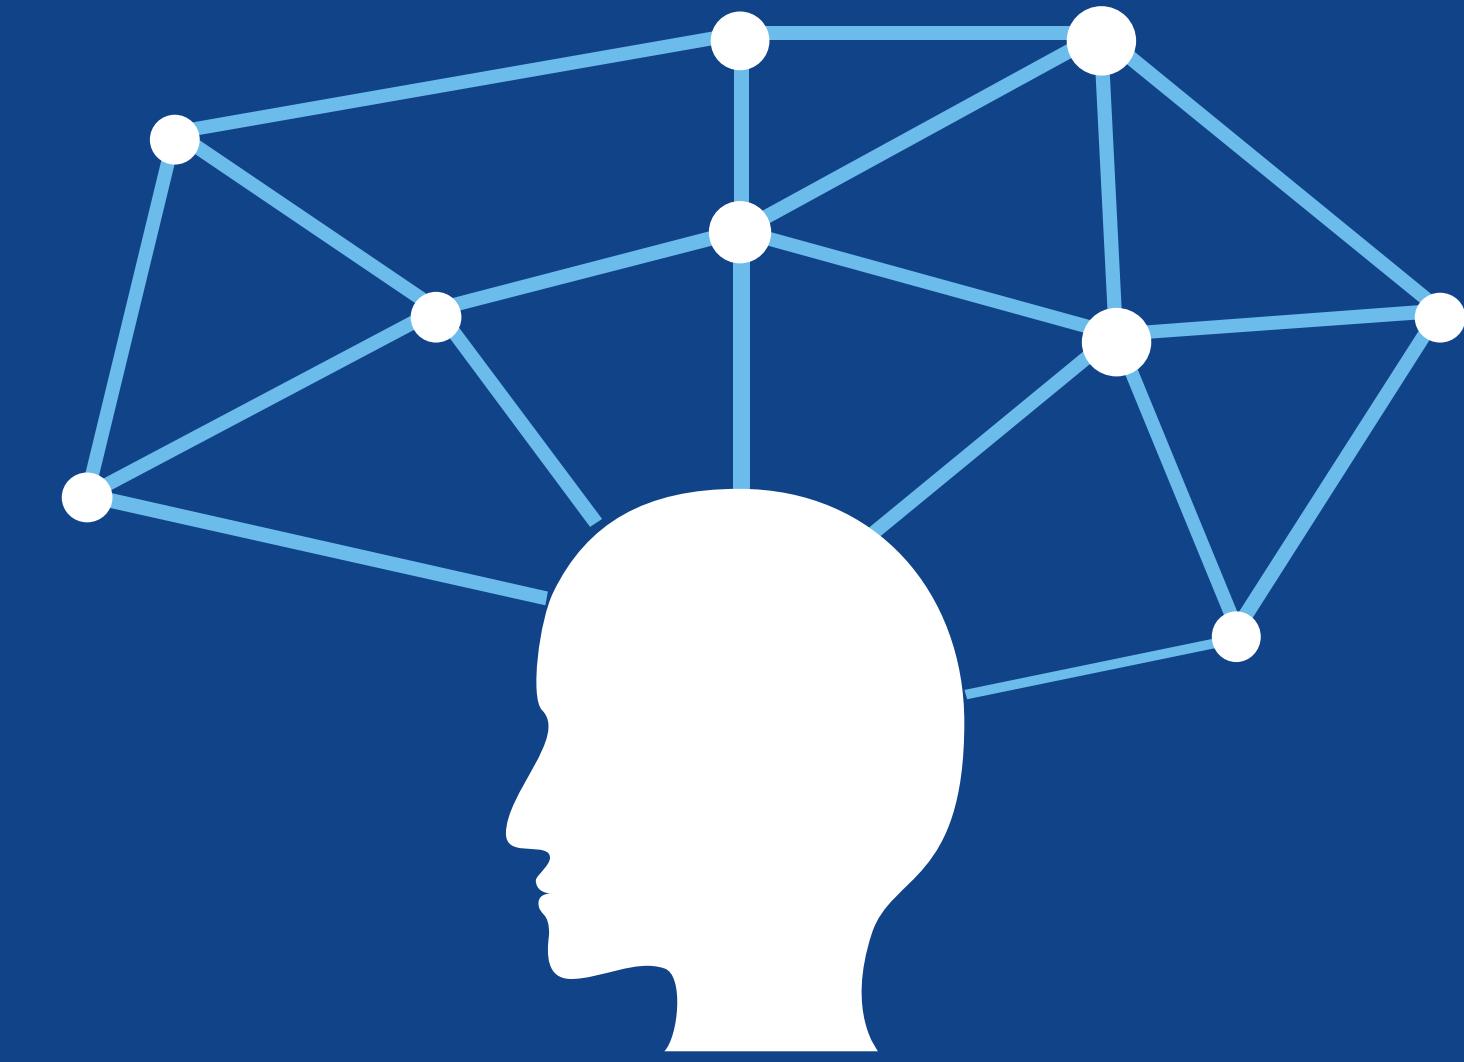
# Over Complete Hidden Layers Denoising AutoEncoders



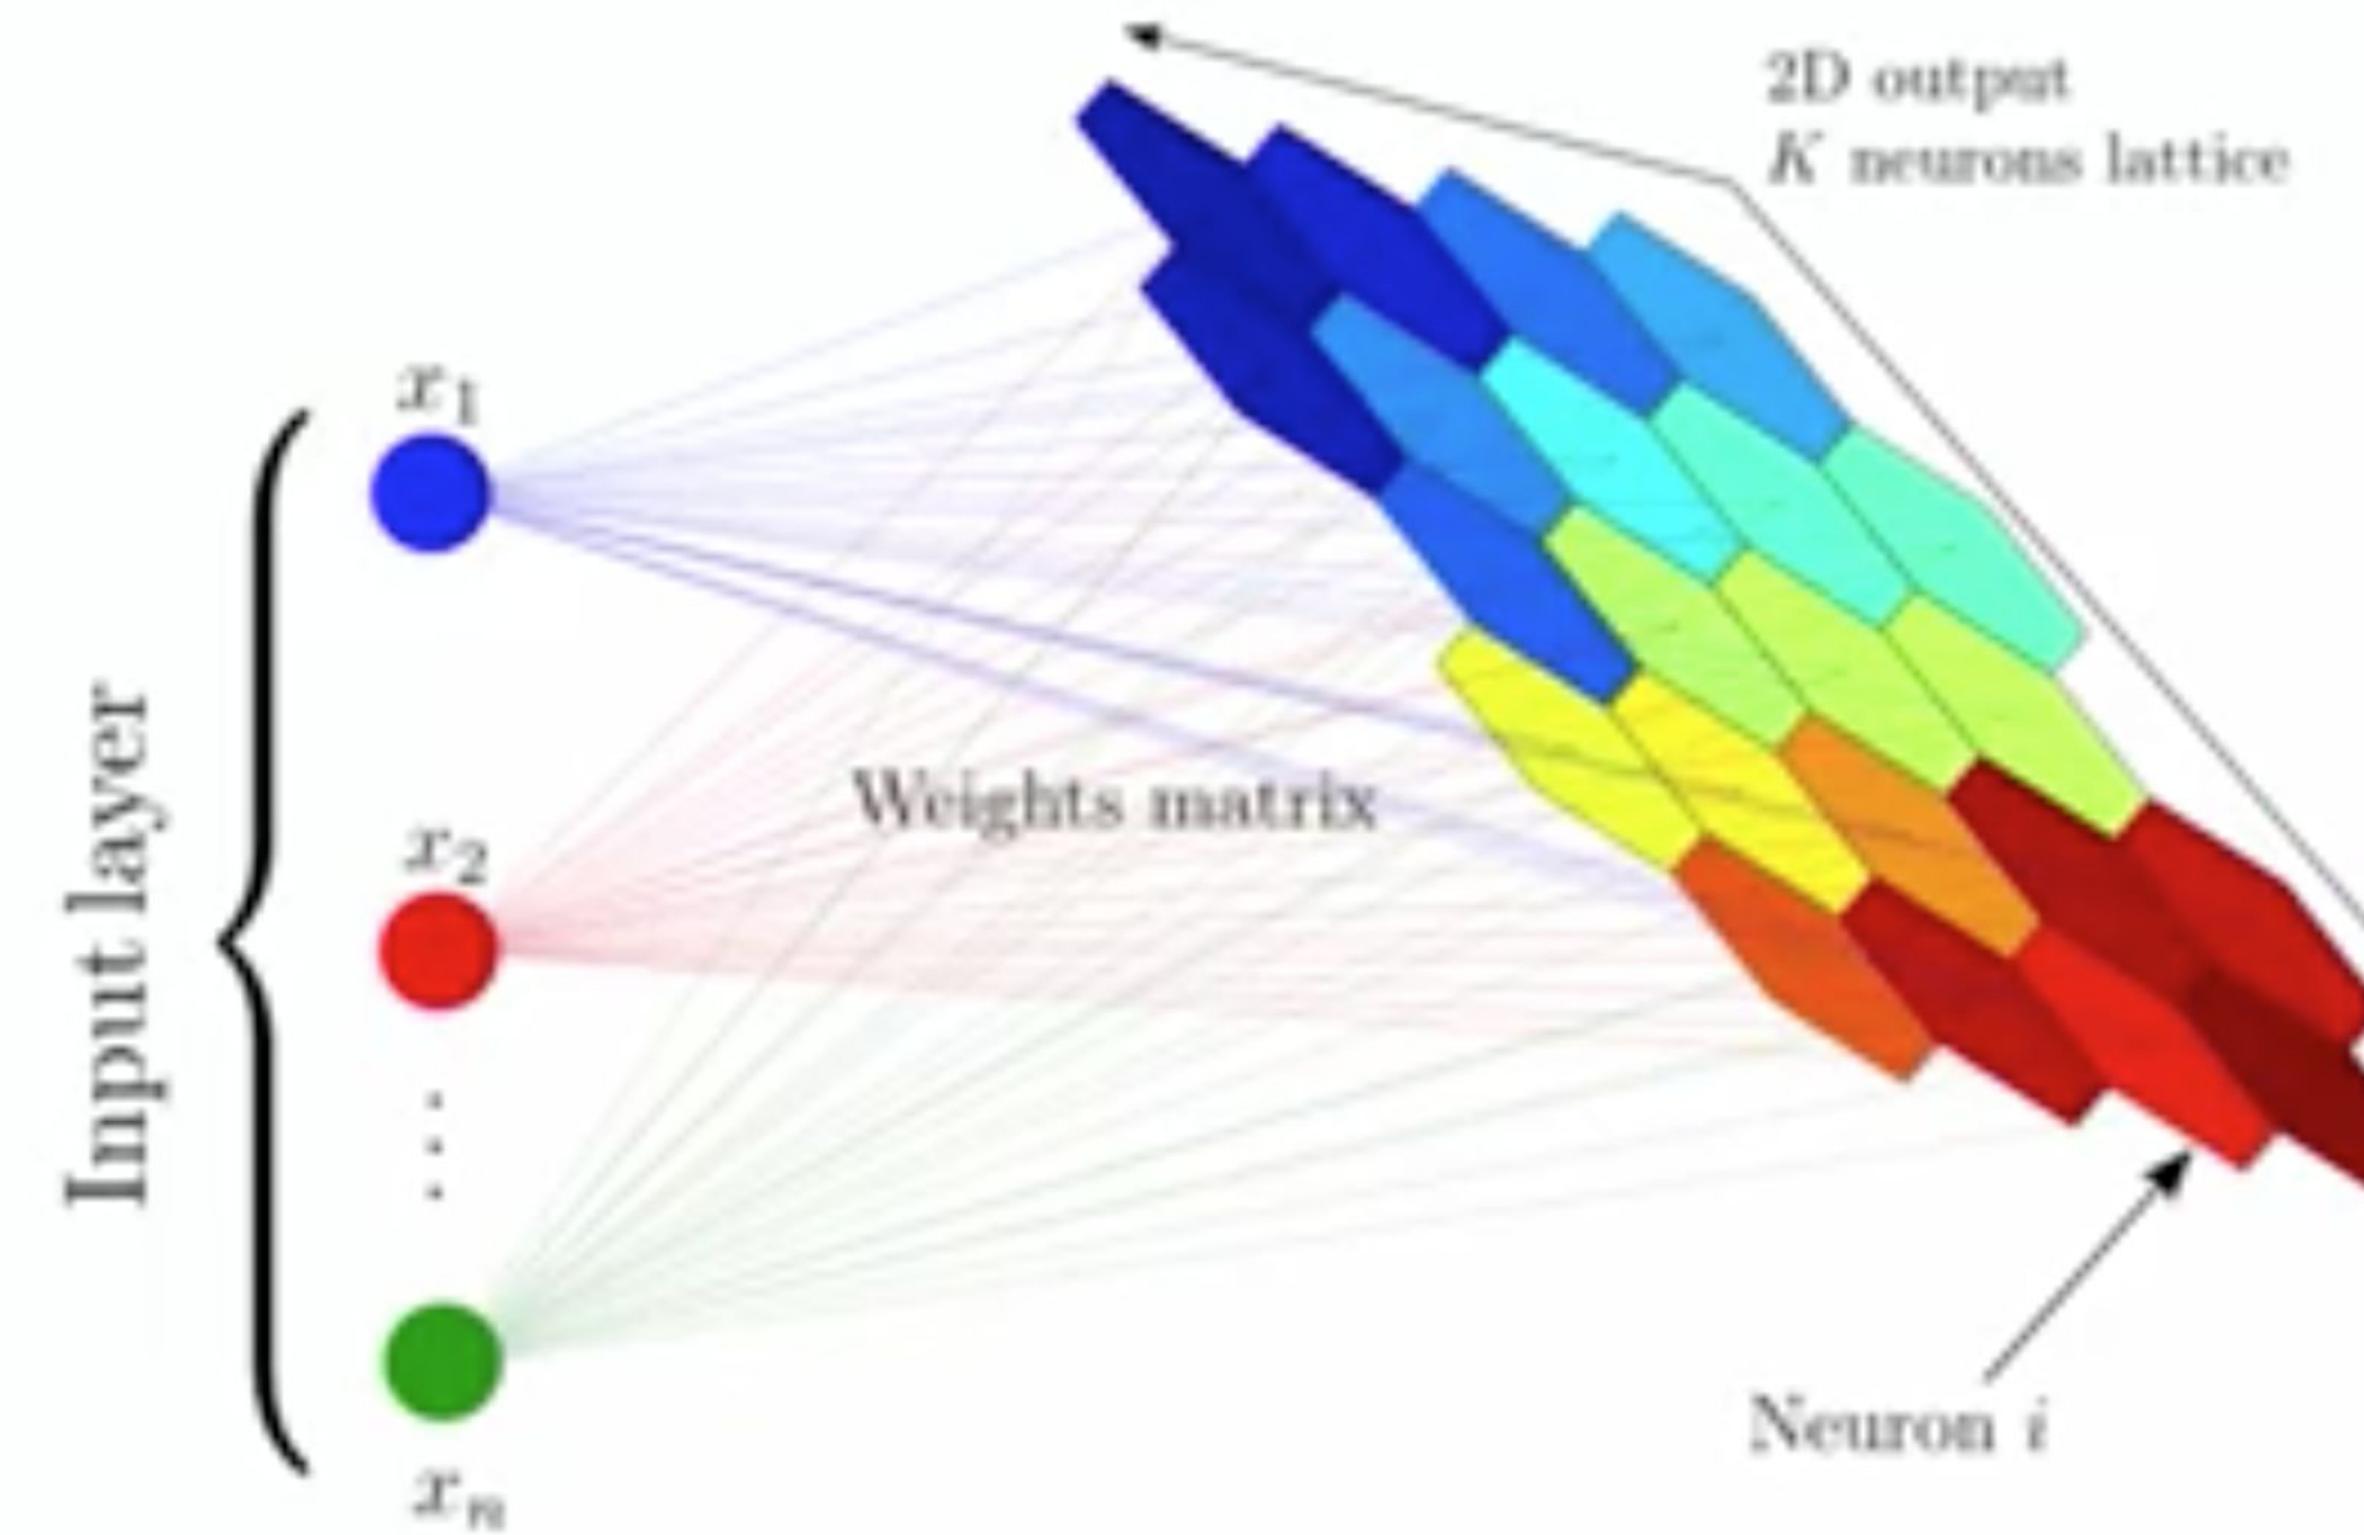
# Over Complete Hidden Layers Stacked AutoEncoders



# Self Organizing Maps

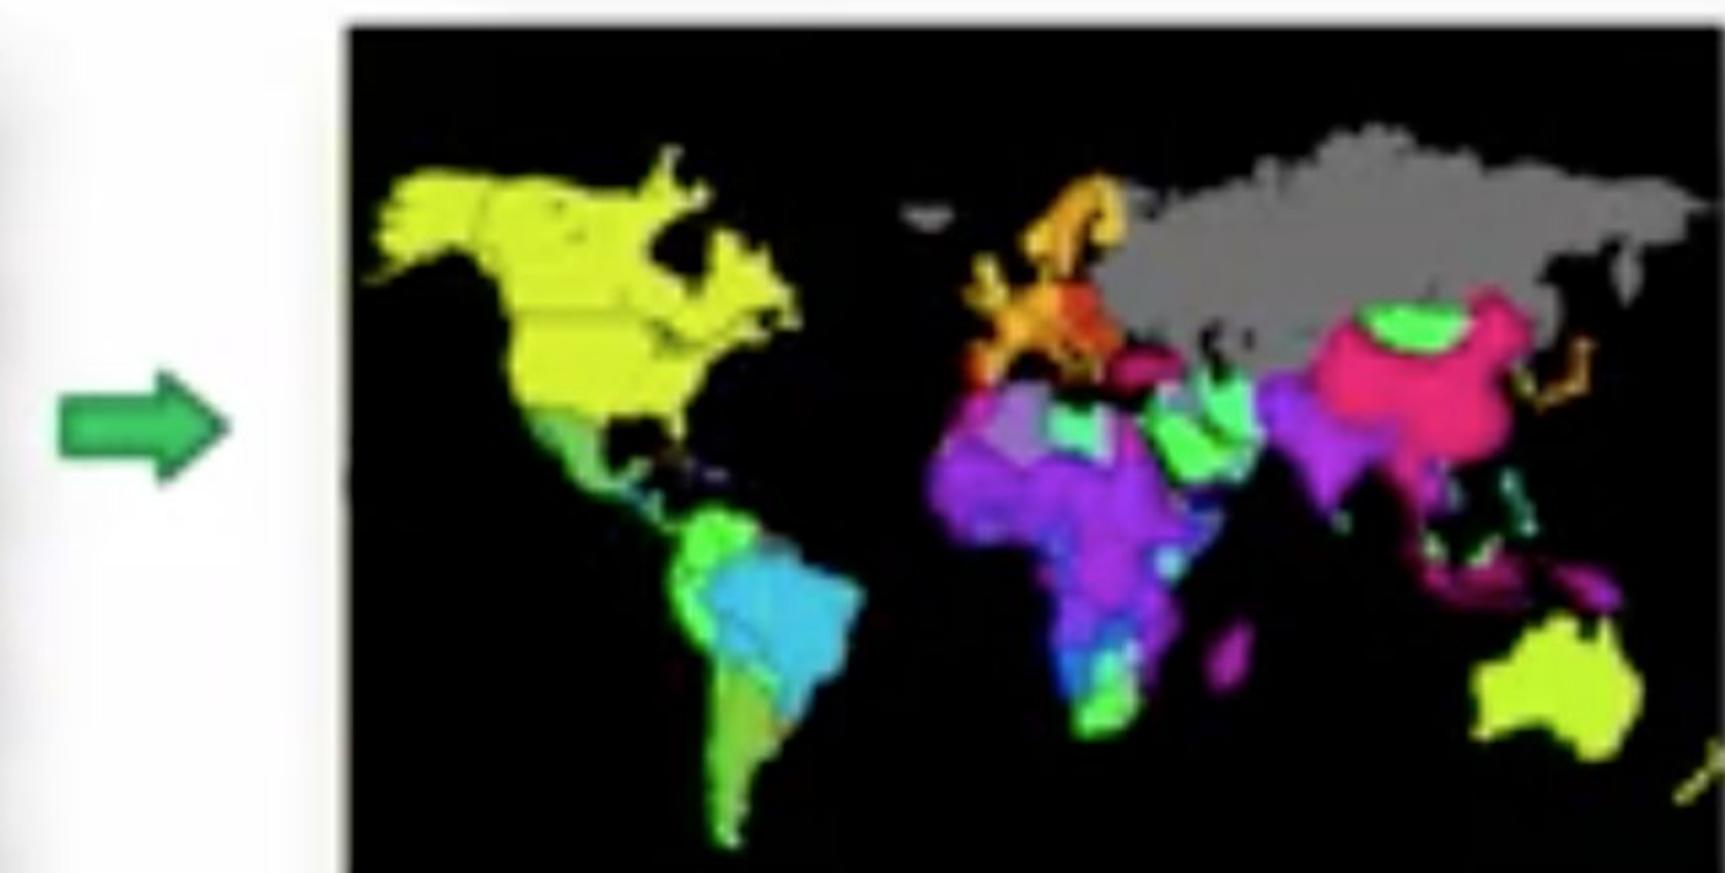
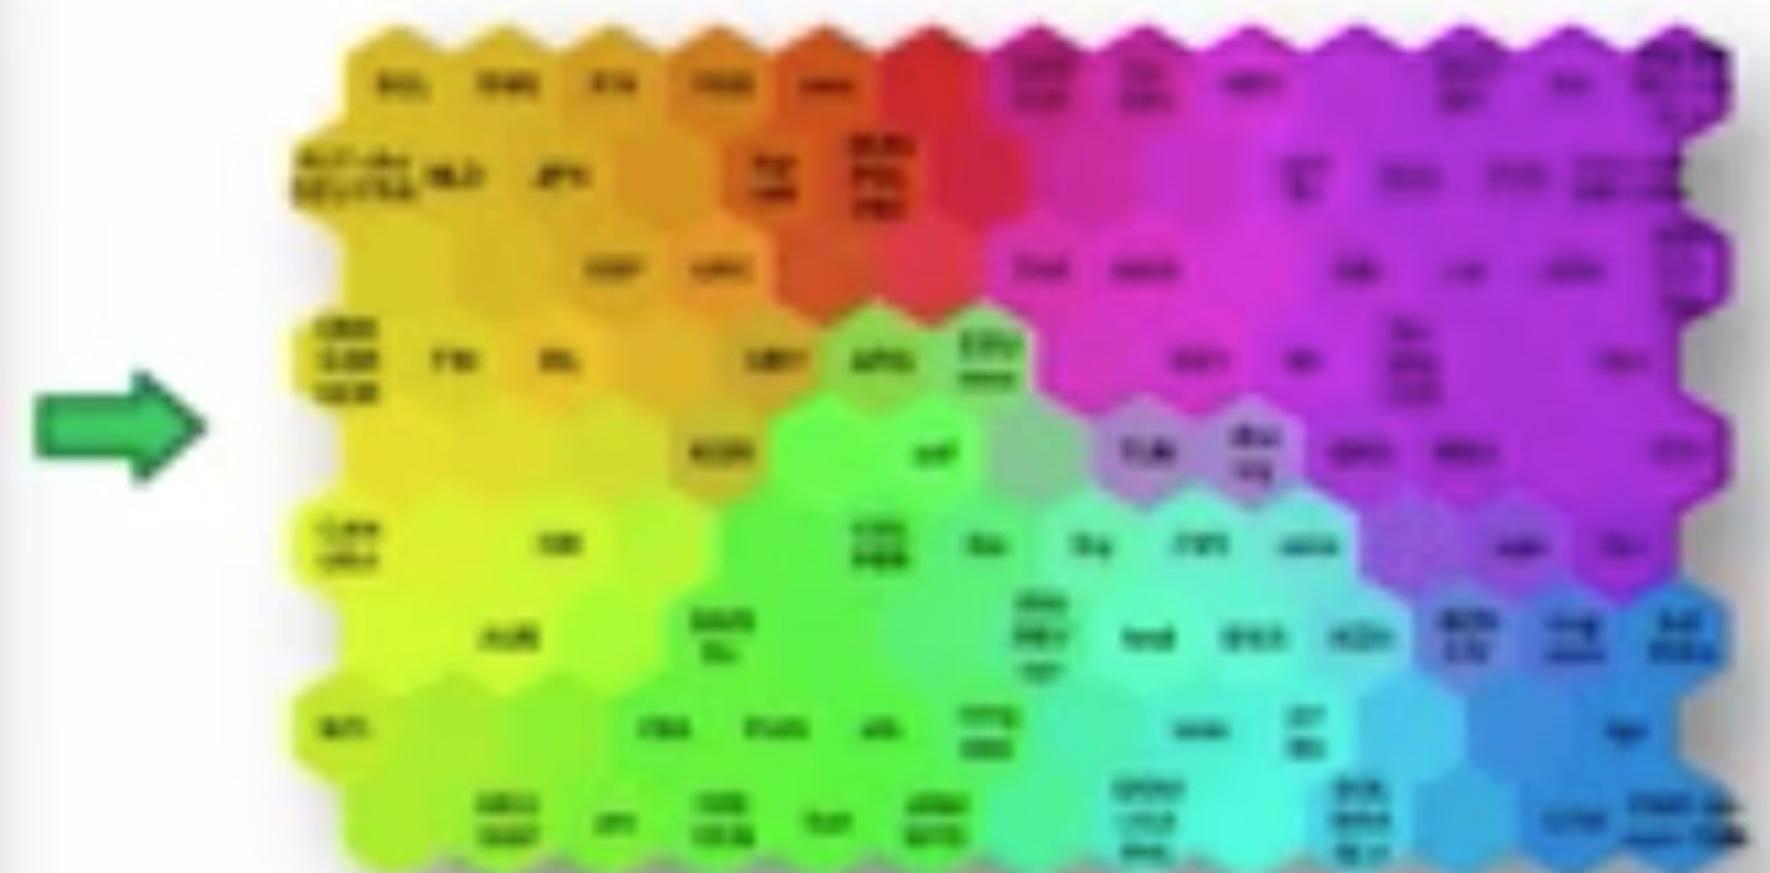


# Self Organizing Maps

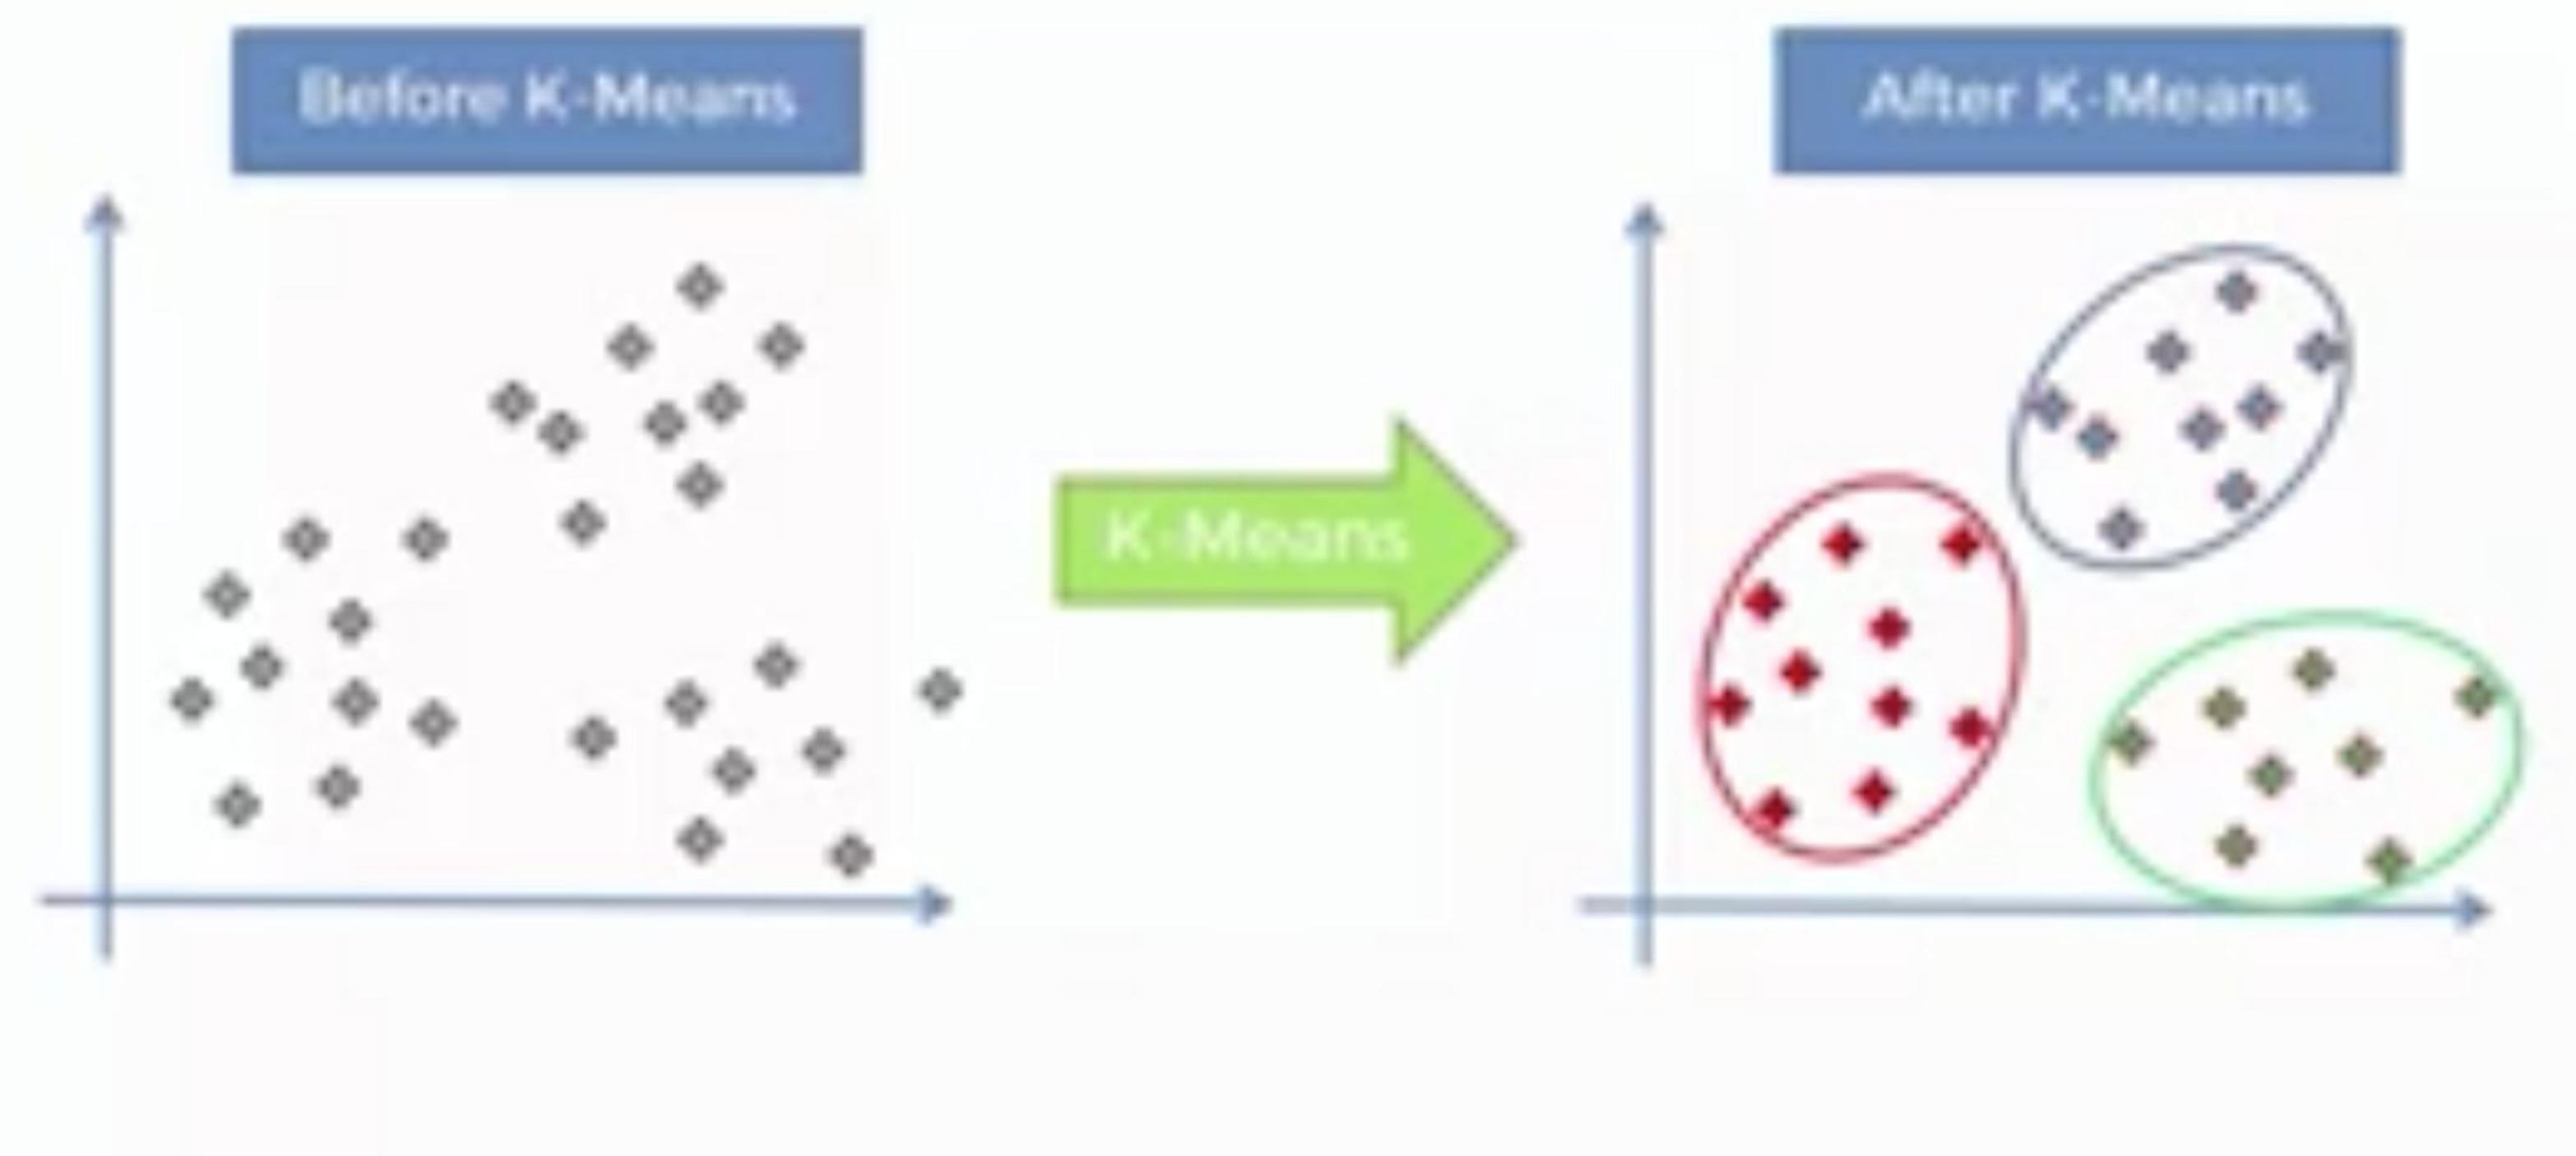


# Self Organising Maps

Rank	Country	Country ID	Health Expenditure (\$)	Inflation (%)
1	Austria	ATW	8.4328972	5.80867002
2	Australia	AUS	4.3712794	-6.28908
3	Angola	AGO	5.7812399	13.75145
4	Azerbaijan	AZB	4.73969	2.280502
5	Amazonia	AND	4.57058	5.14008705
6	Armenia	ARM	4.049924	5.324854
7	Aruba	ABW	7.884758	
8	United Arab Emirates	ARE	4.840323	4.888807986
9	Argentina	ARG	4.840323	4.282774
10	Armenia	ARM	3.84079005	5.408767
11	American Samoa	ASM	4.840323	
12	Antigua & Barbuda	ATG	9.040556	2.504407058
13	Australia	AUS	5.394444	5.80867002
14	Austria	AUT	5.80867002	0.308113
15	Azerbaijan	AZB	4.944387	1.254009992
16	Burundi	BDI	10.39404	6.30879889
17	Brunei Darussalam	BRU	4.484032	6.402259807
18	Bosnia	BHR	7.408431	4.200940518



# K-Means ~ Self Organising Maps



# K-Means ~ Self Organising Maps

STEP 1: Choose the number K of clusters



STEP 2: Select at random K points, the centroids (not necessarily from your dataset)



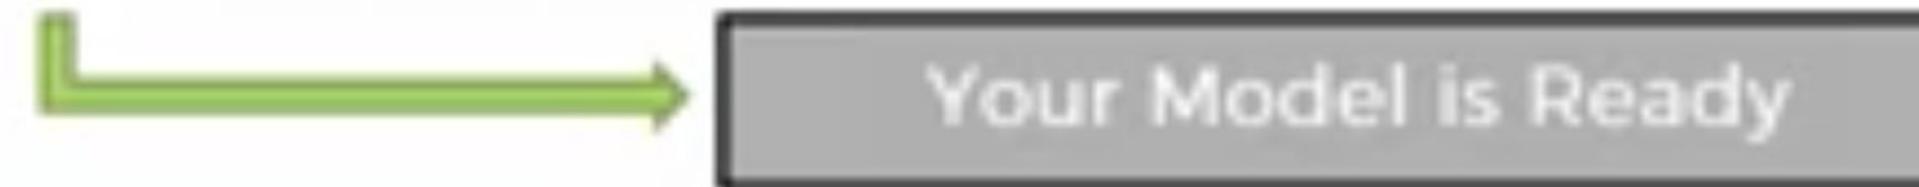
STEP 3: Assign each data point to the closest centroid ➔ That forms K clusters



STEP 4: Compute and place the new centroid of each cluster

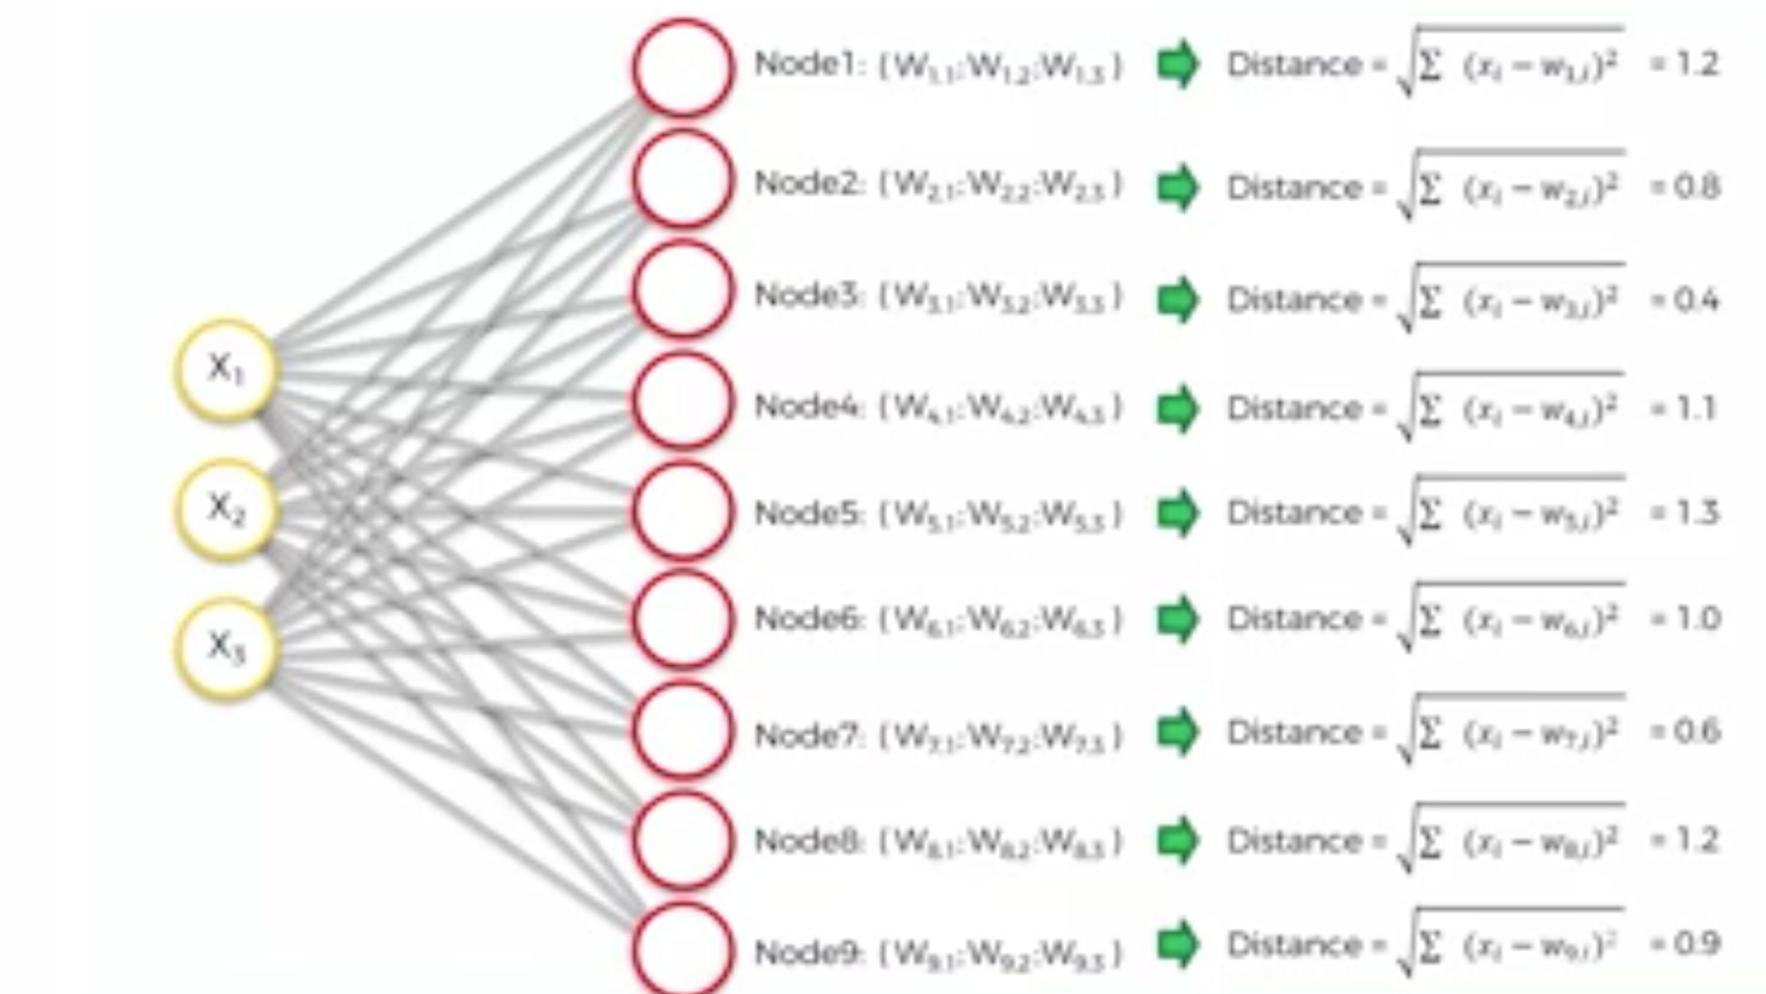
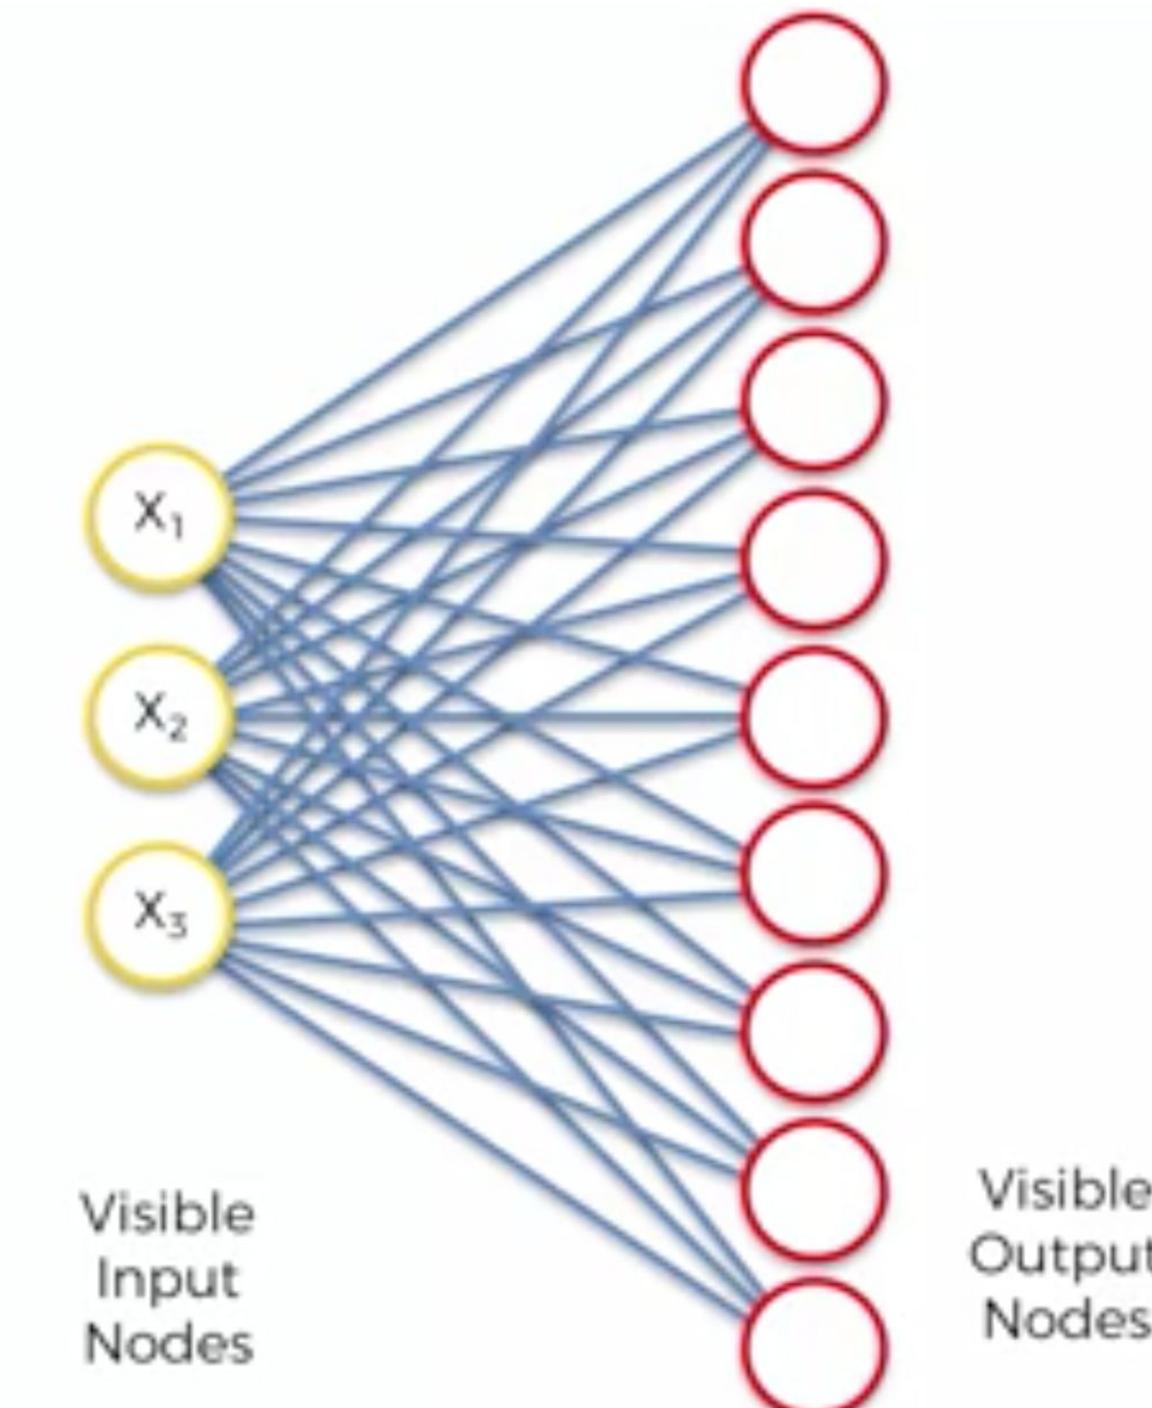
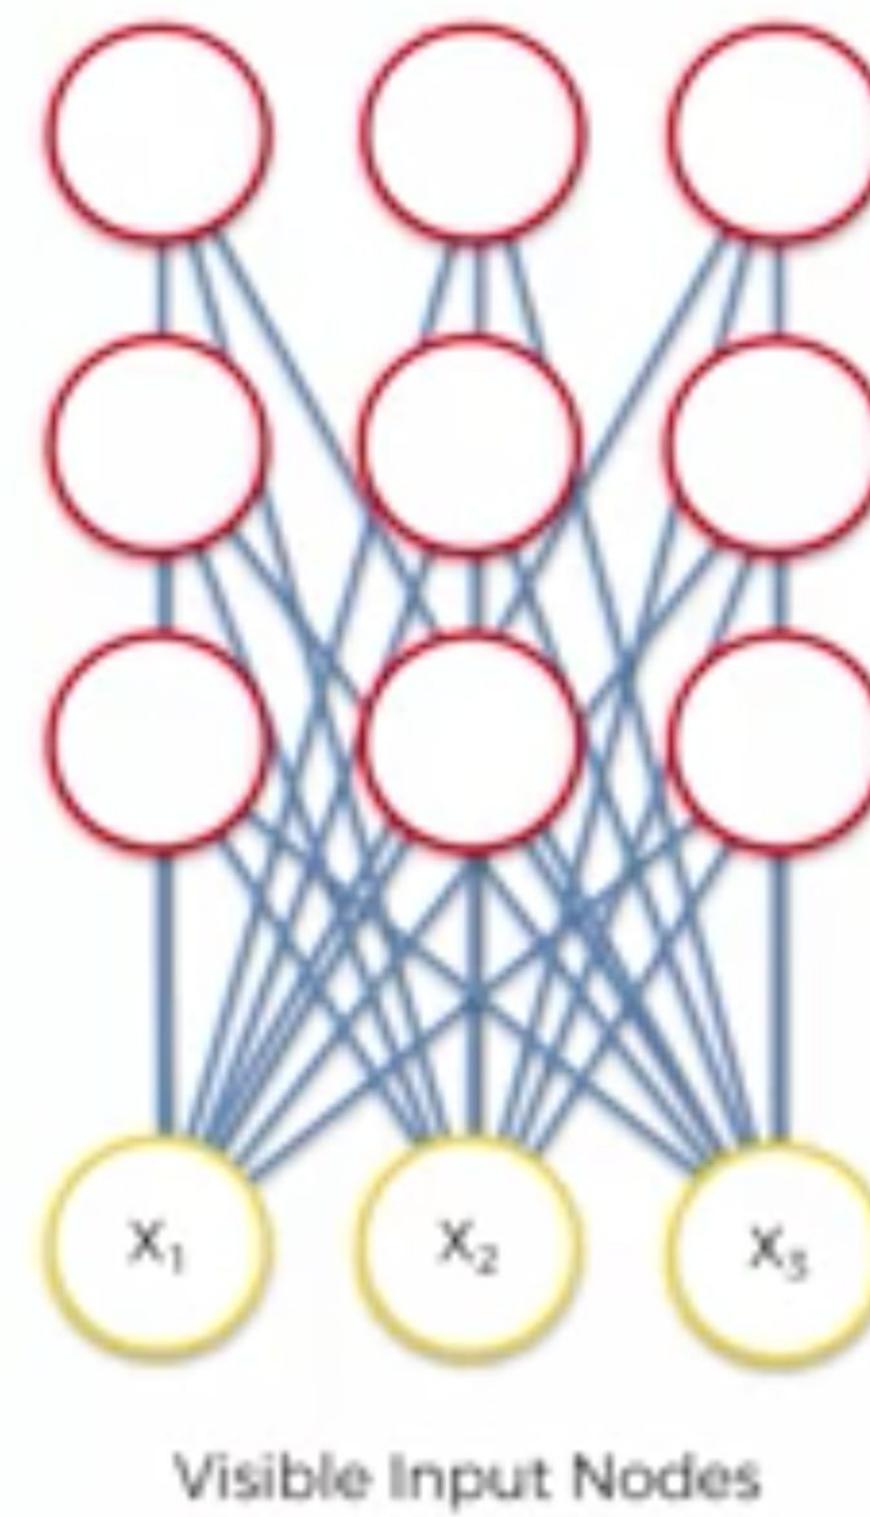


STEP 5: Reassign each data point to the new closest centroid.  
If any reassignment took place, go to STEP 4, otherwise go to FIN.

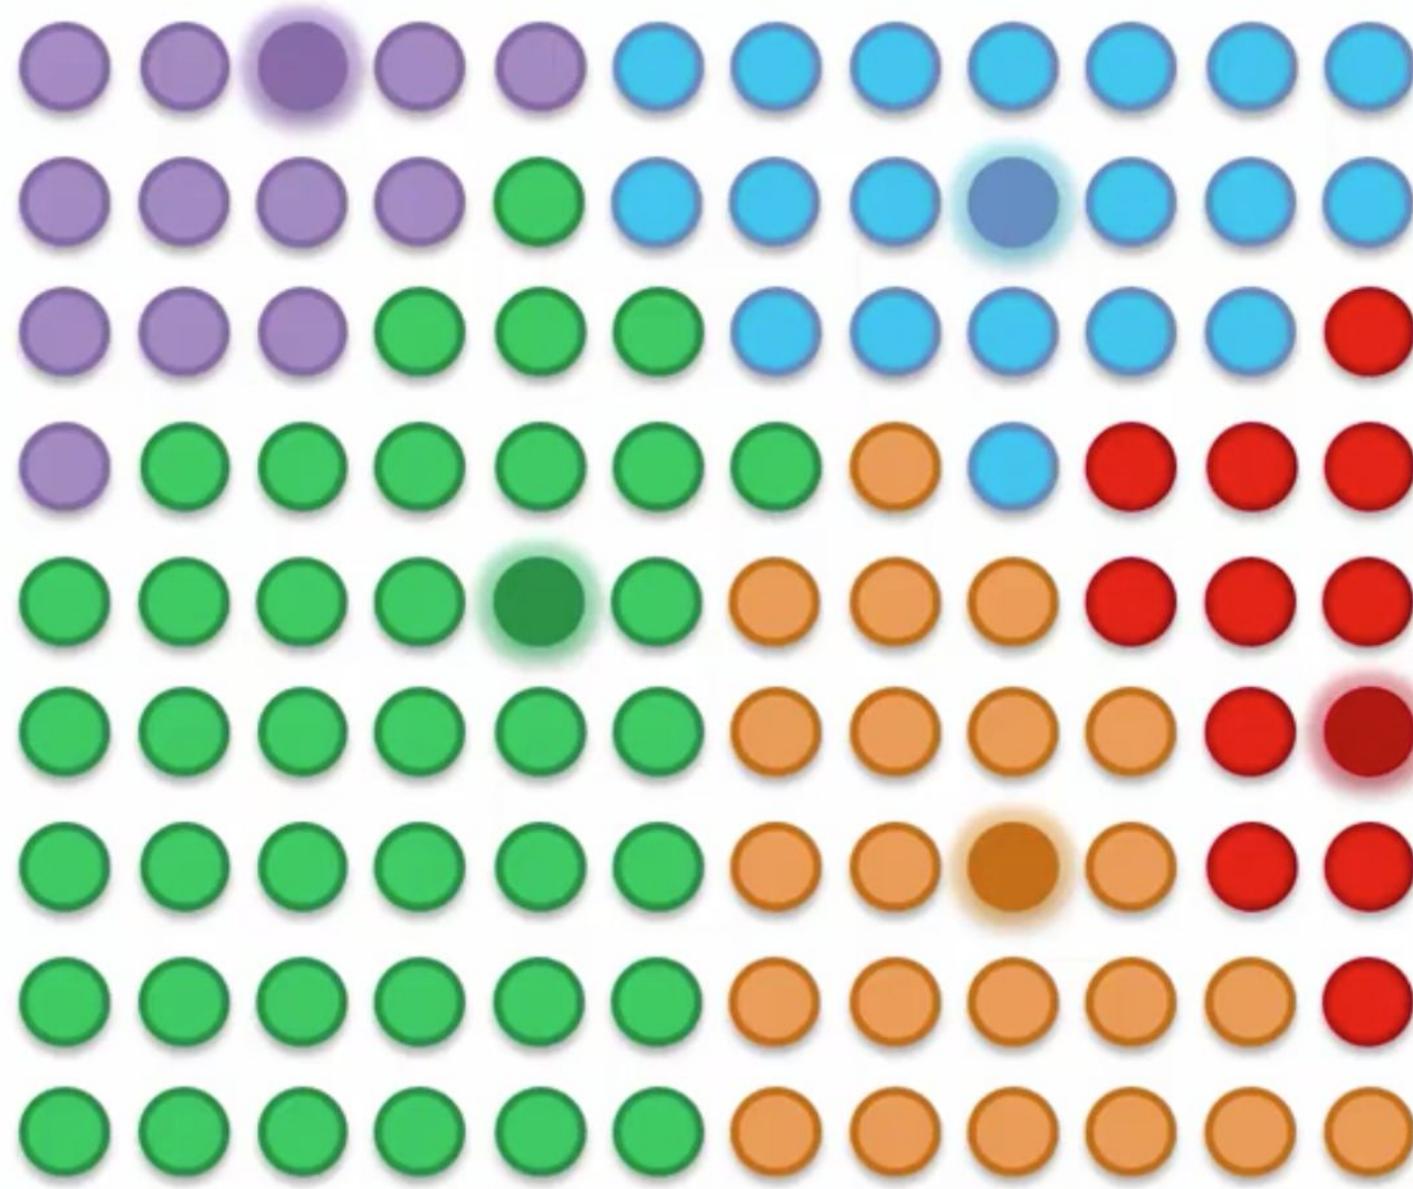
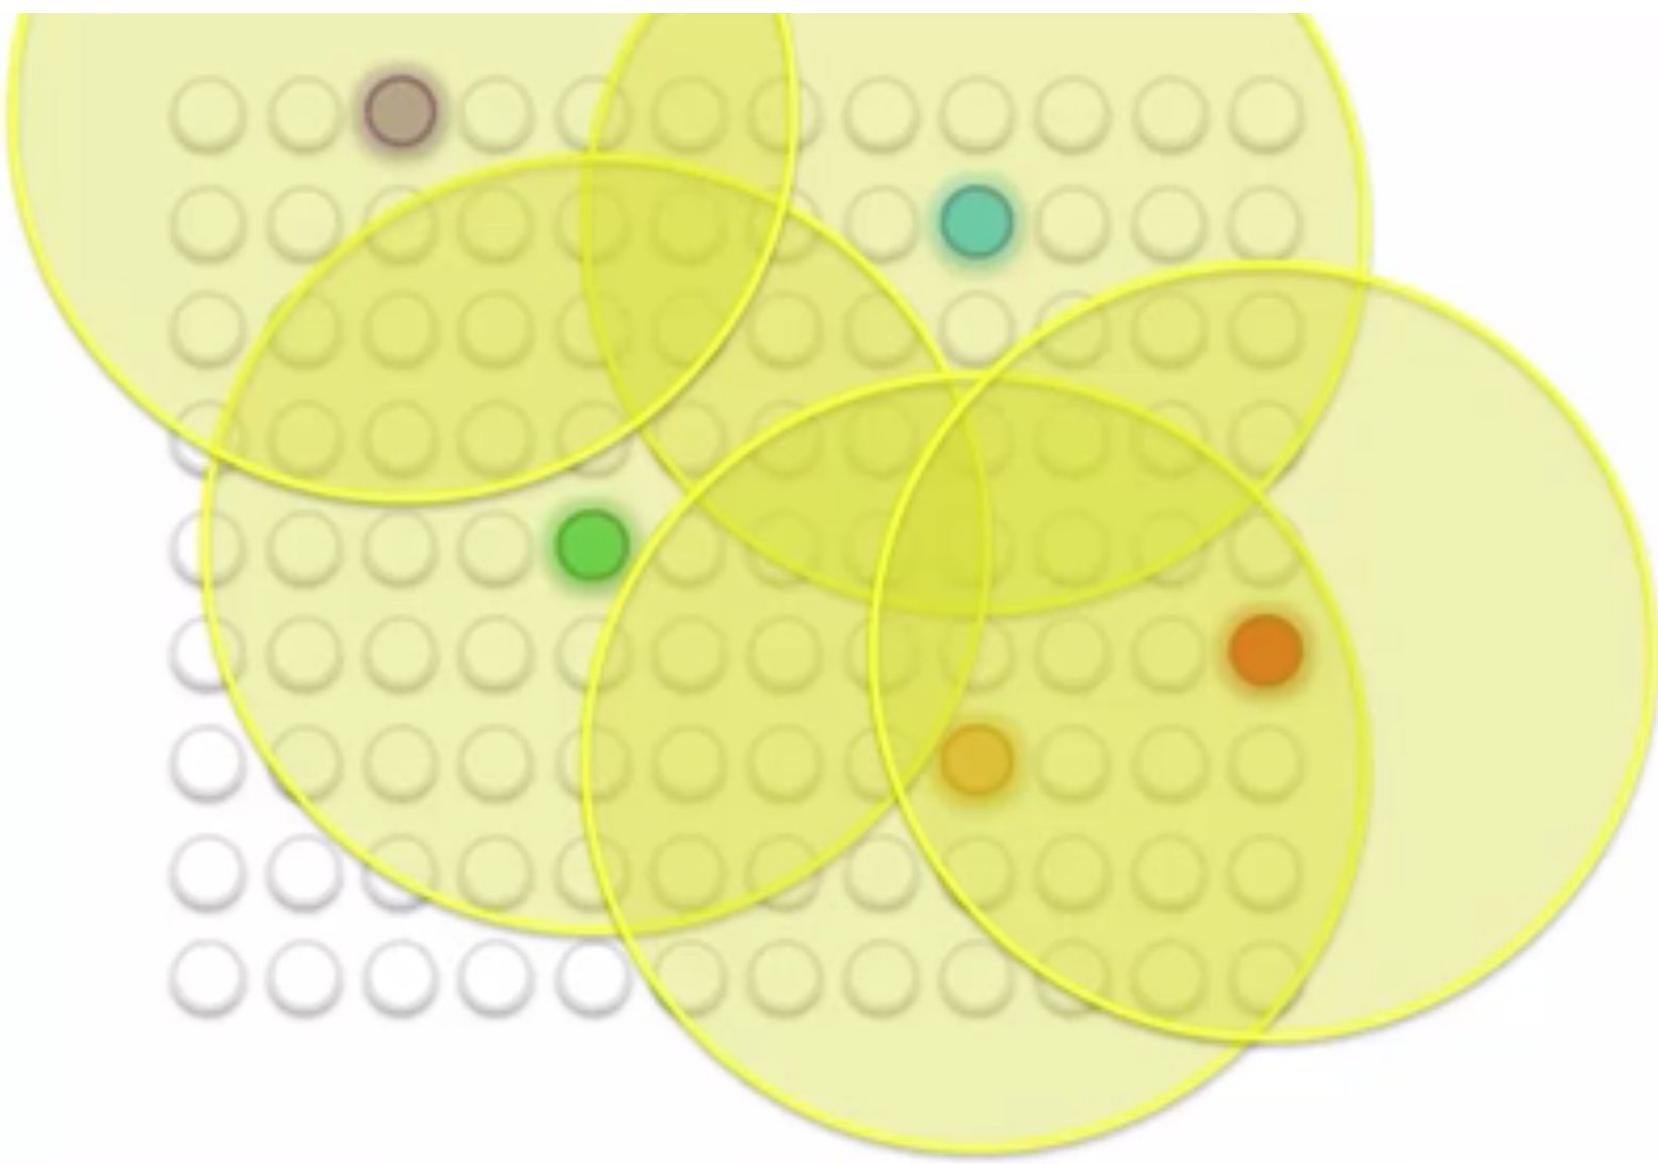


Your Model is Ready

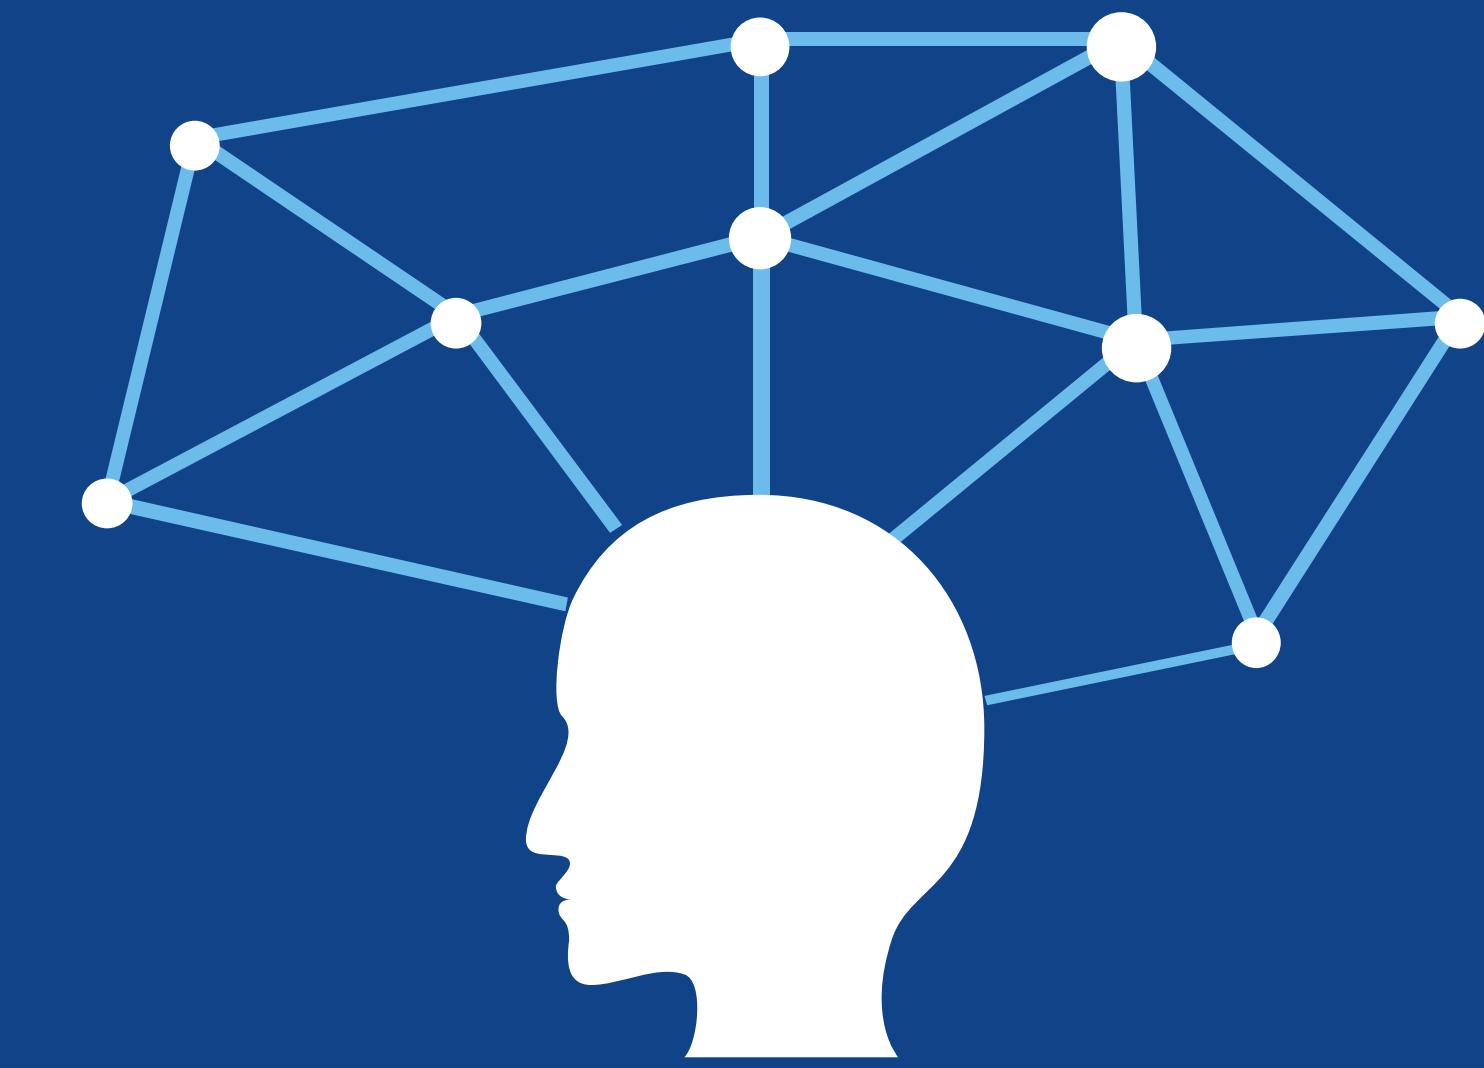
# Self Organising Maps



# Self Organising Maps



# Boltzmann Machines



# Boltzmann Machines

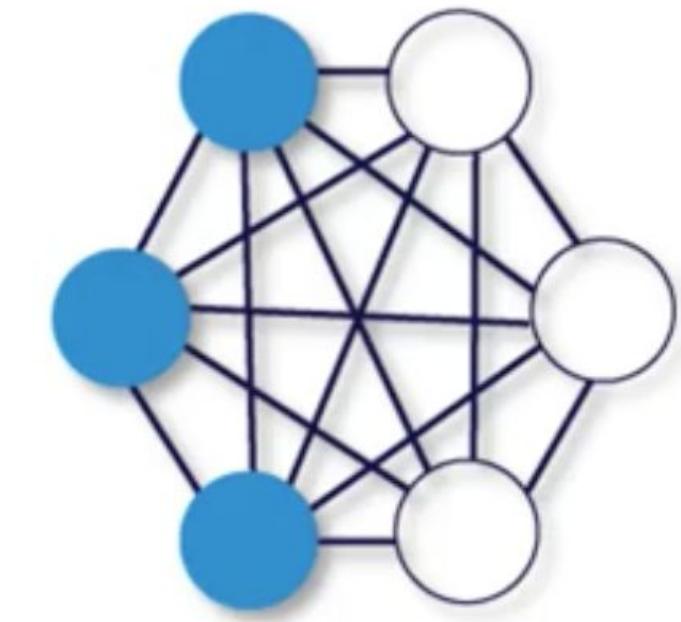
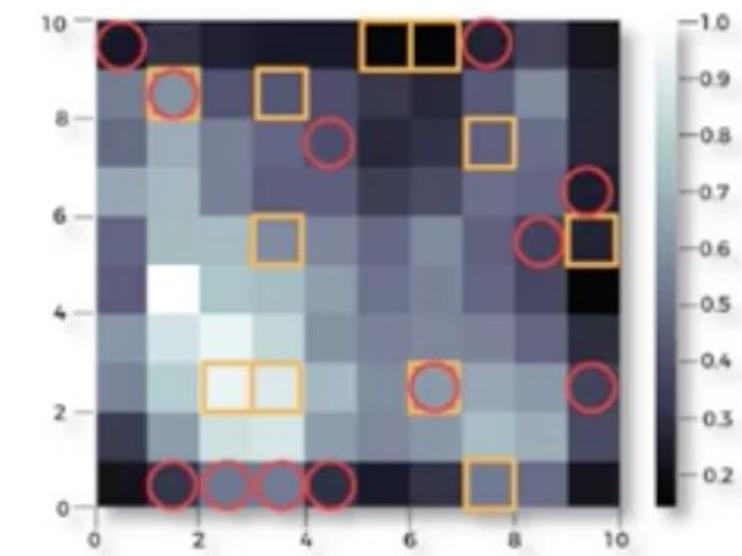
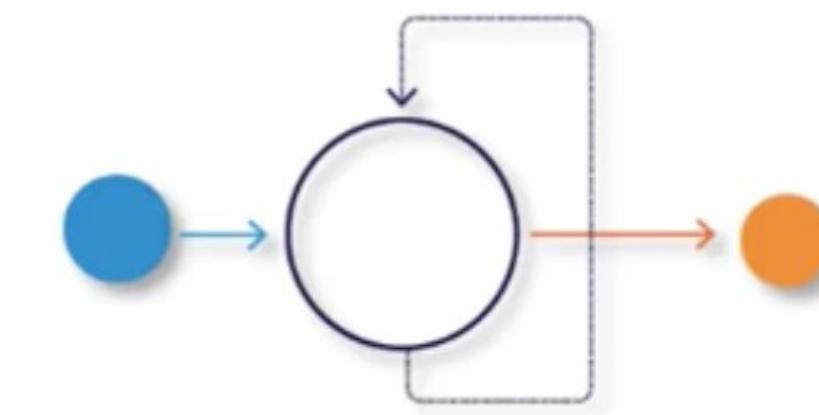
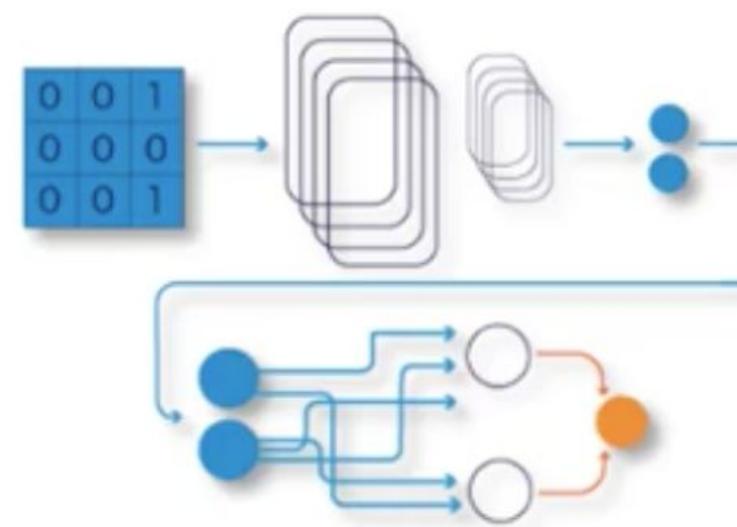
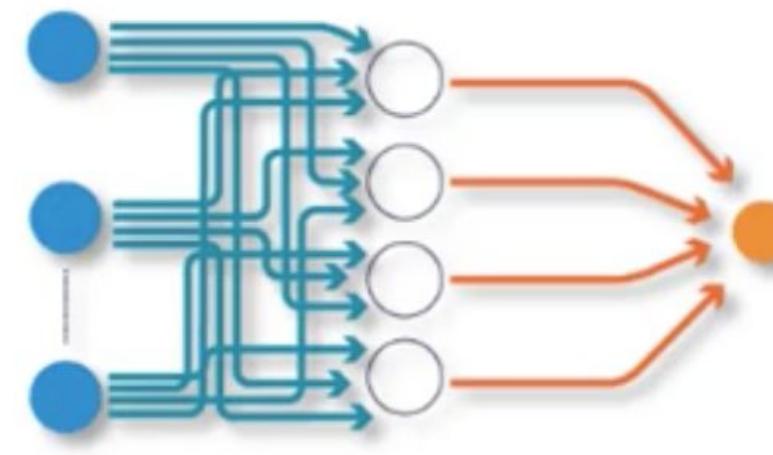
Restricted Boltzmann machines for collaborative filtering

Authors: Ruslan Salakhutdinov University of Toronto, Toronto, Ontario, Canada  
Andriy Mnih University of Toronto, Toronto, Ontario, Canada  
Geoffrey Hinton University of Toronto, Toronto, Ontario, Canada

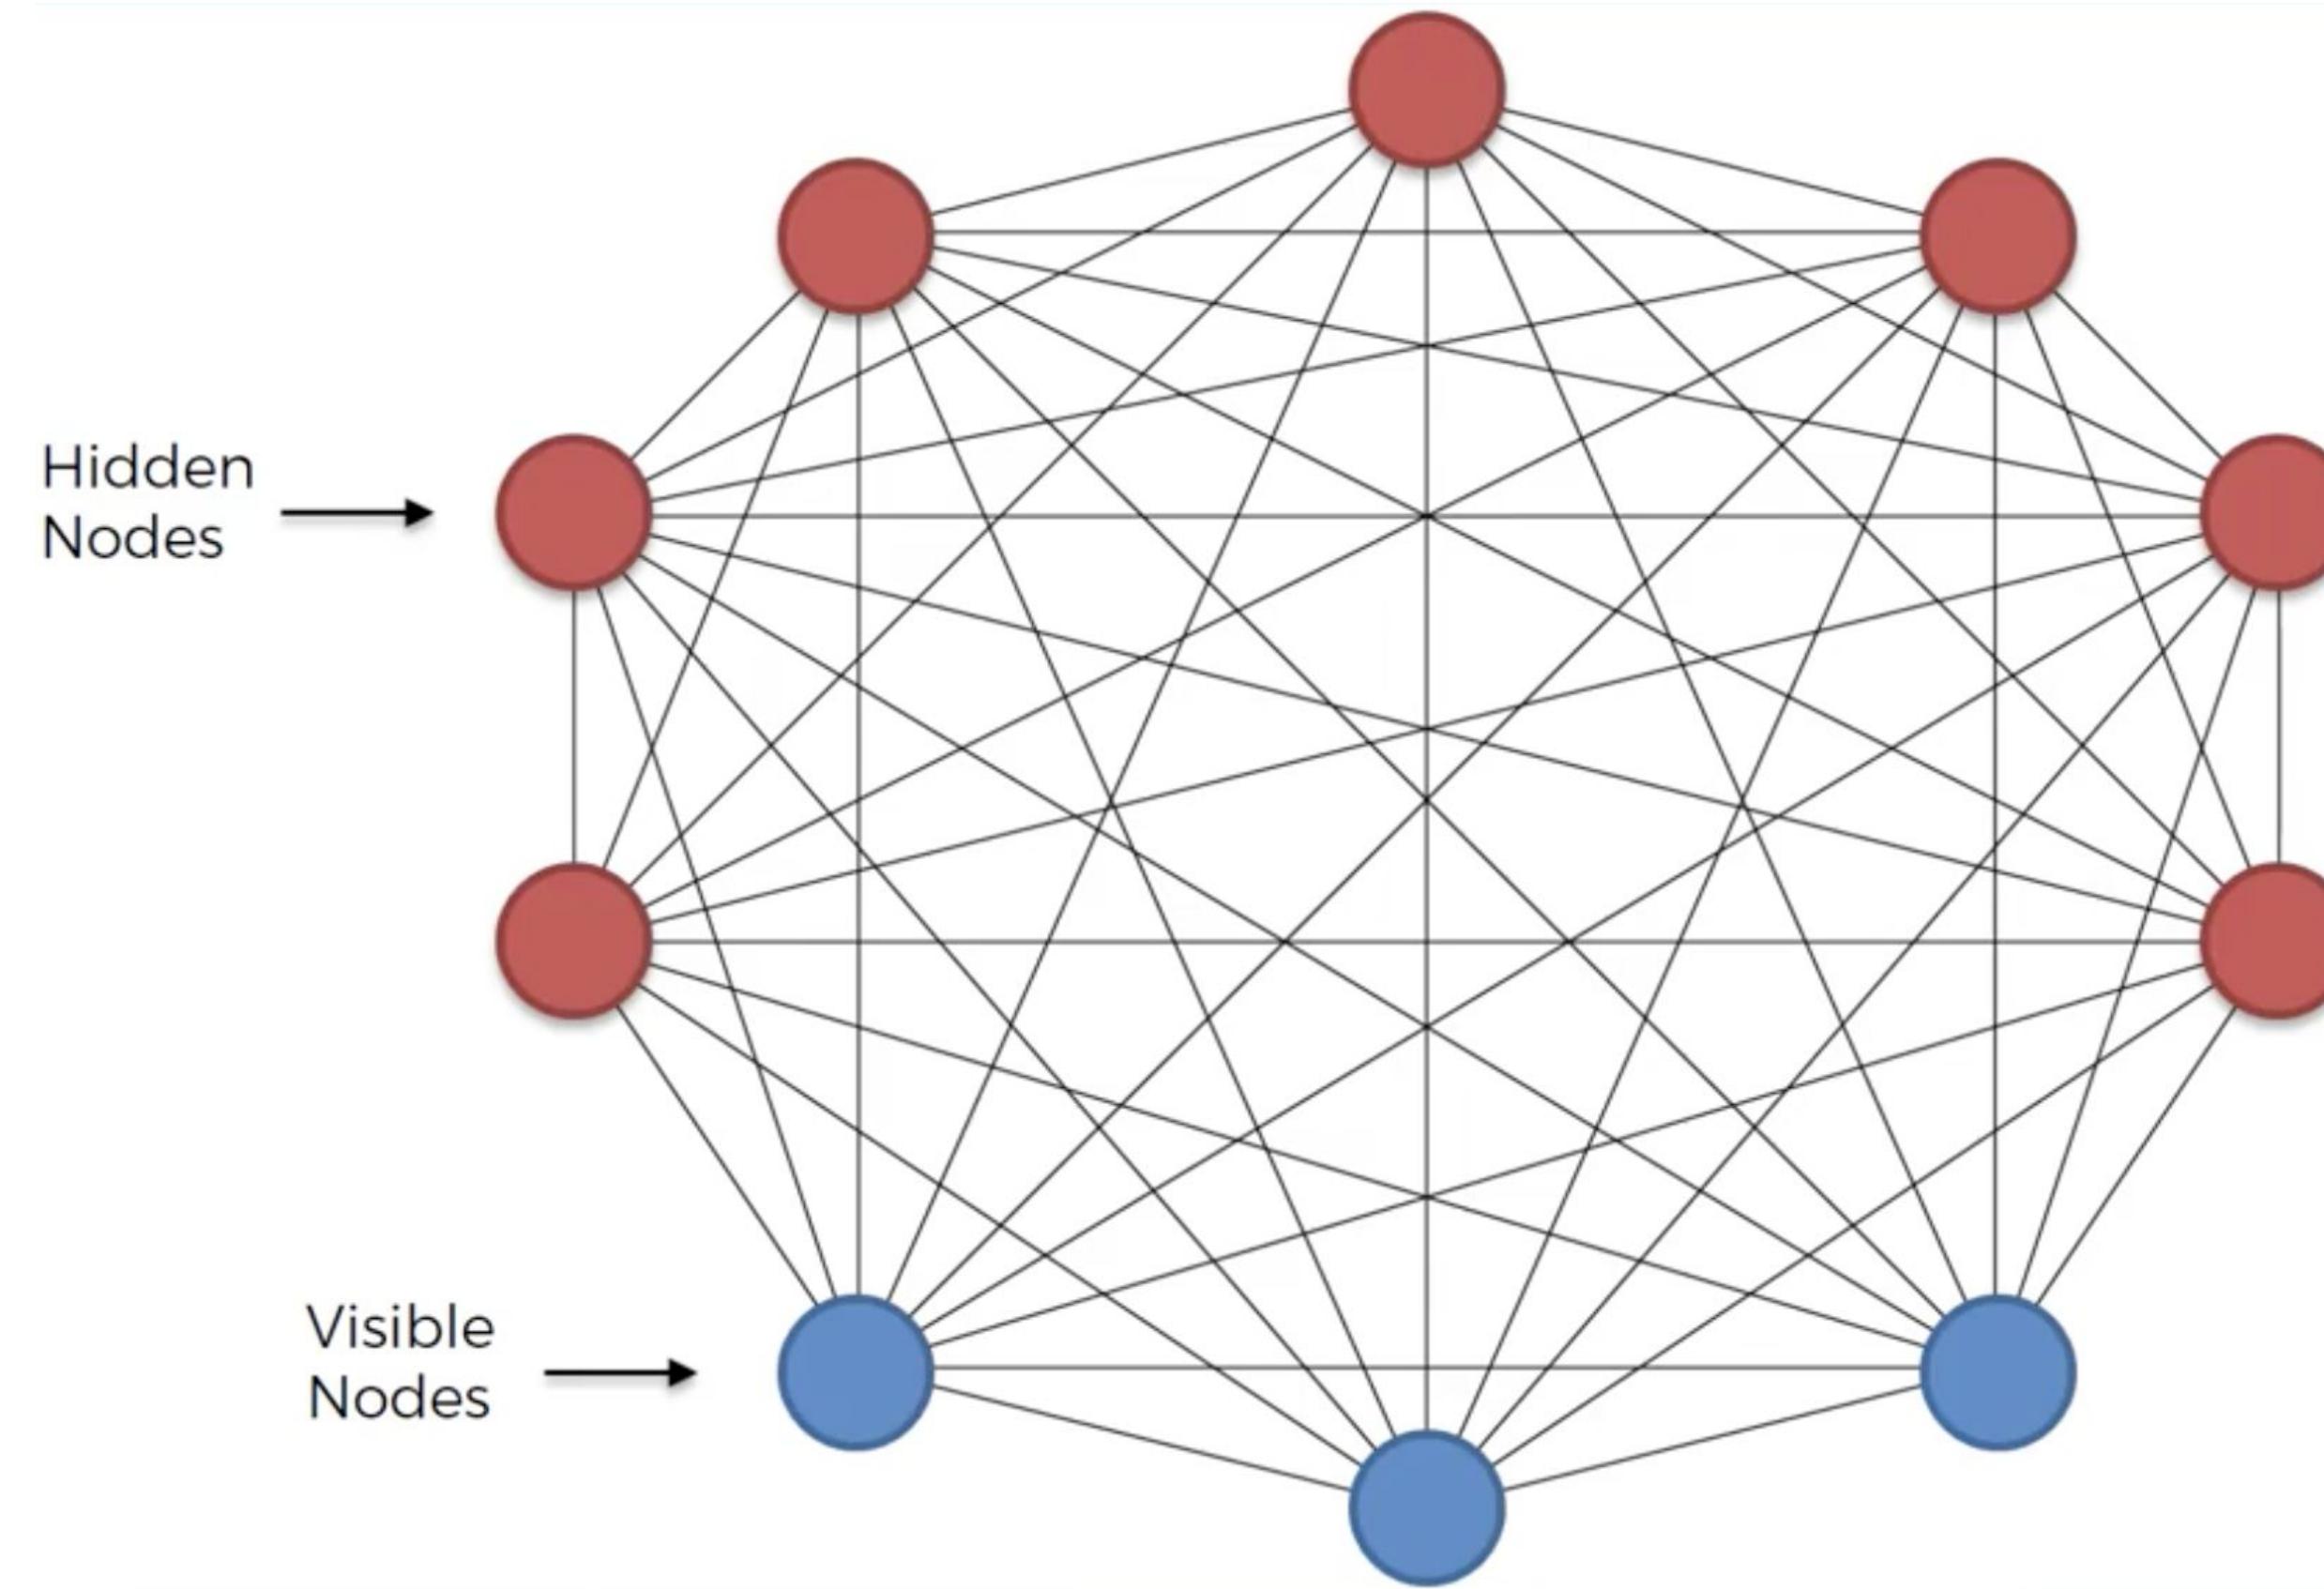
Most of the existing approaches to collaborative filtering cannot handle very large data sets. In this paper we show how a class of two-layer undirected graphical models, called Restricted Boltzmann Machines (RBM's), can be used to model tabular data, such as user's ratings of movies. We present efficient learning and inference procedures for this class of models and demonstrate that RBM's can be successfully applied to the Netflix data set, containing over 100 million user/movie ratings. We also show that RBM's slightly outperform carefully-tuned SVD models. When the predictions of multiple RBM models and multiple SVD models are linearly combined, we achieve an error rate that is well over 6% better than the score of Netflix's own system.



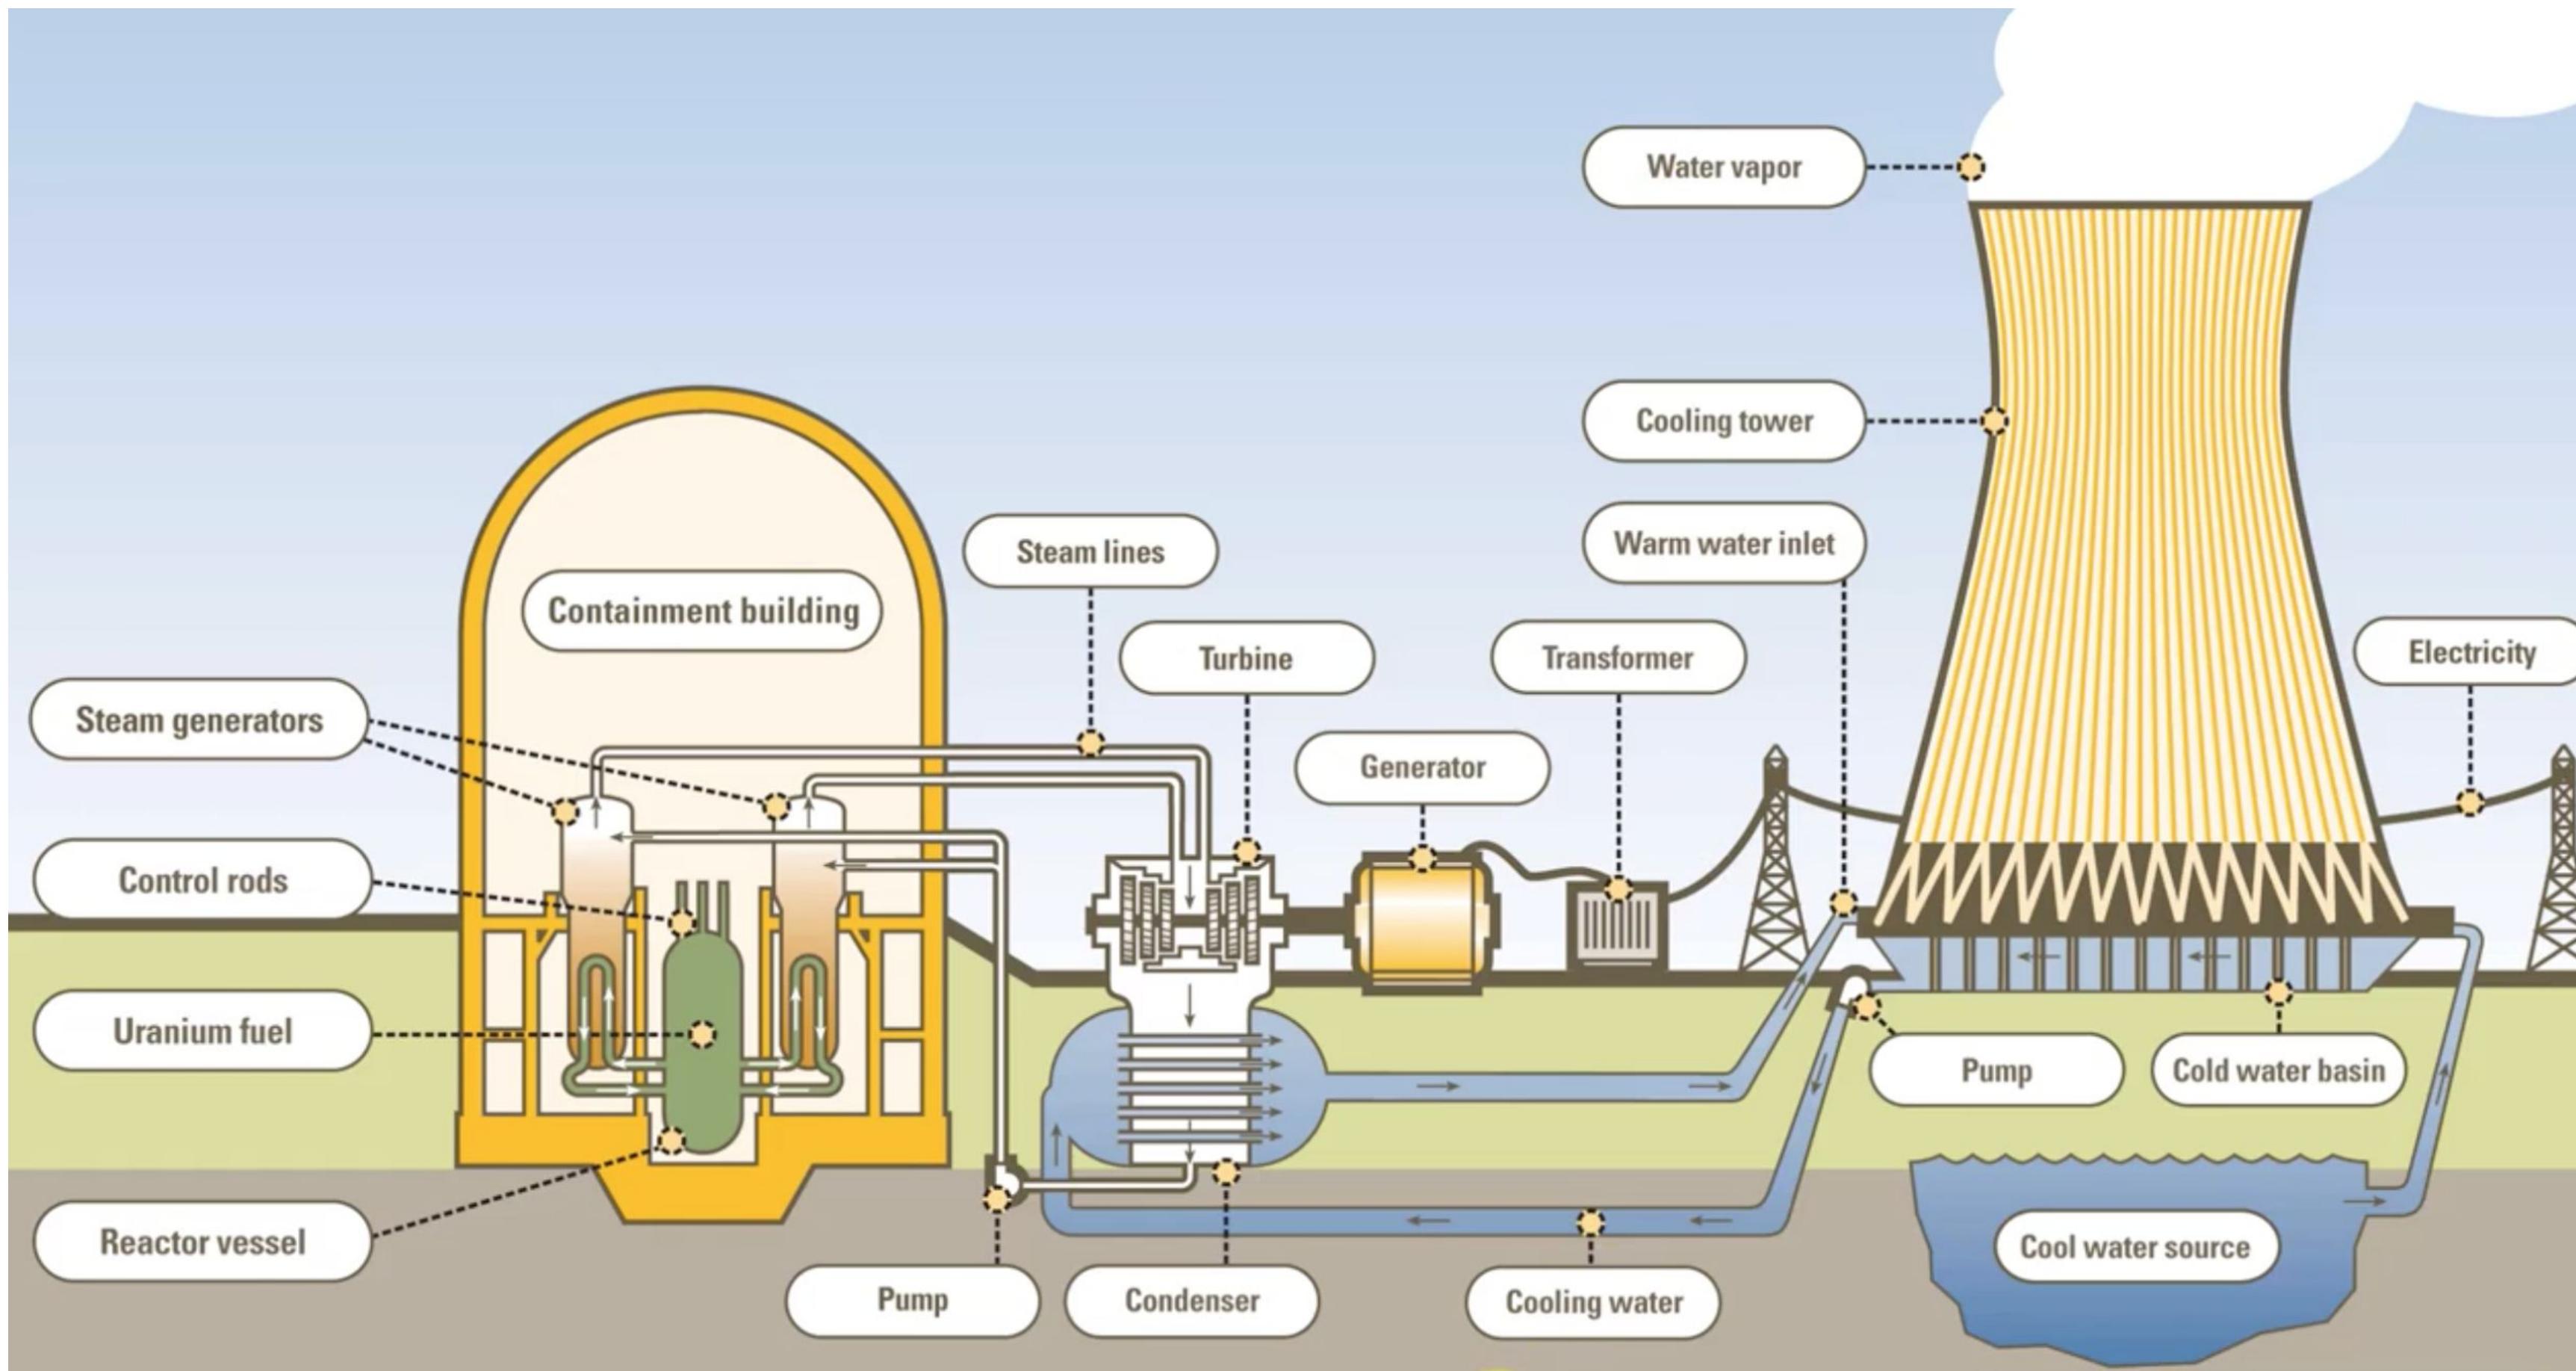
# Boltzmann Machines



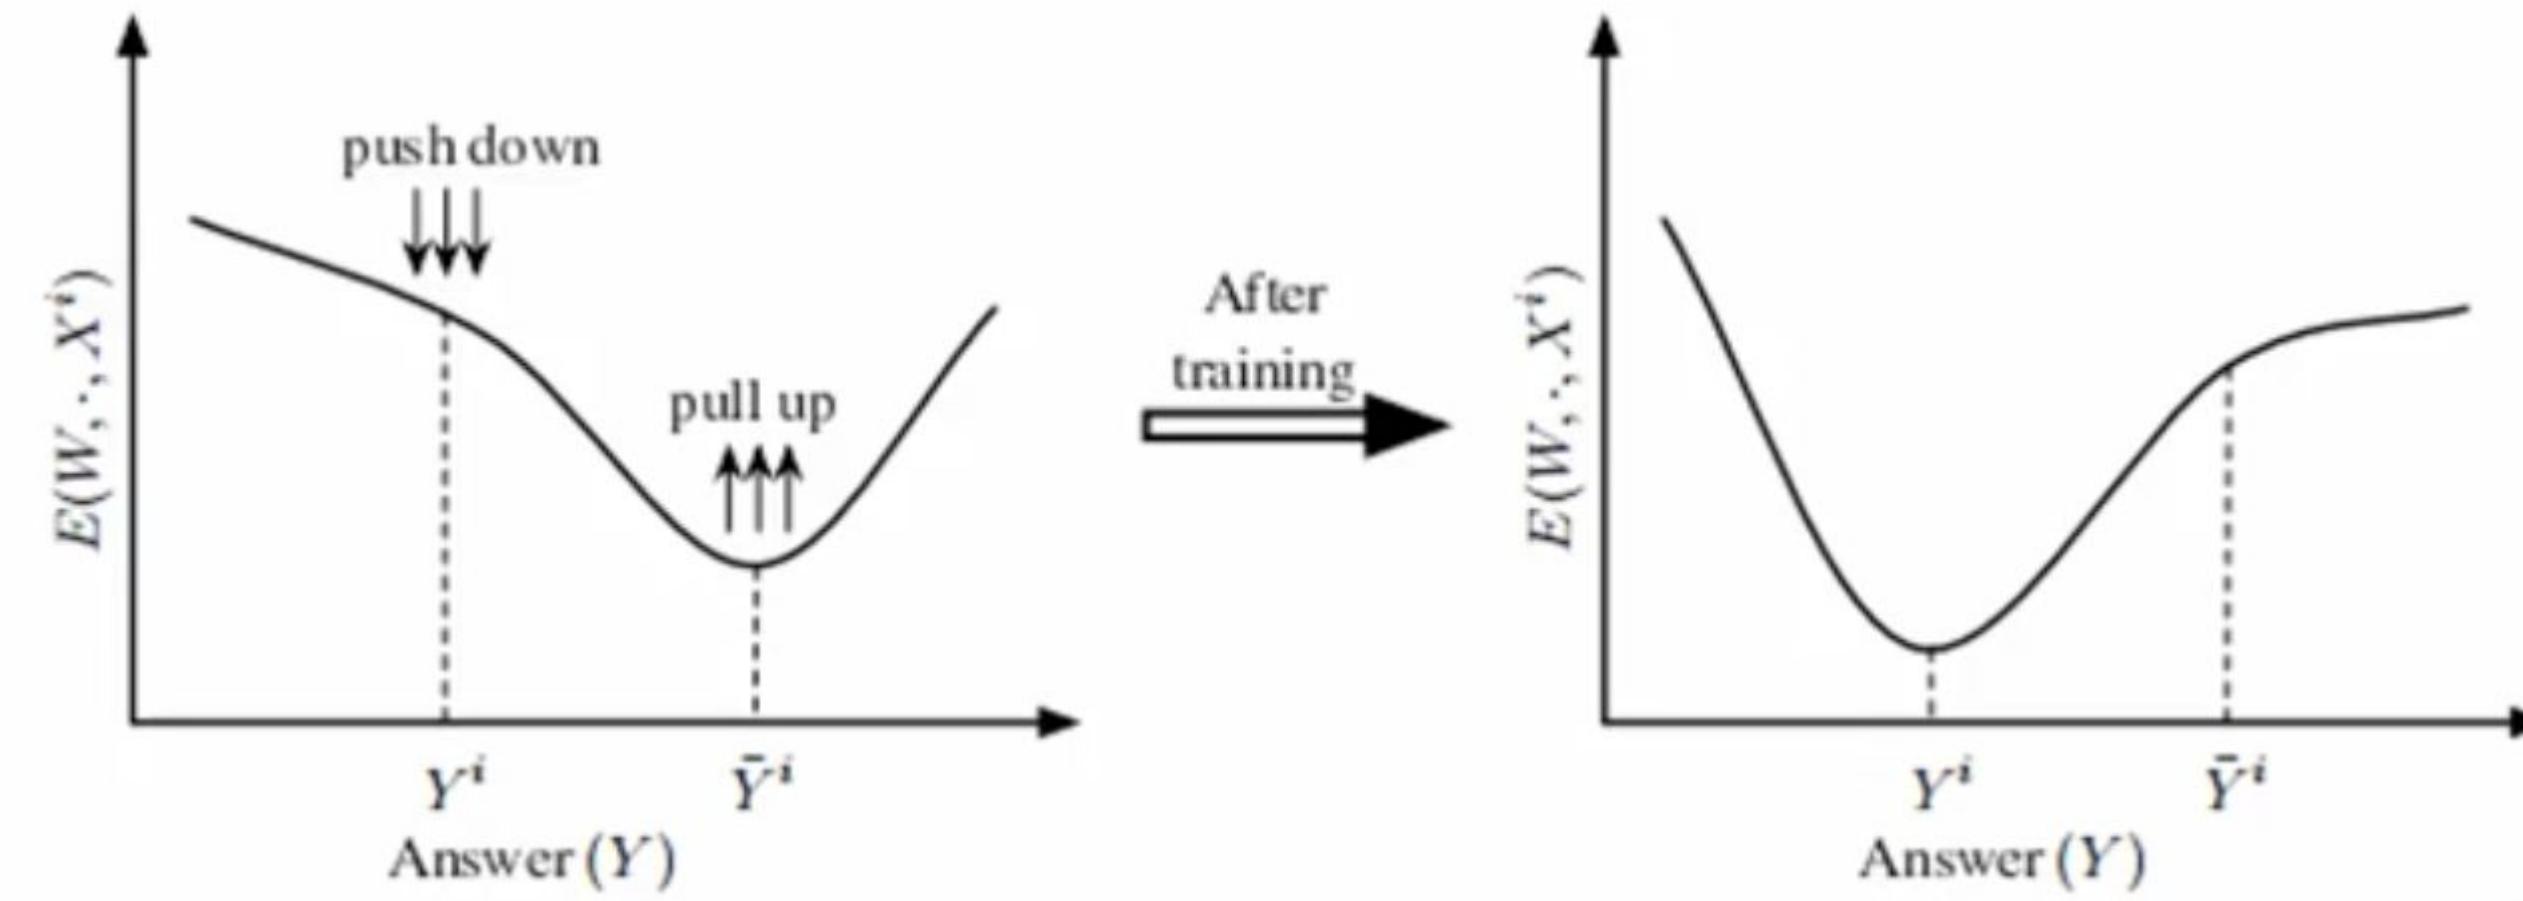
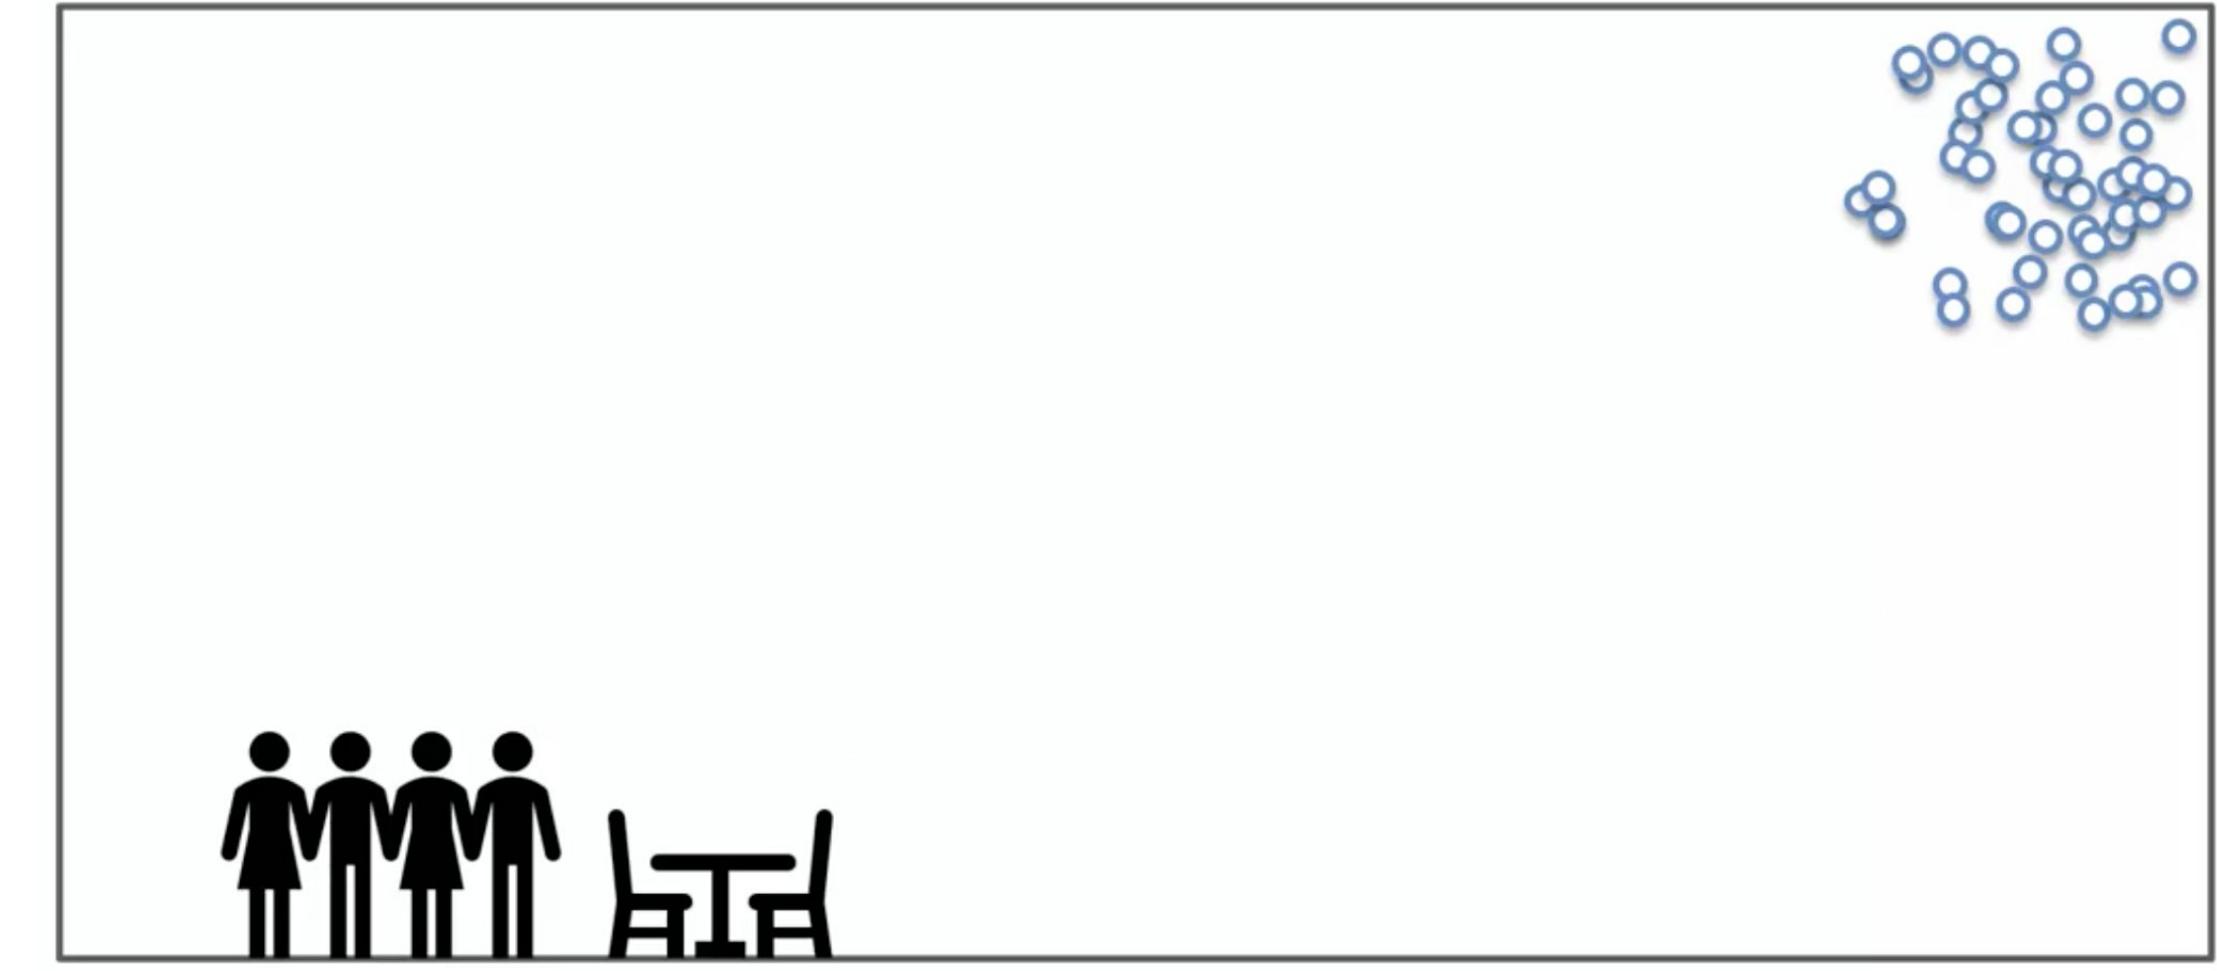
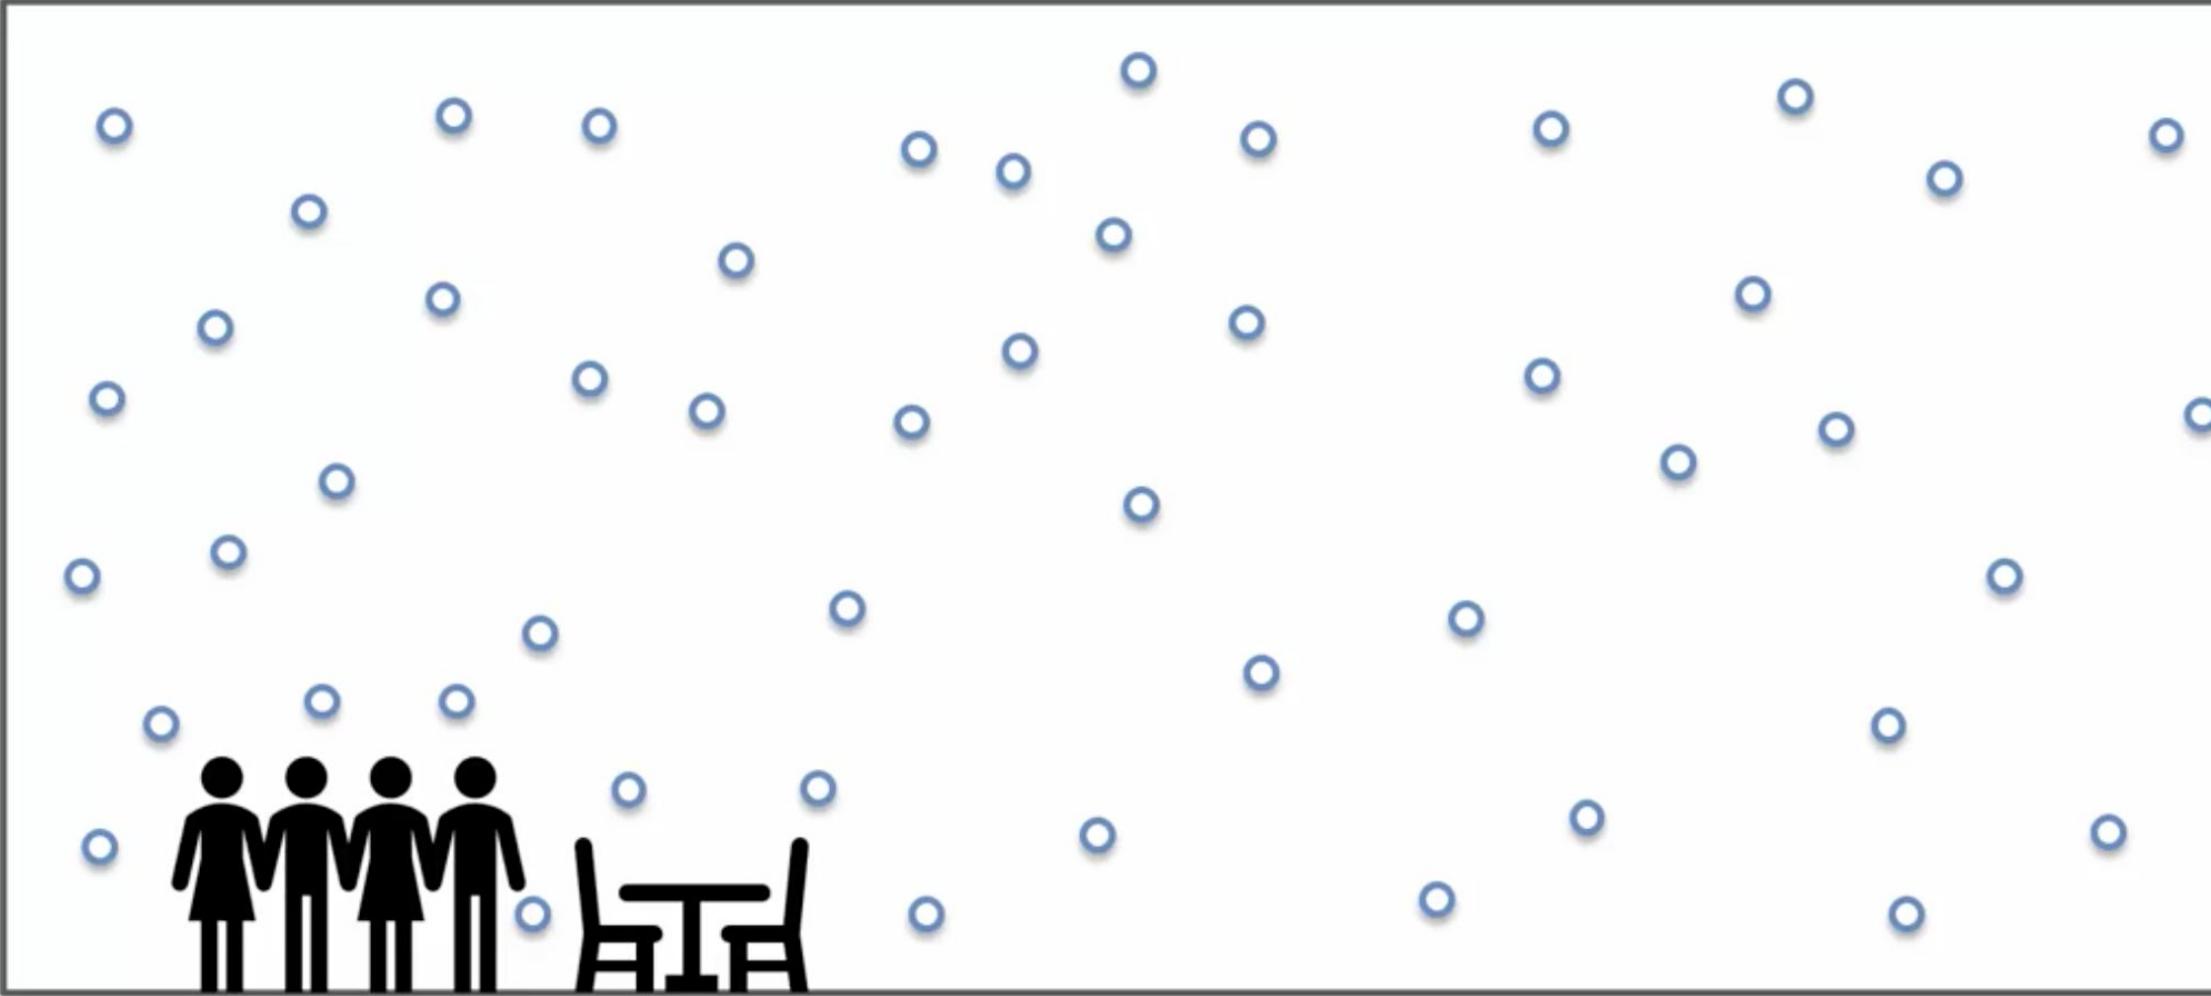
# Boltzmann Machines



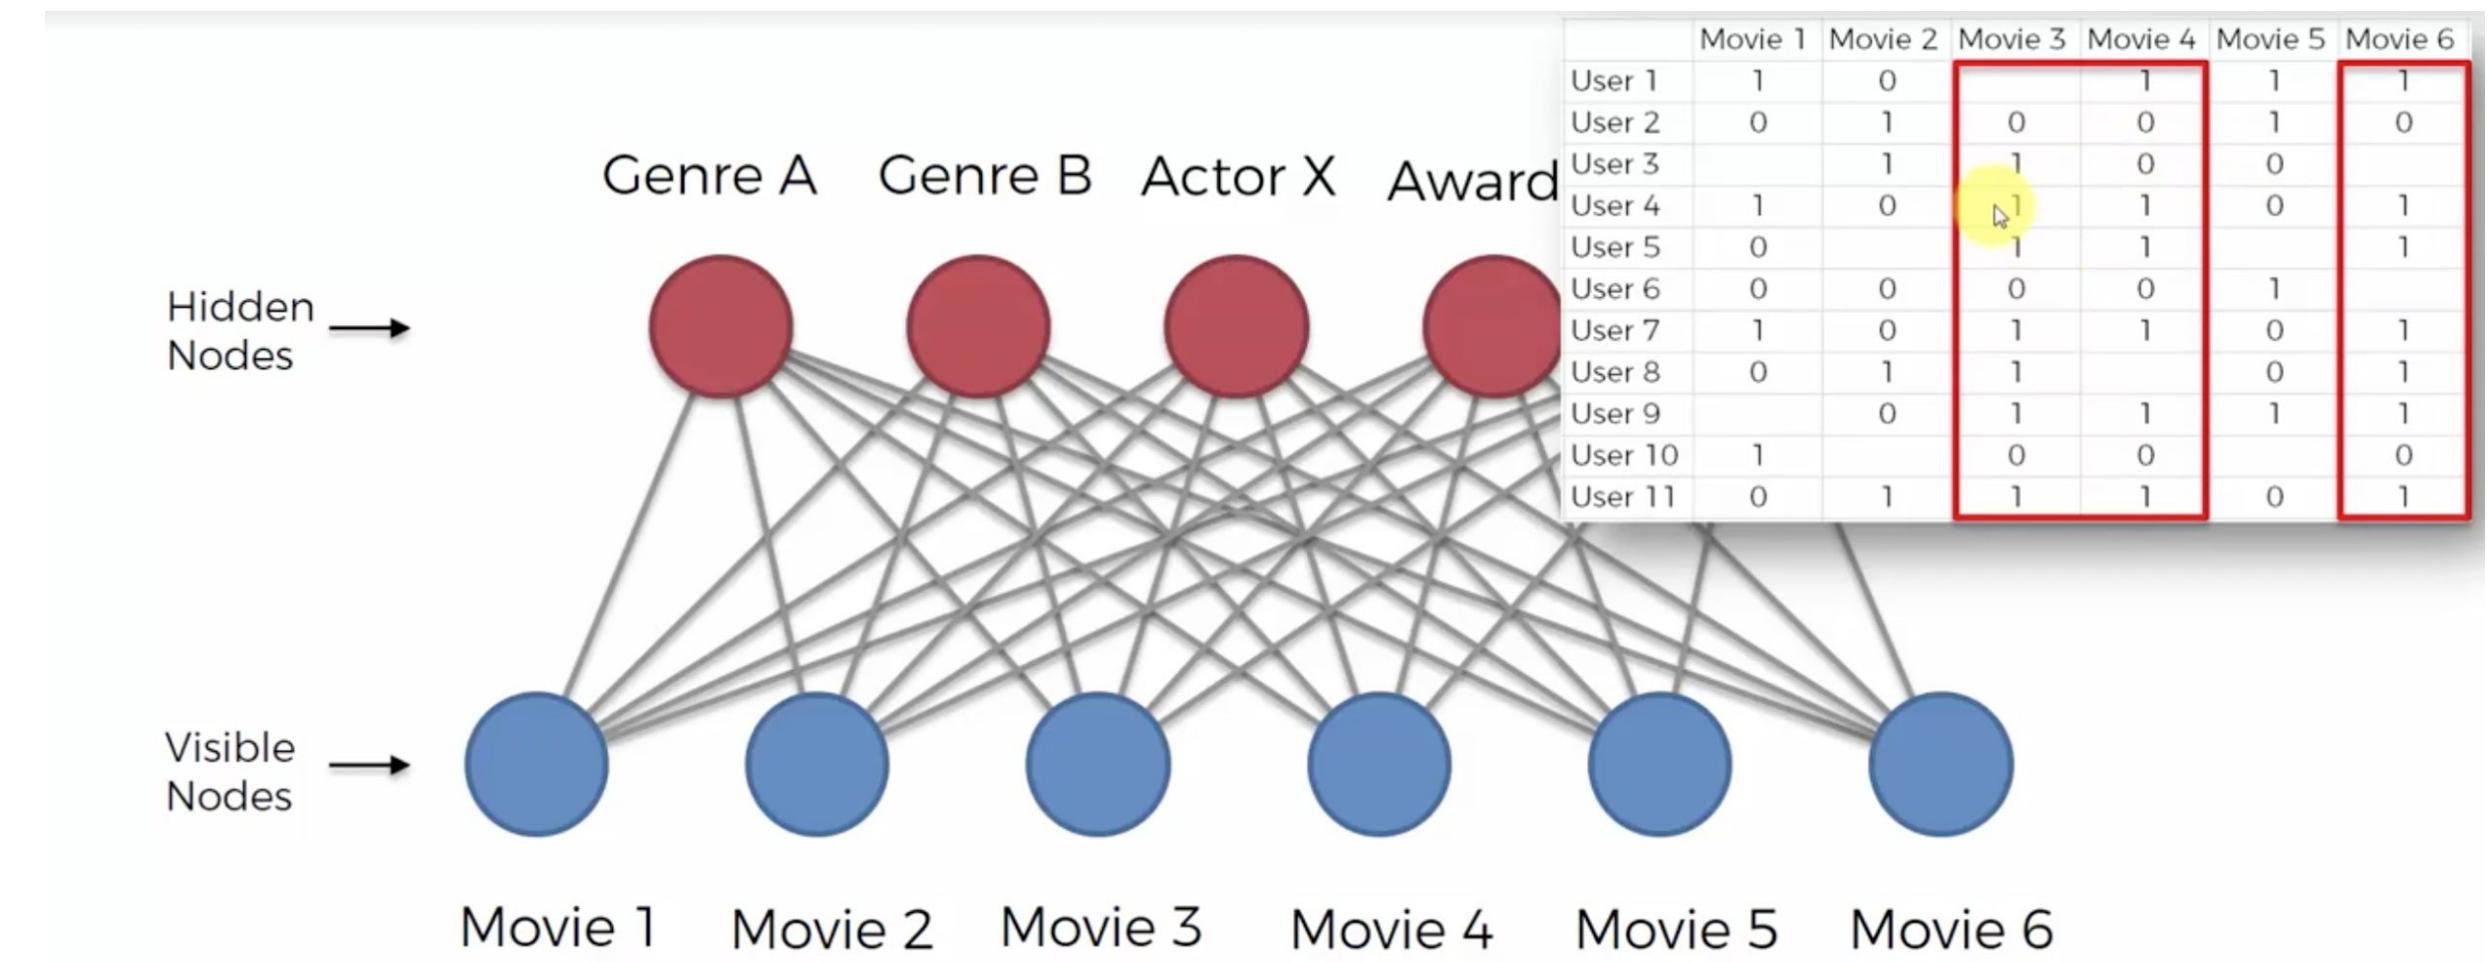
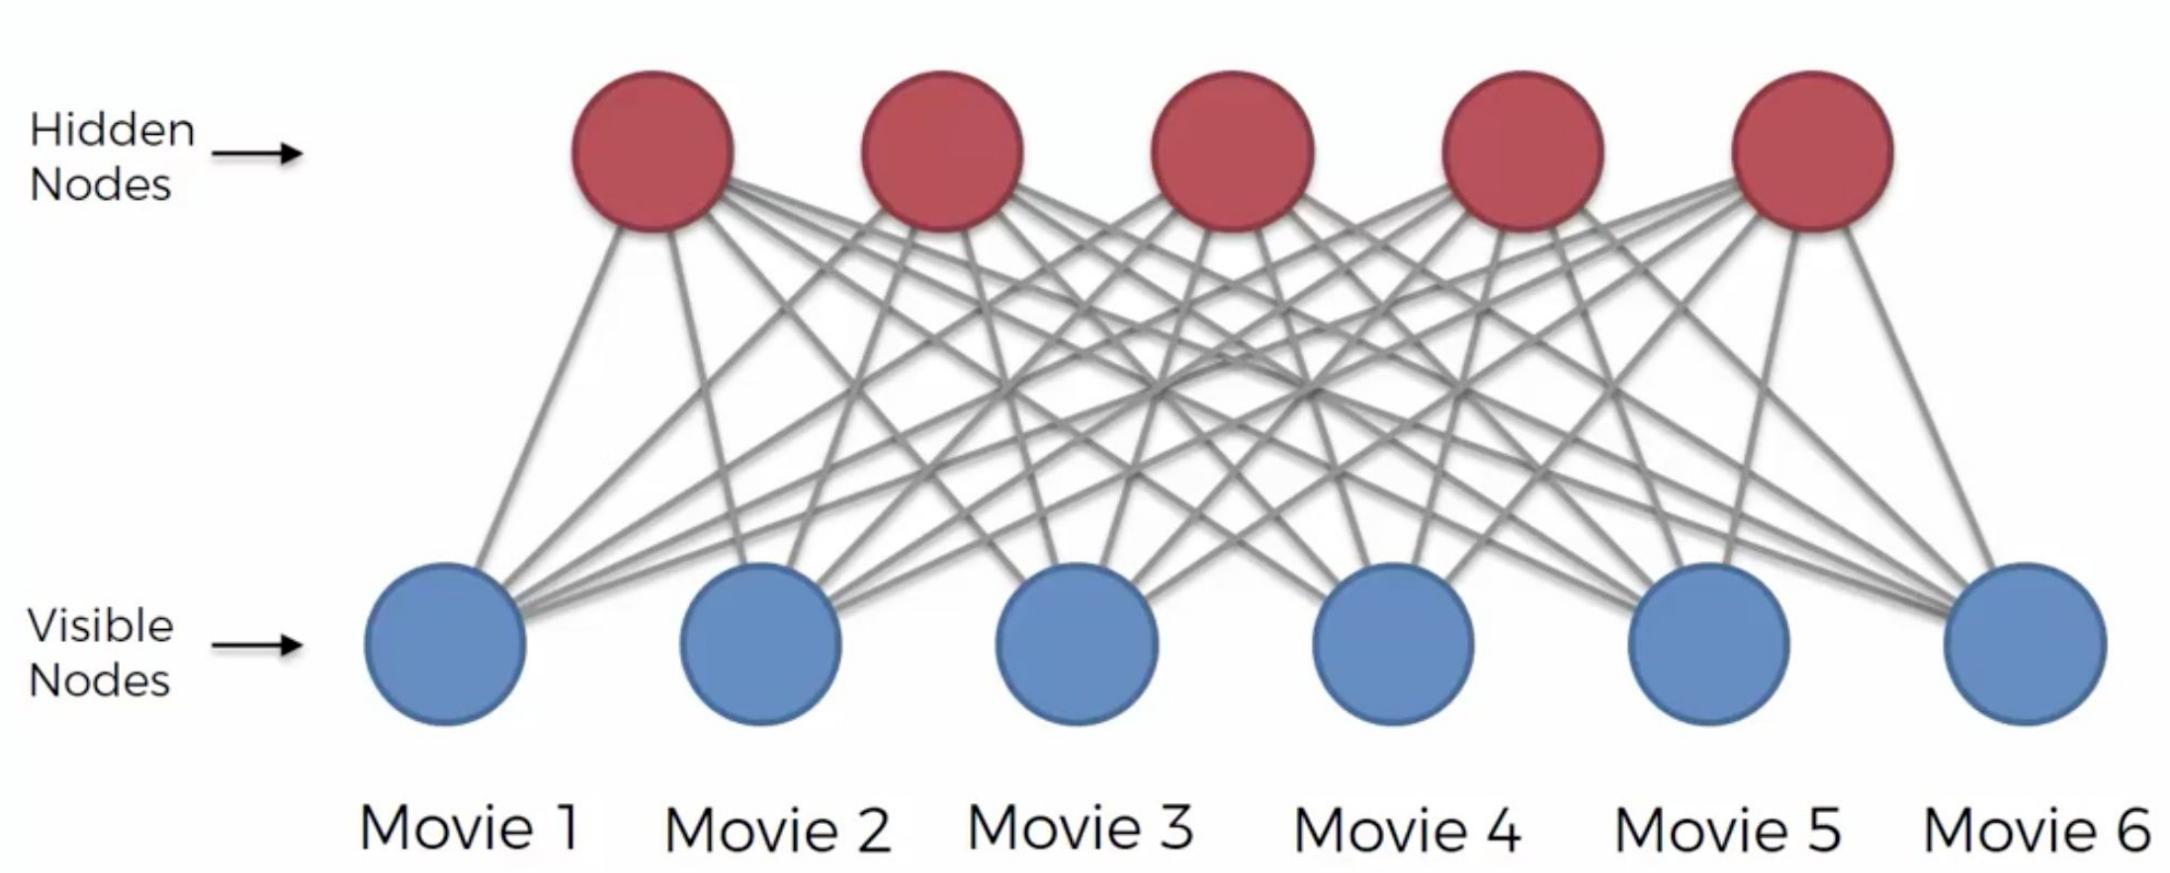
# Boltzmann Machines



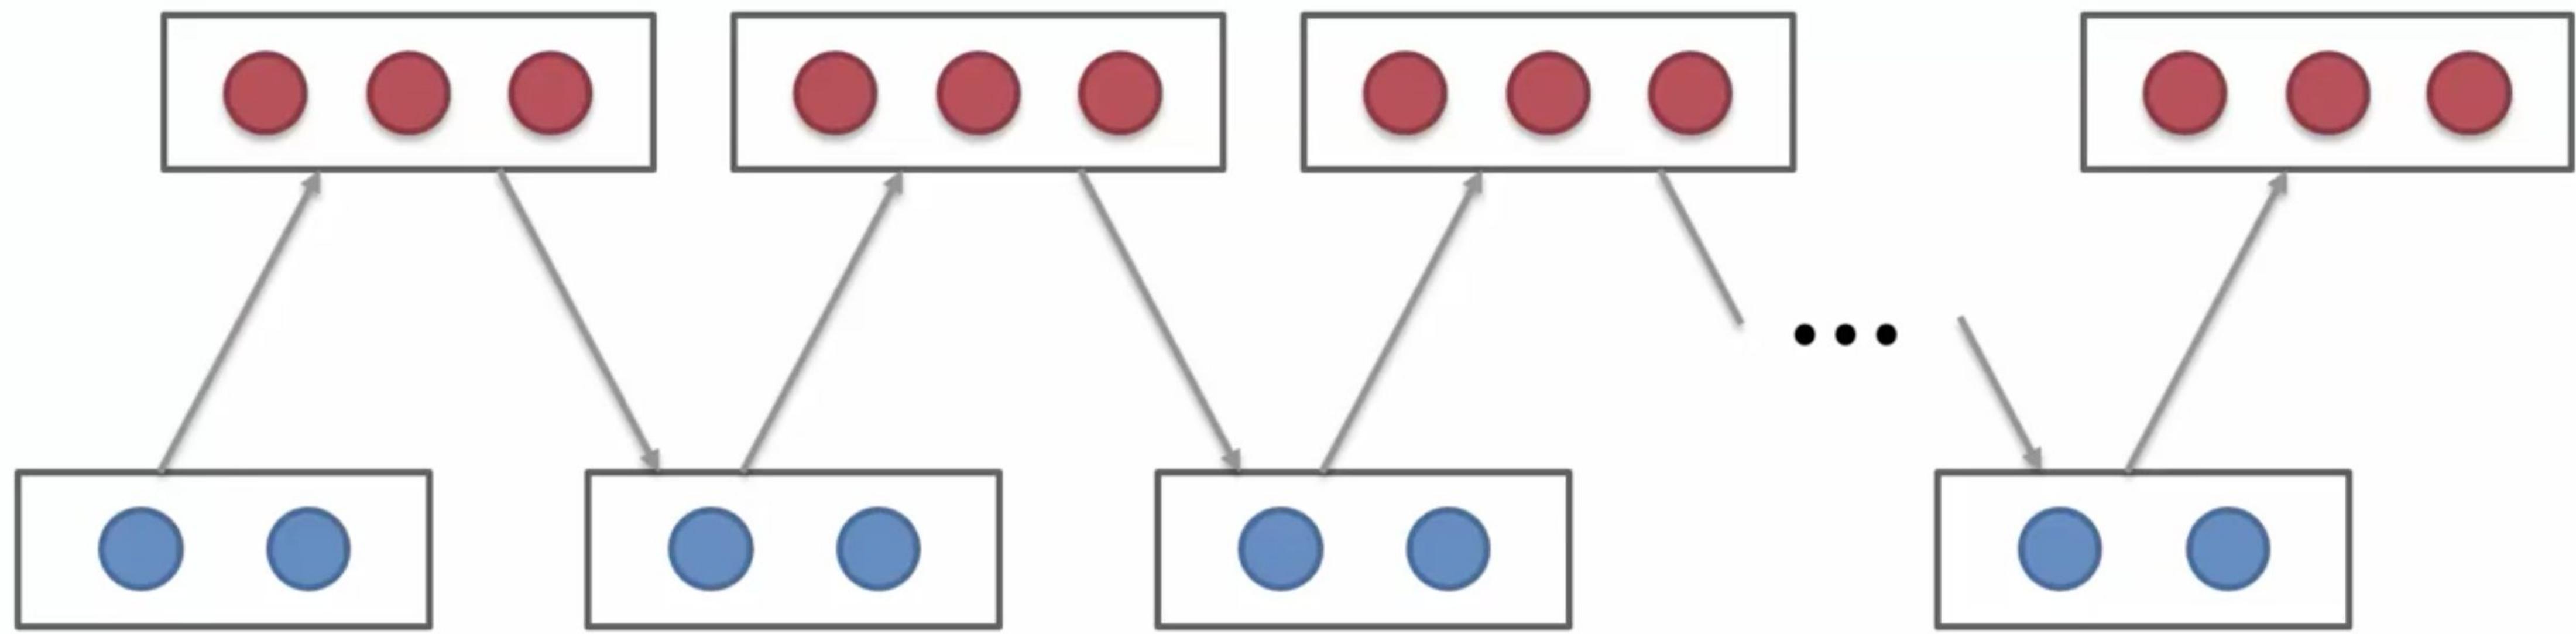
# Boltzmann Machines



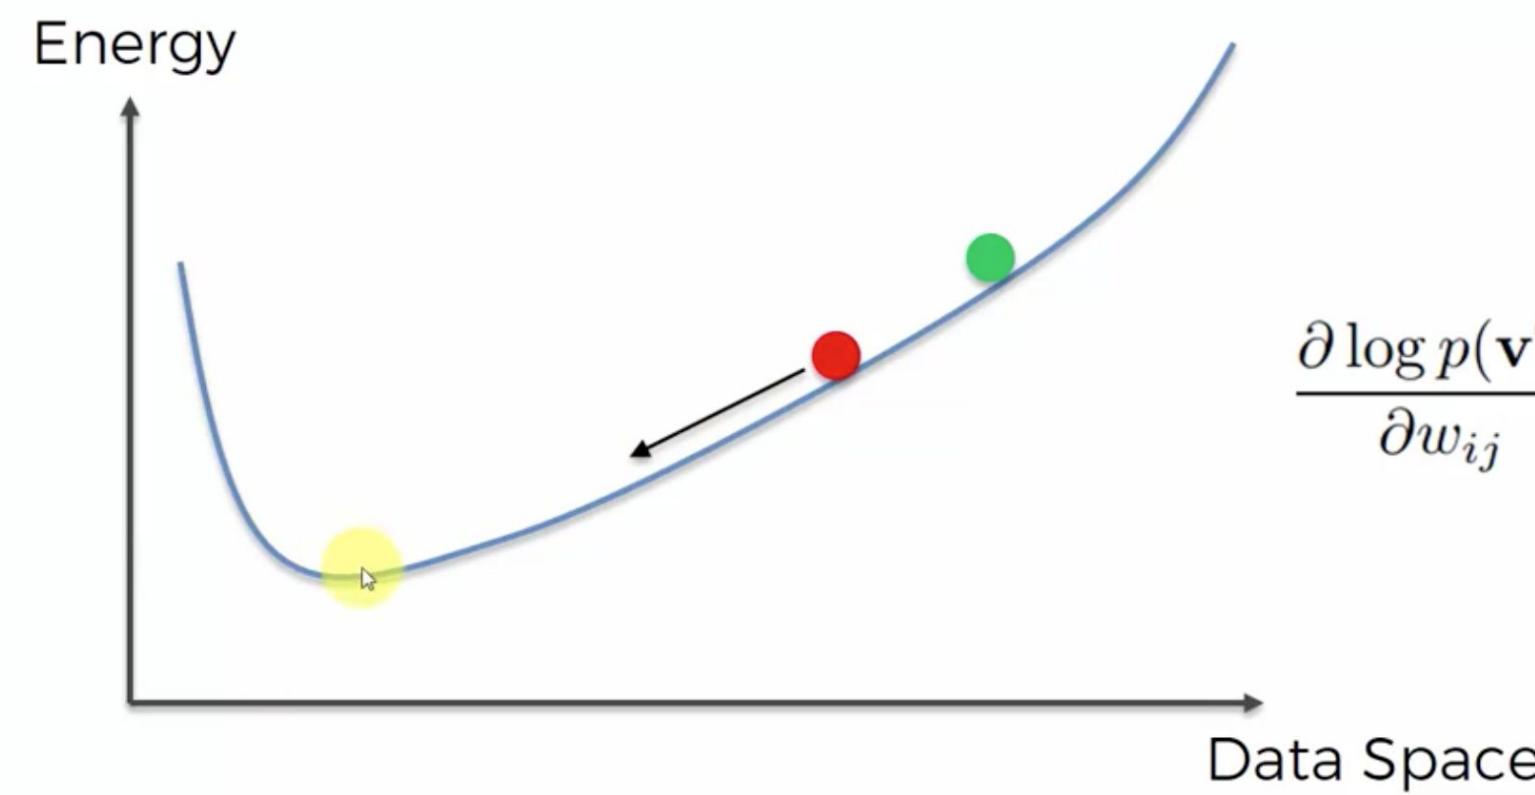
# Boltzmann Machines



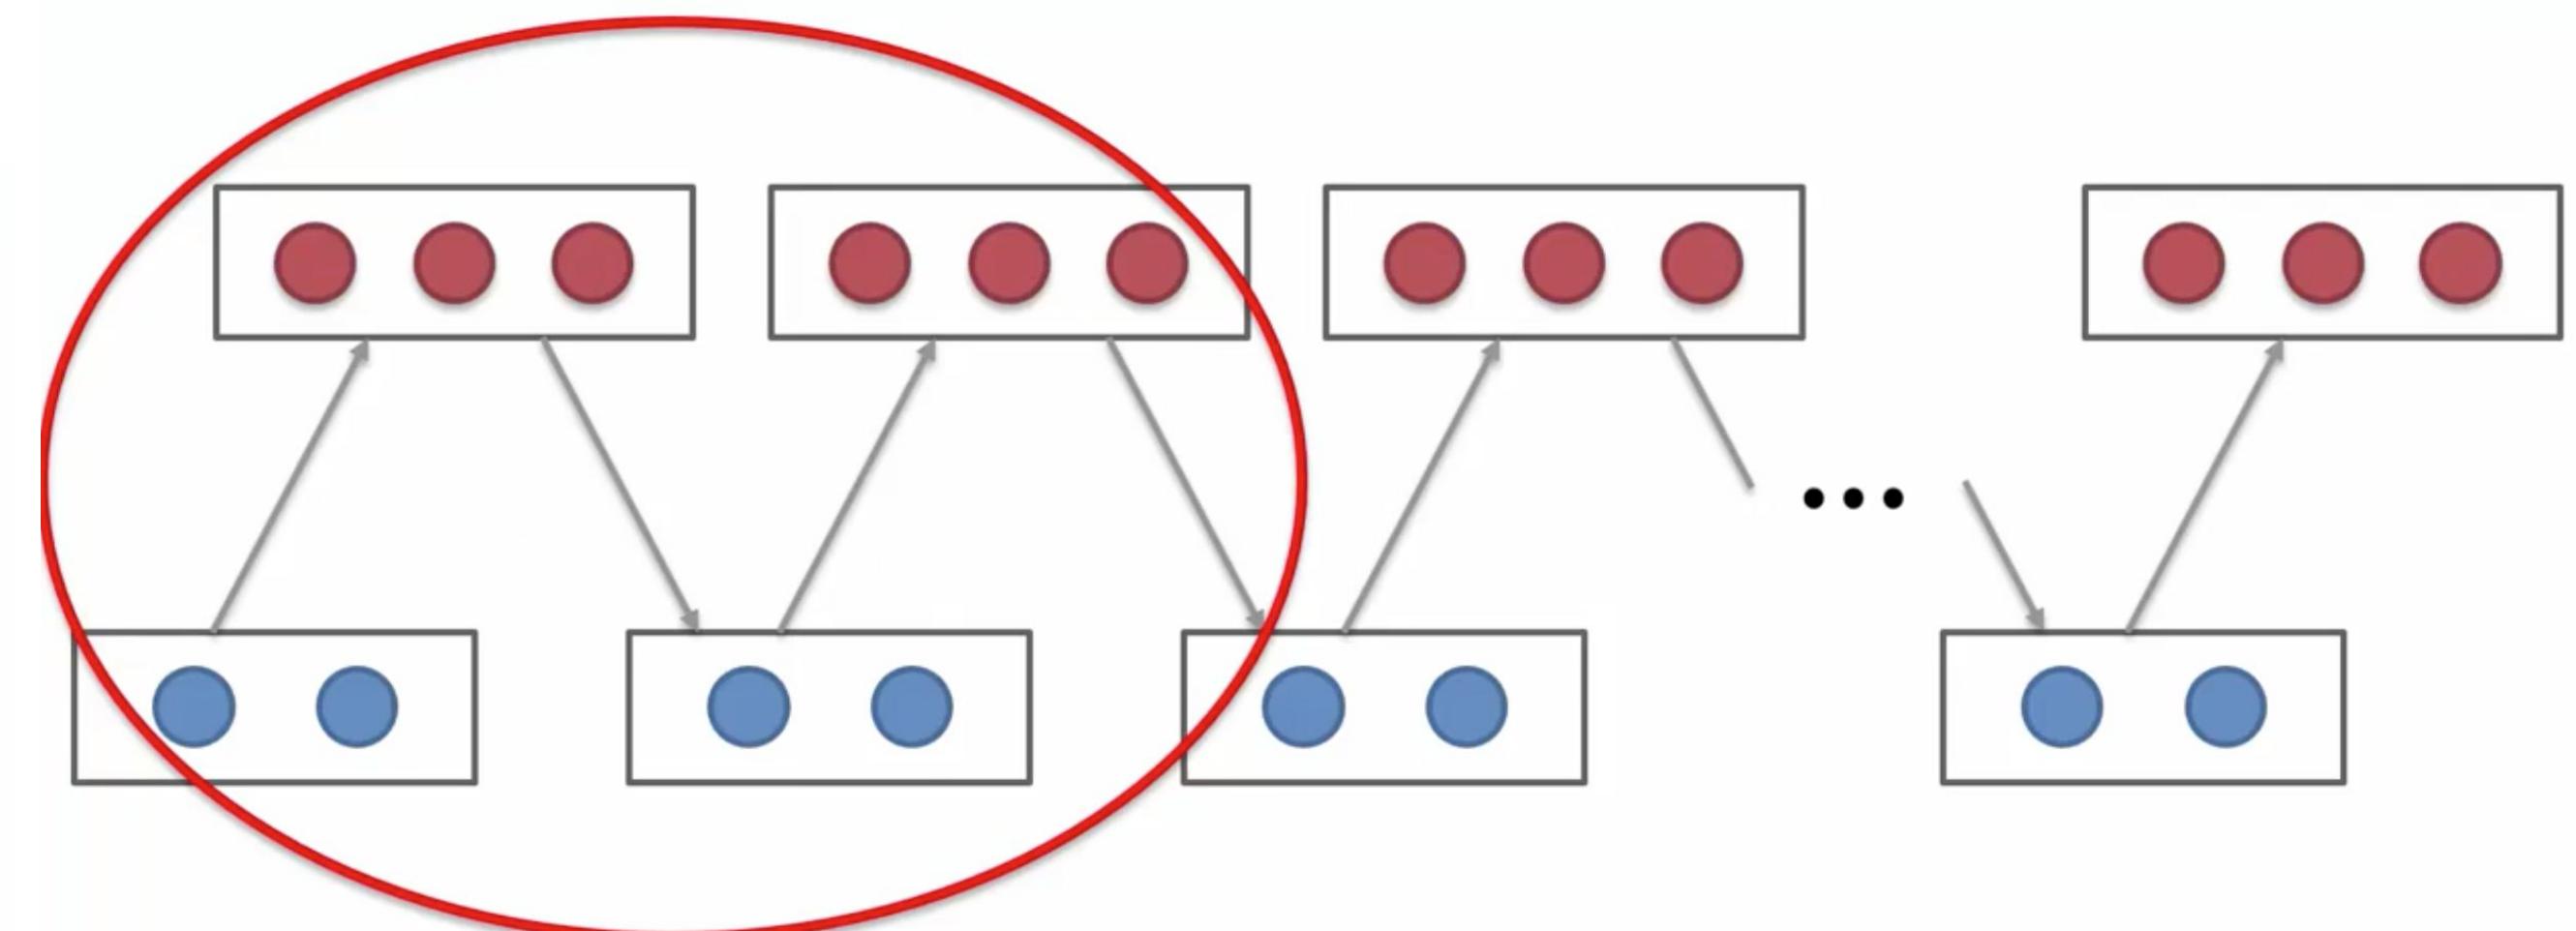
# Contrastive Divergence



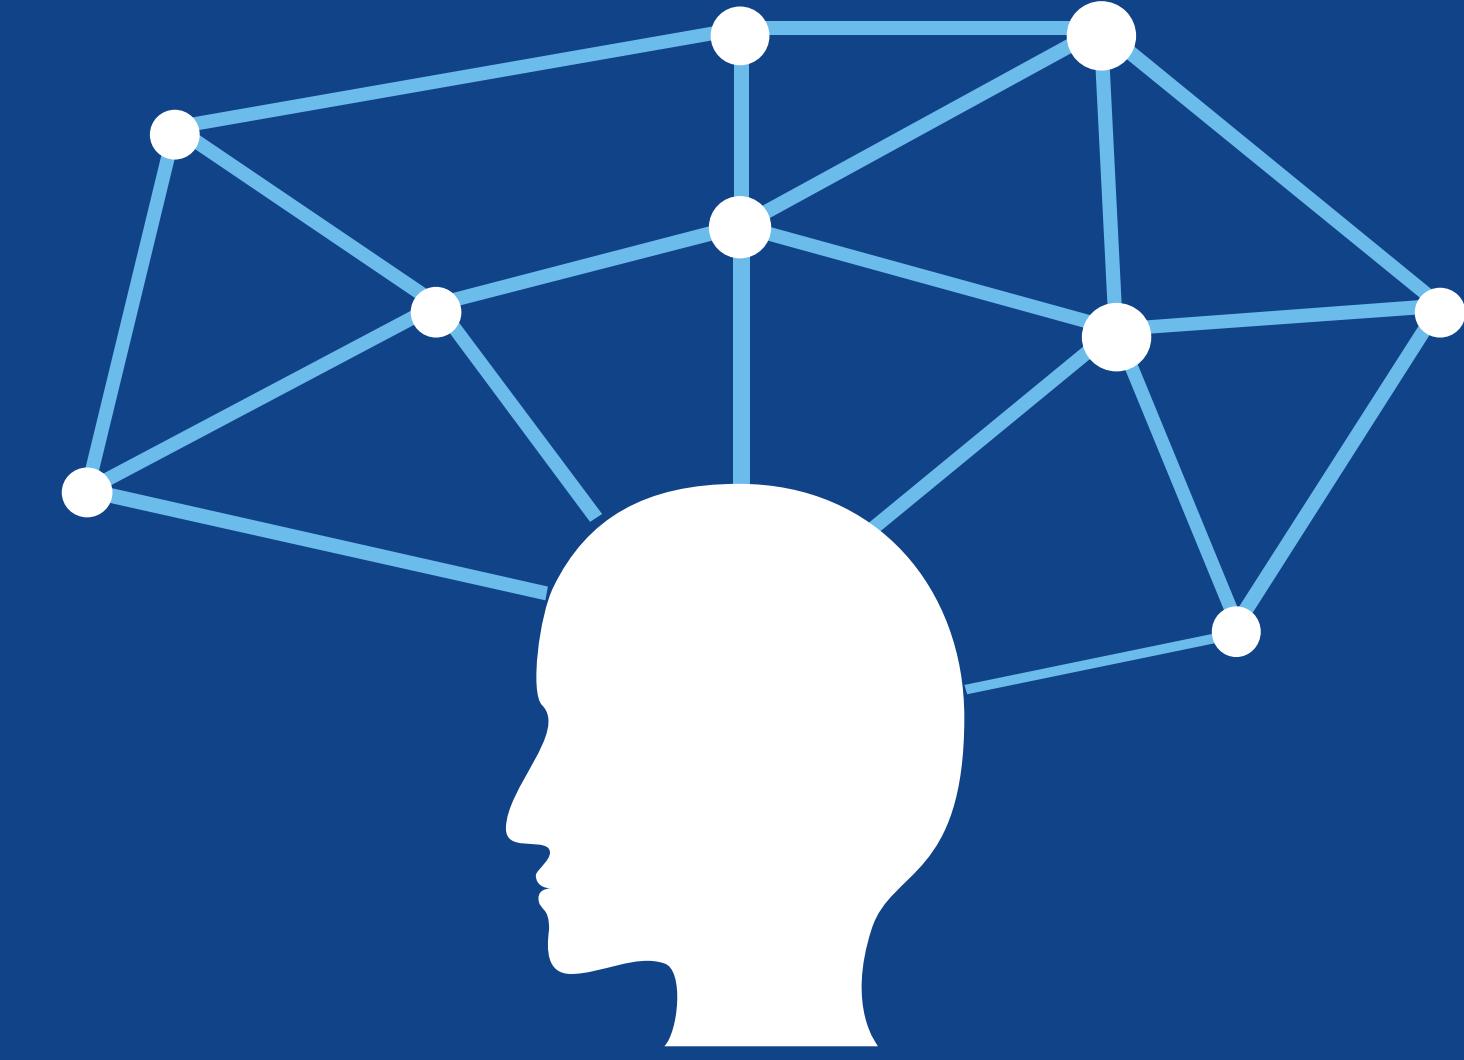
# Contrastive Divergence



$$\frac{\partial \log p(\mathbf{v}^0)}{\partial w_{ij}} = \langle v_i^0 h_j^0 \rangle - \langle v_i^\infty h_j^\infty \rangle$$



# Deep Learning Frameworks





# Frameworks

1

Open Source Deep Learning Libraries

# LIBRARIES

TensorFlow	Google Brain, 2015 (rewritten DistBelief)
Theano	University of Montréal, 2009
Keras	François Chollet, 2015 (now at Google)
Torch	Facebook AI Research, Twitter, Google DeepMind
Caffe	Berkeley Vision and Learning Center (BVLC), 2013

# More Libraries (1/3)

## Deeplearning4j

DeepLearning4j (or DL4J) is a popular deep learning framework developed in Java and supports other JVM languages as well. It is very slick, and is very widely used as a commercial, industry-focused distributed deep learning platform. The advantage of using DL4j is that you can bring together the power of the whole Java ecosystem to perform efficient deep learning, as it can be implemented on top of the popular Big Data tools such as Apache Hadoop and Apache Spark.

## MXNet

MXNet is one of the most languages-supported deep learning frameworks, with support for languages such as R, Python, C++ and Julia. This is helpful because if you know any of these languages, you won't need to step out of your comfort zone at all, to train your deep learning models. Its backend is written in C++ and cuda, and is able to manage its own memory like Theano. MXNet is also popular because it scales very well and is able to work with multiple GPUs and computers, which makes it very useful for the enterprises. This is also one of the reasons why Amazon made MXNet its reference library for Deep Learning too.

# More Libraries (2/3)

## Microsoft Cognitive Toolkit

Microsoft Cognitive Toolkit, previously known by its acronym CNTK, is an open-source deep learning toolkit to train deep learning models. It is highly optimized, and has support for languages such as Python and C++. Known for its efficient resource utilization, you can easily implement efficient Reinforcement Learning models or Generative Adversarial Networks (GANs) using the Cognitive Toolkit. It is designed to achieve high scalability and performance, and is known to provide high performance gains when compared to other toolkits like Theano and Tensorflow, when running on multiple machines.

## Lasagne

Lasagne is a high-level deep learning library that runs on top of Theano. It has been around for quite some time now, and was developed with the aim of abstracting the complexities of Theano, and provide a more friendly interface to the users to build and train neural networks. It requires Python, and finds many similarities to Keras, which we just saw above. However, if we are to find differences between the two, Keras is faster, and has a better documentation in place.

# More Libraries (3/3)

## BigDL

BigDL is distributed deep learning library for Apache Spark, and is designed to scale very well. With the help of BigDL, you can run your deep learning applications directly on Spark or Hadoop clusters, by writing them as Spark programs. It has a rich deep learning support, and uses Intel's Math Kernel Library (MKL) to ensure high performance. Using BigDL, you can also load your pre-trained Torch or Caffe models into Spark. If you want to add deep learning functionalities to a massive set of data stored on your cluster, this is a very good library to use.

## Many More ...

There are many other deep learning libraries and frameworks available for use today – DSSTNE, Apache Singa, Veles are just a few worth an honourable mention.



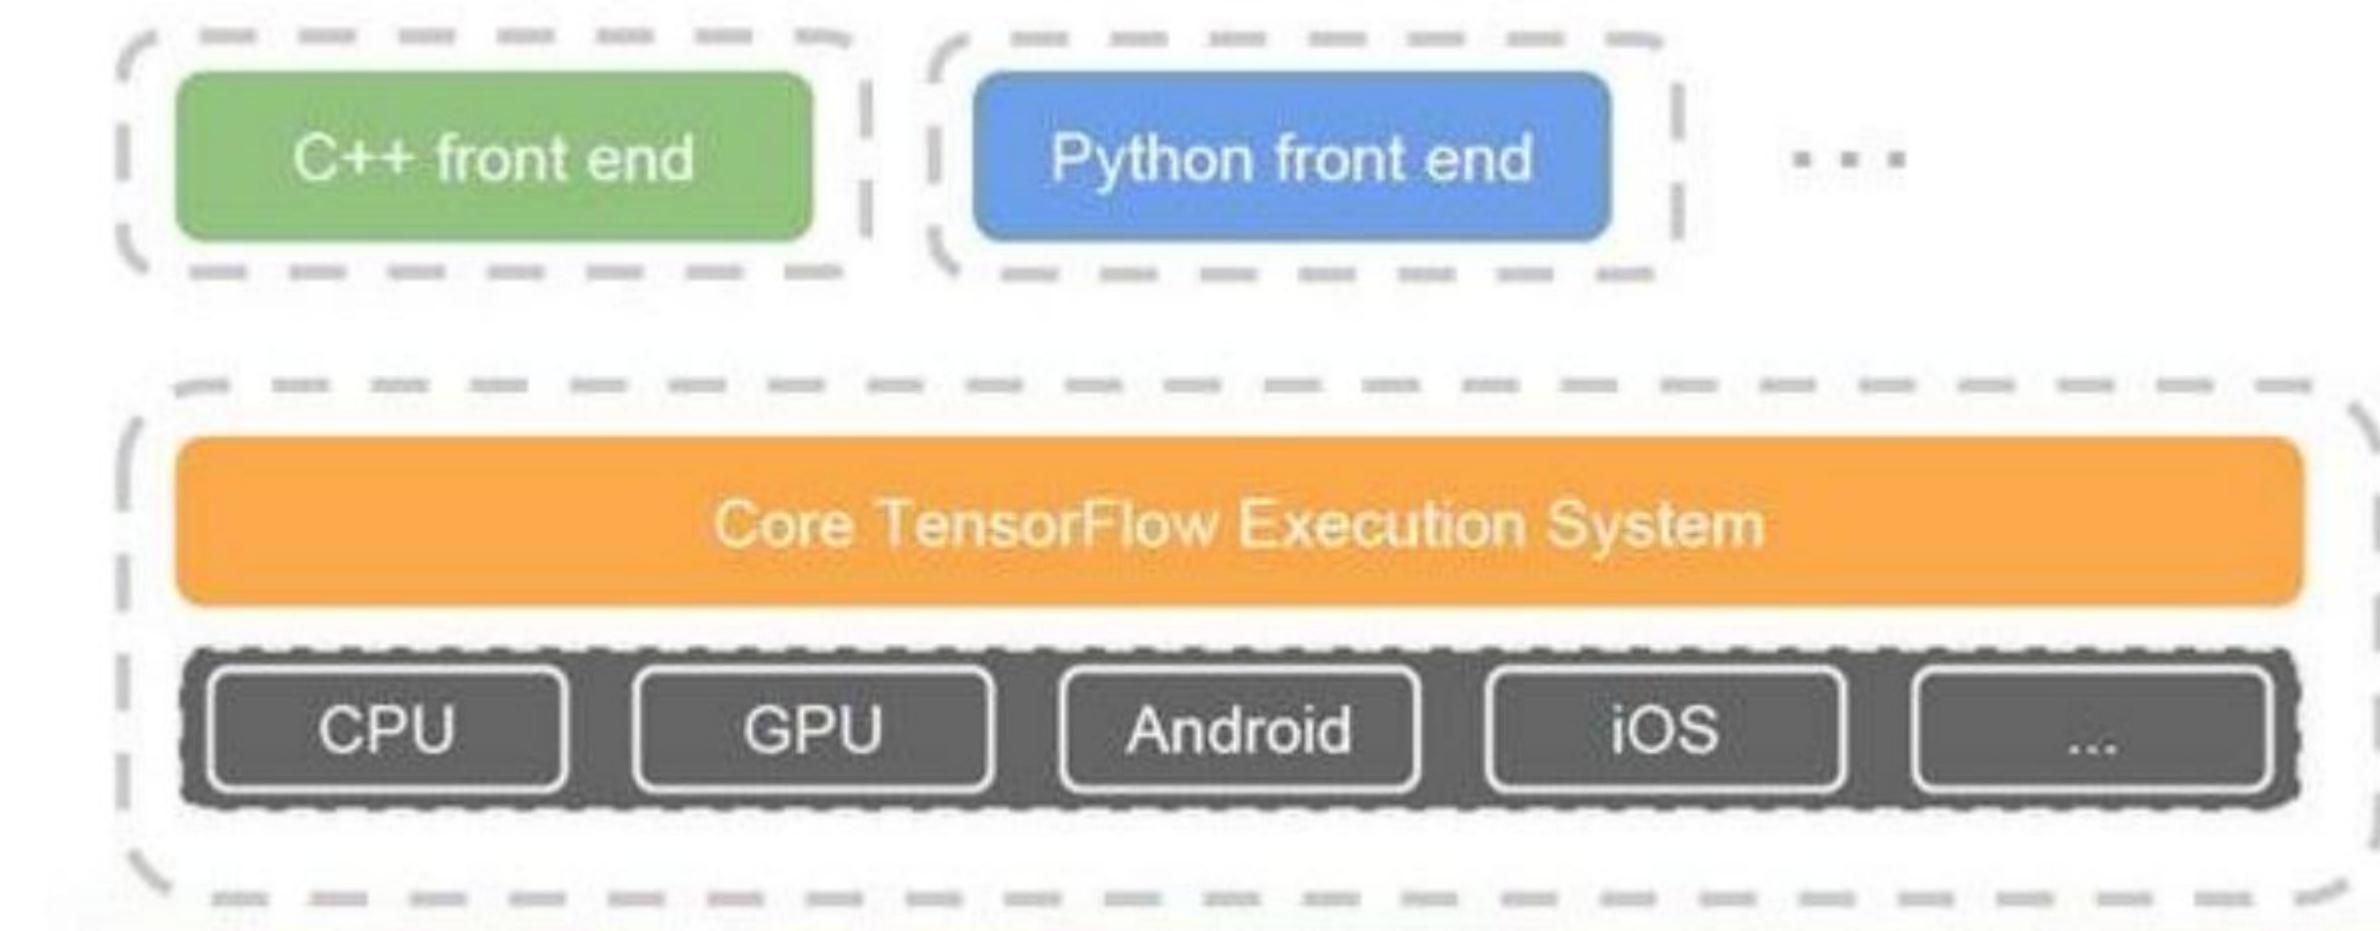
# TensorFlow, Theano, Keras, Torch, Caffe

# 2

Open Source Deep Learning Libraries

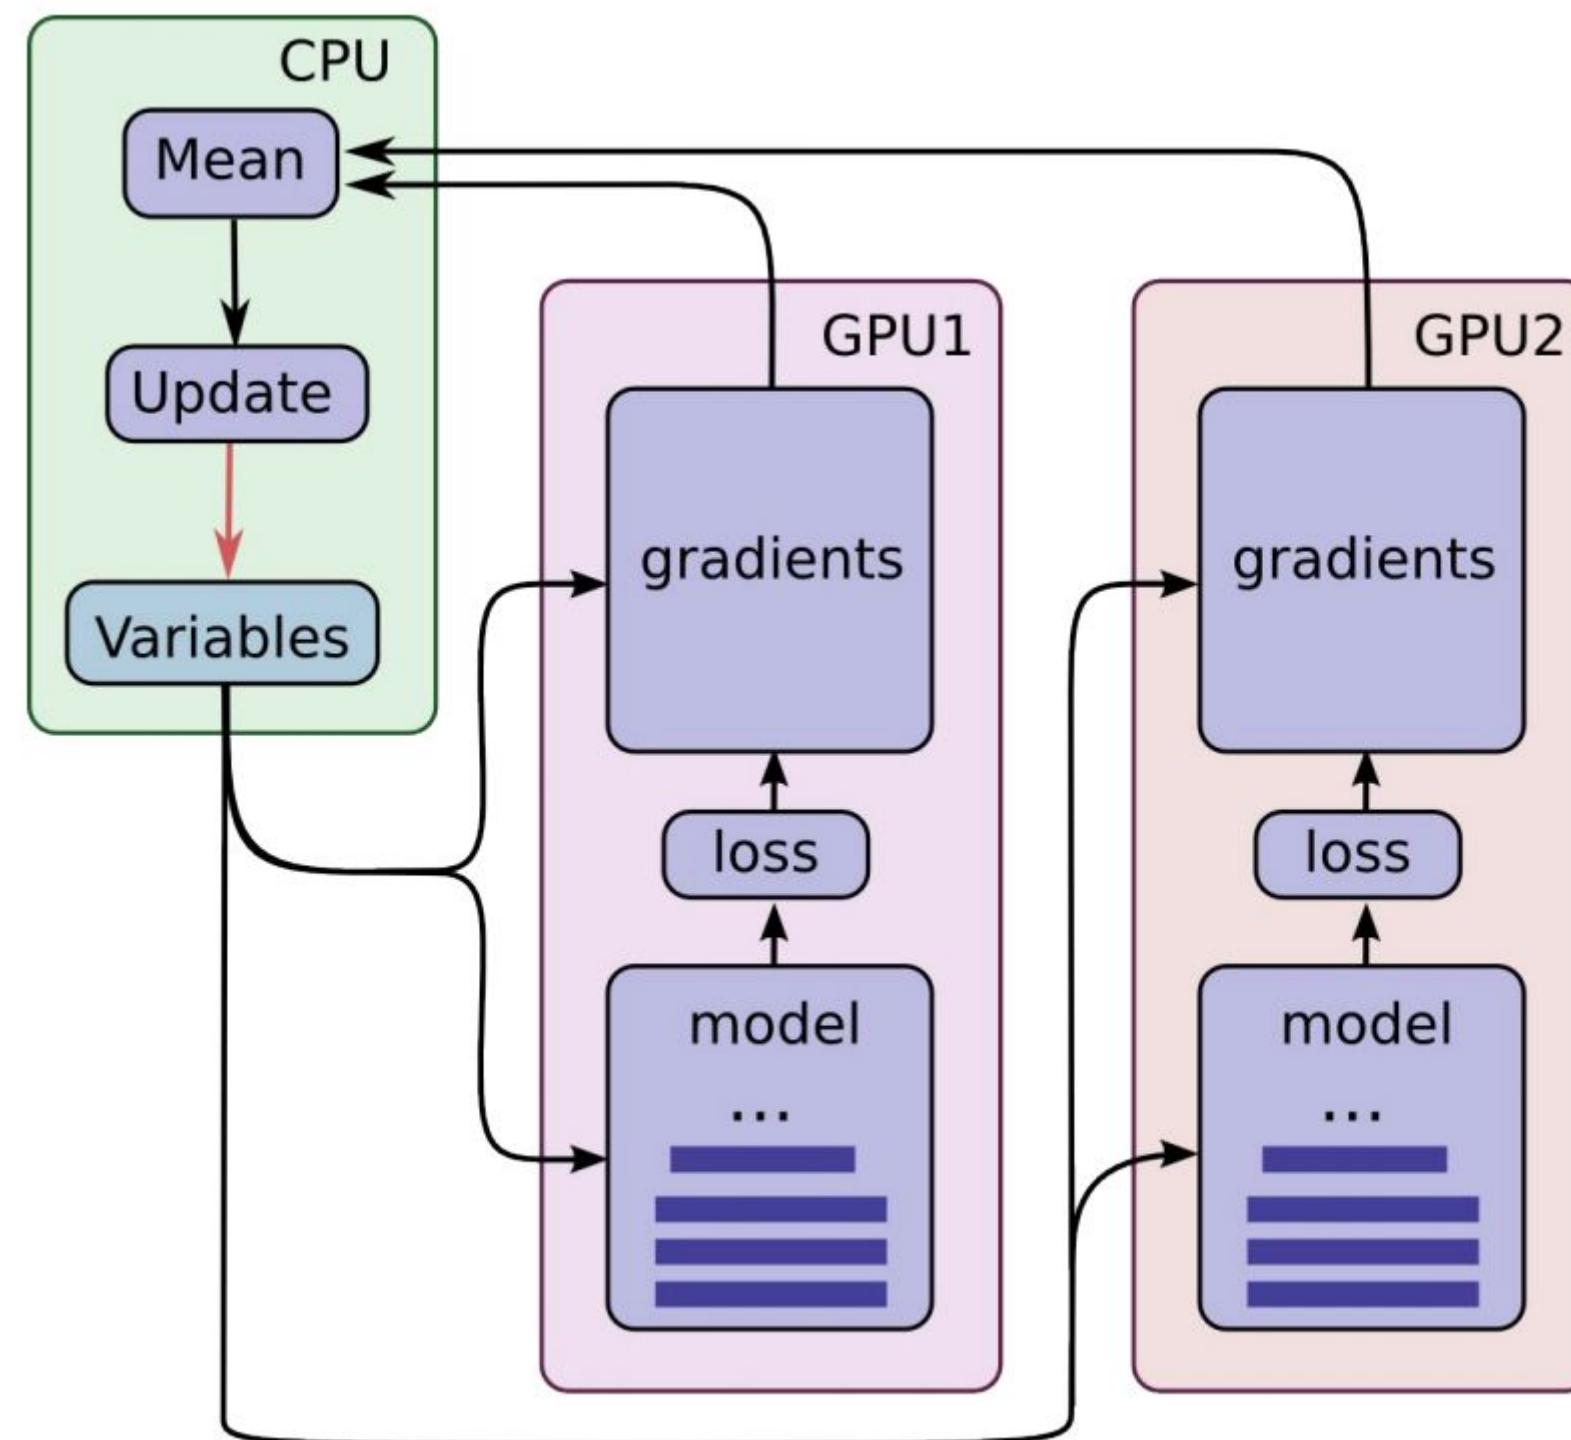
# TensorFlow Architecture

- 1) Low-level core (C++/CUDA)
- 2) Simple Python API to define the computational graph
- 3) High-level API (TF-Learn, TF-Slim, Keras...)

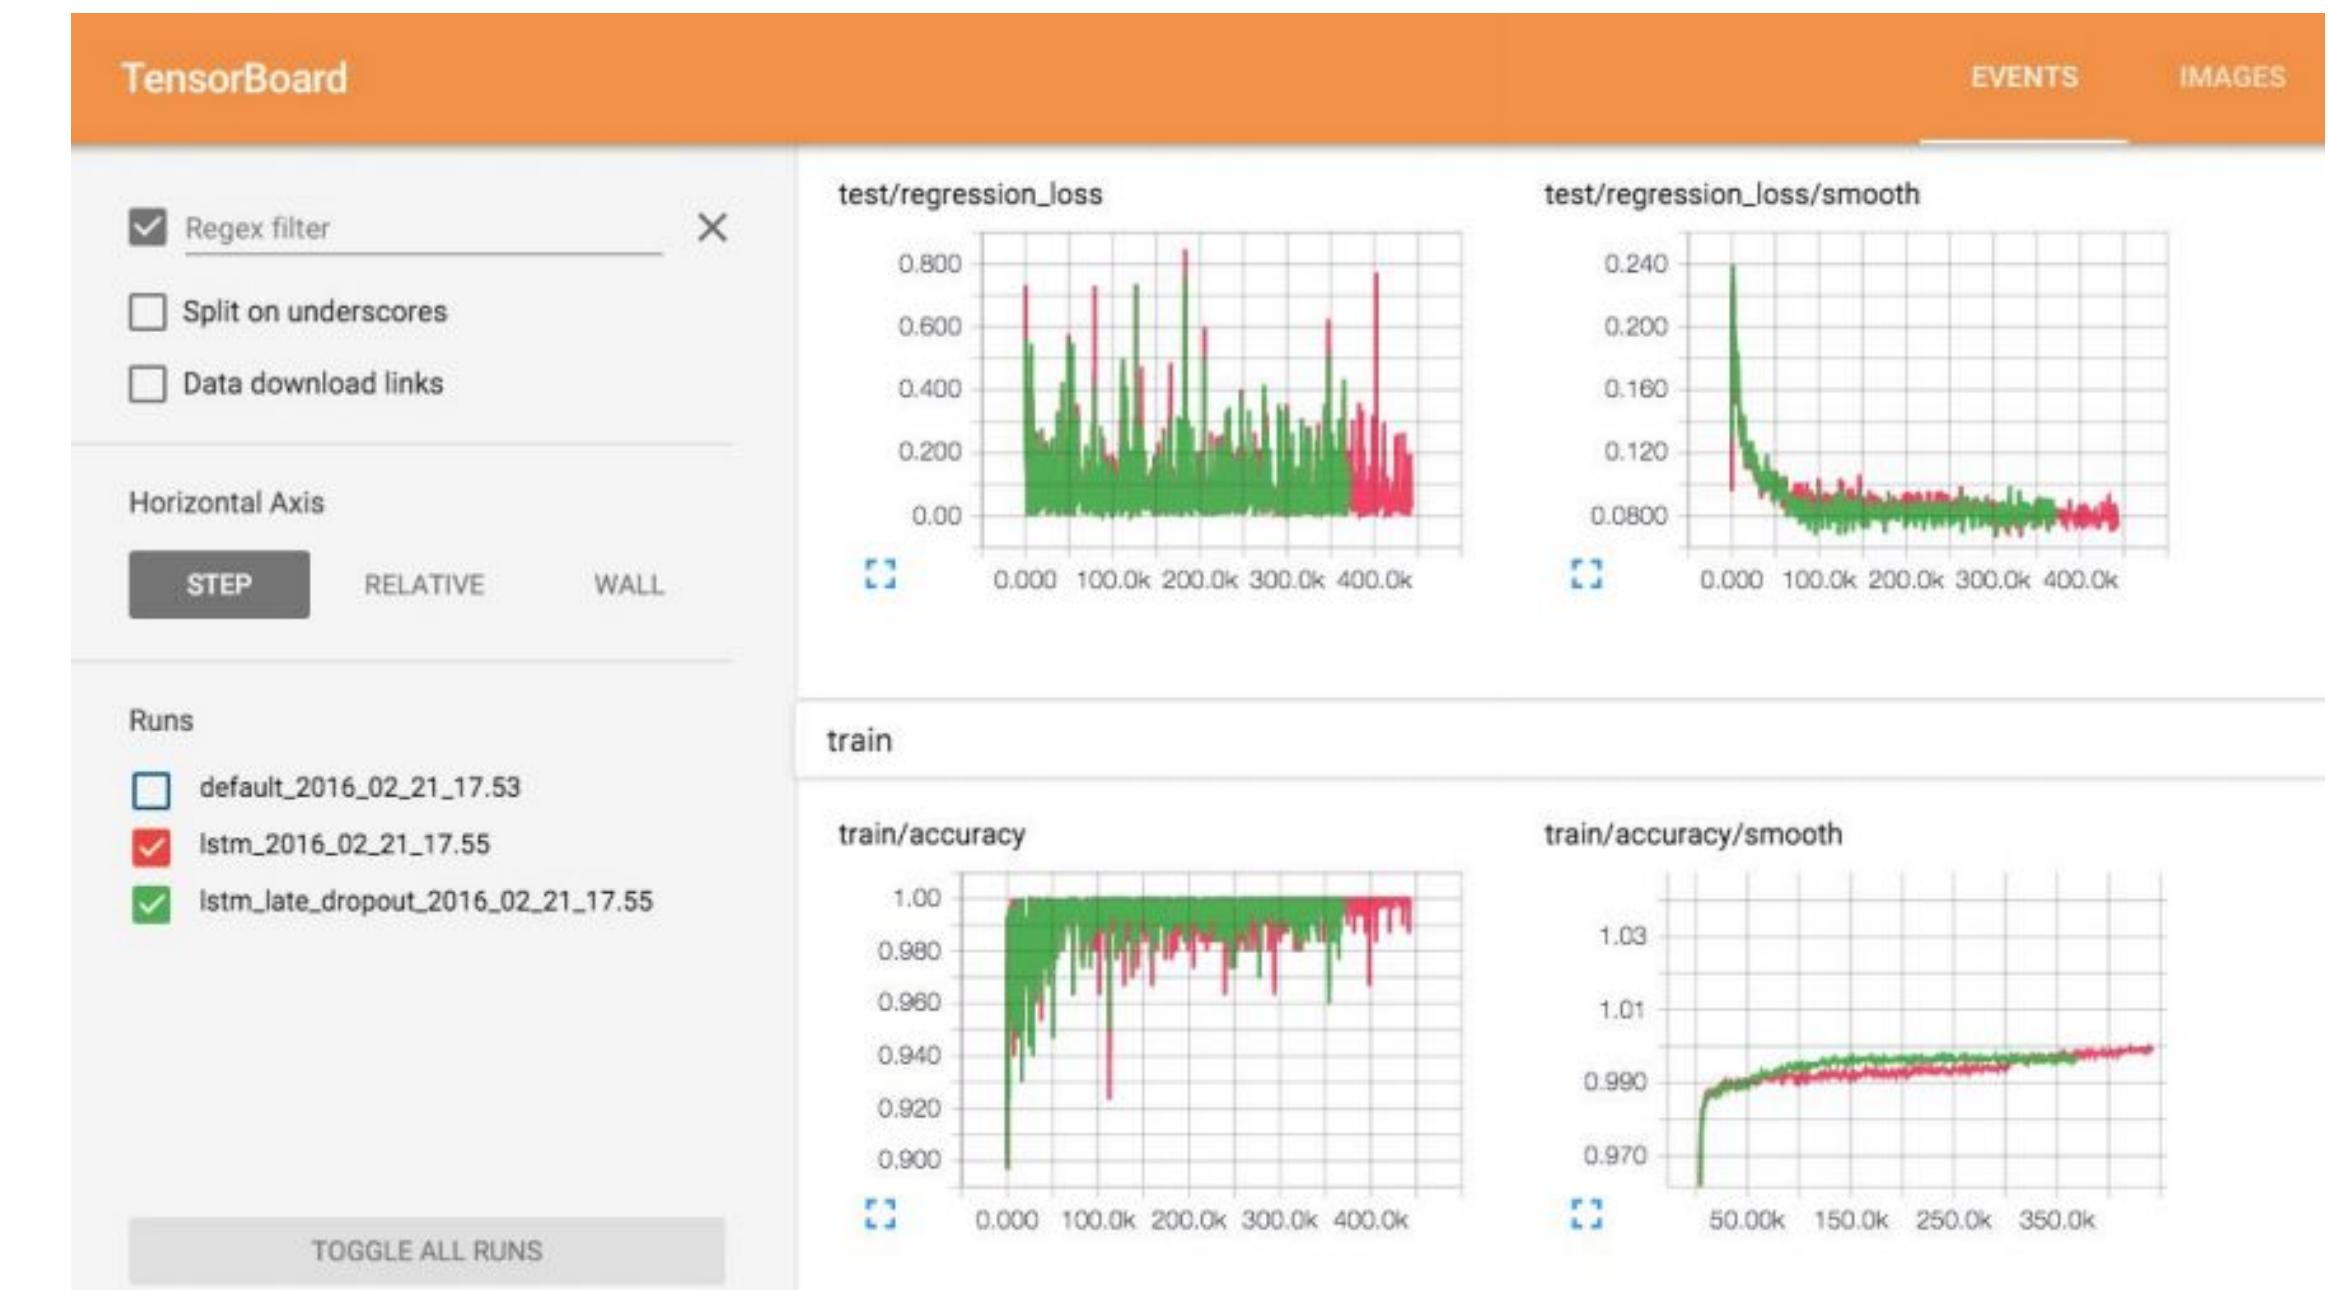
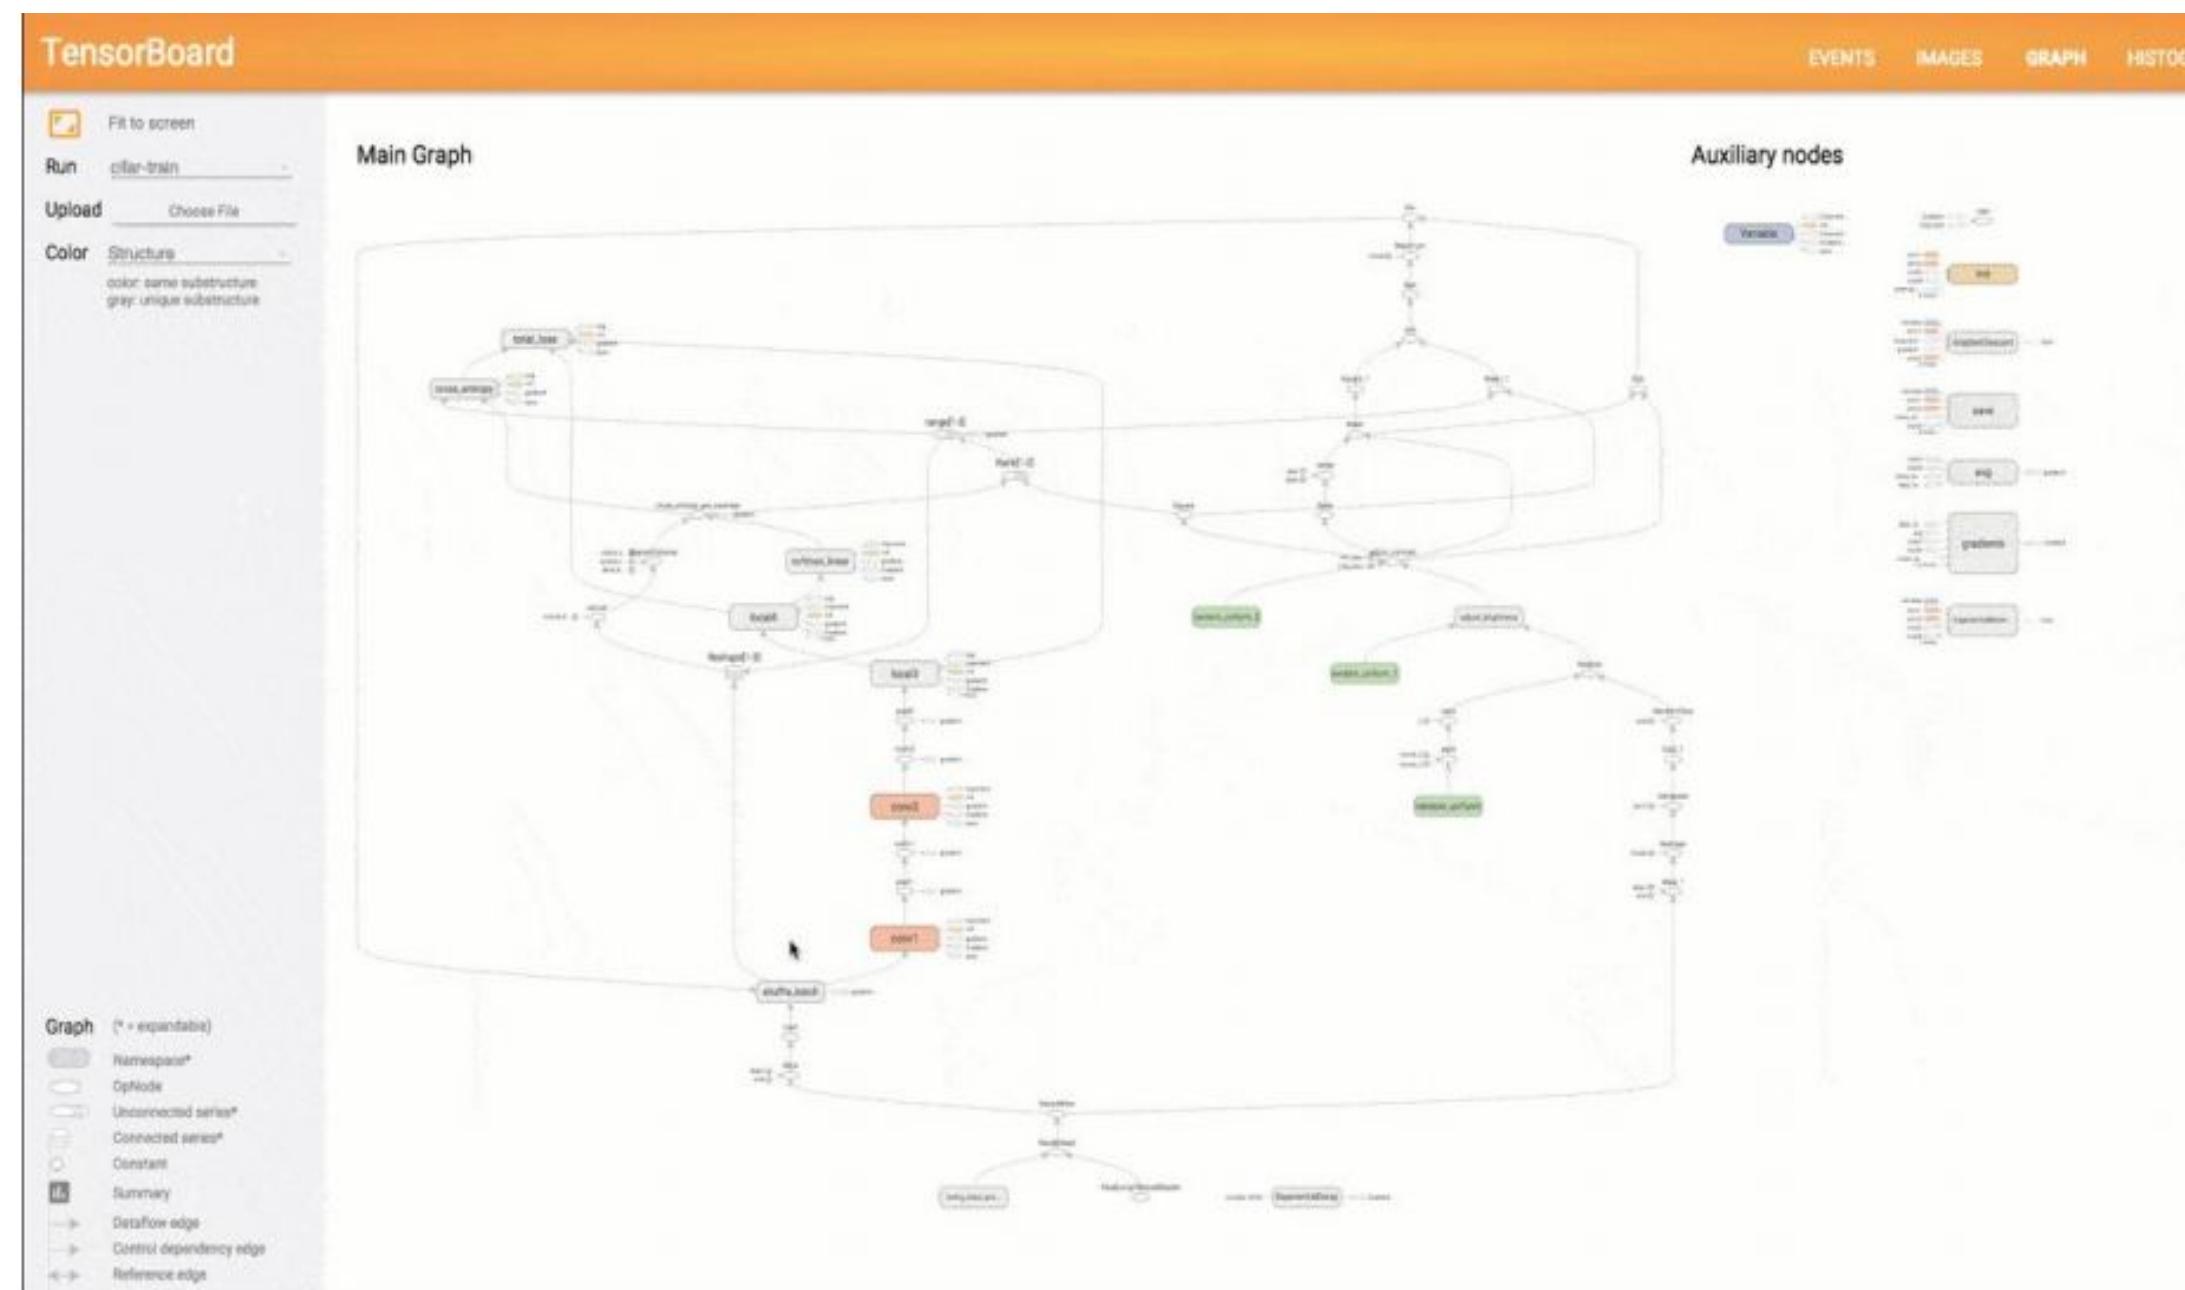


# TensorFlow computational graph

- auto-differentiation!
- easy multi-GPU/multi-node
- native C++ multithreading
- device-efficient
- implementation for most ops
- whole pipeline in the graph:
- data loading, preprocessing,
- prefetching...



# TensorBoard



# TensorFlow Development

- + bleeding edge (GitHub yay!)
- + division in core and contrib => very quick merging of new hotness
- + a lot of new related API: CRF, BayesFlow, SparseTensor, audio IO, CTC, seq2seq
- + so it can easily handle images, videos, audio, text...
- + if you really need a new native op, you can load a dynamic lib
- sometimes contrib stuff disappears or moves
- recently introduced bells and whistles are barely documented



# Code + Models, Performance, Model deployment

3

Open Source Deep Learning Libraries

# Code and Models

TensorFlow	Theano	Keras	Torch	Caffe
More code and models available (hype), including a lot provided by Google, some of open sourced code is already unusable due to very fast development, caffe models loaded in semi-manual way	<ul style="list-style-type: none"><li>- Lots of code to get started (tutorials and older models).</li><li>- Loss of momentum and so recent models appear more slowly.</li><li>- State of the art models can always be found, pretrained models can be slow to appear.</li><li>- Community based: new code can be unstable.</li></ul>	Models from TensorFlow and Theano + converted model from caffe	State of the art classification models available in the <a href="#">ModelZoo</a> . Significant proportion of research projects in torch. <a href="#">Loadcaffe</a> convert caffe models to torch (often involves 0.5 to 2 days of work when non-sequential architecture or “non-standard layers”)	<ul style="list-style-type: none"><li>- <b>Extensive code</b> available online (even from the 1st year: &gt;1K forked + significant changes since then)</li><li>- <b>heaven</b>: state of the art pre-trained models available for a variety of domains</li><li>- great for <b>fine-tuning</b> networks even without writing code</li></ul>

# Community and Documentation

## Community : (Github, groups, discussions...)

- Caffe had the largest community
- TensorFlow's is already large and growing.
- Keras' community is growing, while Theano's and Lasagne's are declining

## Documentation

- Great documentation for Theano, Lasagne, Keras and Torch
- Most recent API is not documented for TensorFlow. Tutorials are often outdated.



## Community

- Caffe has a large community
- TensorFlow has a large community
- Keras' community is declining

Computation using data flow graphs for scalable machine learning <http://tensorflow.org>

tensorflow machine-learning python deep-learning deep-neural-networks neural-network ml distributed

25,567 commits 16 branches 43 releases 1,182 contributors Apache-2.0

Branch: master New pull request Create new file Upload files Find file Clone or download

Theano / Theano

Watch 566 Star 7,389 Fork 2,349

Code Issues 542 Pull requests 101 Projects 0 Wiki Insights

## Documentation

- Great documentation
- Most recent documentation is often outdated

Theano is a Python library that allows you to define, optimize, and evaluate mathematical expressions involving multi-dimensional arrays efficiently. It can use GPUs and perform efficient symbolic differentiation.

[http://www.deeplearning.net/software/...](http://www.deeplearning.net/software/)

27,893 commits 3 branches 30 releases 326 contributors

Watch 221 Star 3,324 Fork 912

Code Issues 107 Pull requests 28 Projects 0 Wiki Insights

Lightweight library to build and train neural networks in Theano <http://lasagne.readthedocs.org/>

deep-learning-library neural-networks python theano

1,136 commits 2 branches 1 release 62 contributors

# Performance

TensorFlow	Theano	Keras	Torch	Caffe
Optimized for big models, can be memory hungry, but as fast as others (cuDNN).	<ul style="list-style-type: none"><li>- Run-time and memory competitive (efficient RNNs),</li><li>- Compile time can be a huge pain (Strong -).</li><li>- Multi-GPU support, not multi-machines.</li></ul>	Cf Theano and TensorFlow	All rely on cuDNN. Advantage : no compilation of models, which saves a lot of time during debugging. Memory : some layers are not very efficient because of inner buffers (fixed with OptNet or Pytorch.)	<ul style="list-style-type: none"><li>- Quite fast, eg. NVIDIA TitanX GPU:<ul style="list-style-type: none"><li>• Training: ~20 secs/20 iterations (5K images)</li><li>• Testing: ~70 secs /validation set (50K images)</li></ul></li><li>- quick compilation</li><li>- multi-GPU support but not with python layers</li><li>- no distributed training</li></ul>

# Model Deployment

TensorFlow	Theano	Keras	Torch	Caffe
TF Serving for self-hosted web, Google Cloud Platform for easy web	Although compiled in C++/CUDA, can't be deployed without Python	Cf Theano and TensorFlow	Require LuaJIT to run models. (Can be problematic for integration more than performance)	+ C++ based + stable library - many forks + can be compiled in variety of devices

# Extra Features

TensorFlow	Theano	Keras	Torch	Caffe
+Written in C++ +Bindings for Python, C++, Java...	+Python	Python	+ Written in Lua and C/CUDA	Written in C++ Python and matlab interface
+Multi-GPU +Distributed +Windows, Android support	+Multi-GPU	+Multi-GPU +Android support	+ multi-GPU + distributed learning (torch-distlearn) + Stability : testing libraries	+ allows <b>cross-platform</b> (including windows) + <b>very stable</b>

# Which Framework to Choose When?

1. You are a student on DL itself?
2. You want to use DL only to get features?
3. You work in industry?
4. You started your internship or POC?
5. You want to solve assignments?
6. You are curious about deep learning?
7. You don't even know python?



# Which Framework to Choose When?

1. You are a student on DL itself? TensorFlow, Theano, Torch
2. You want to use DL only to get features? Keras, Caffe
3. You work in industry? TensorFlow, Caffe
4. You started your internship or POC? Keras, Caffe
5. You want to solve assignments? Keras, Caffe
6. You are curious about deep learning? Caffe
7. You don't even know python? Keras, Torch



# References

<https://www.nature.com/articles/nature14539> -- Must read with best papers on Deep Learning (Yann LeCun, Yoshua Bengio & Geoffrey Hinton)

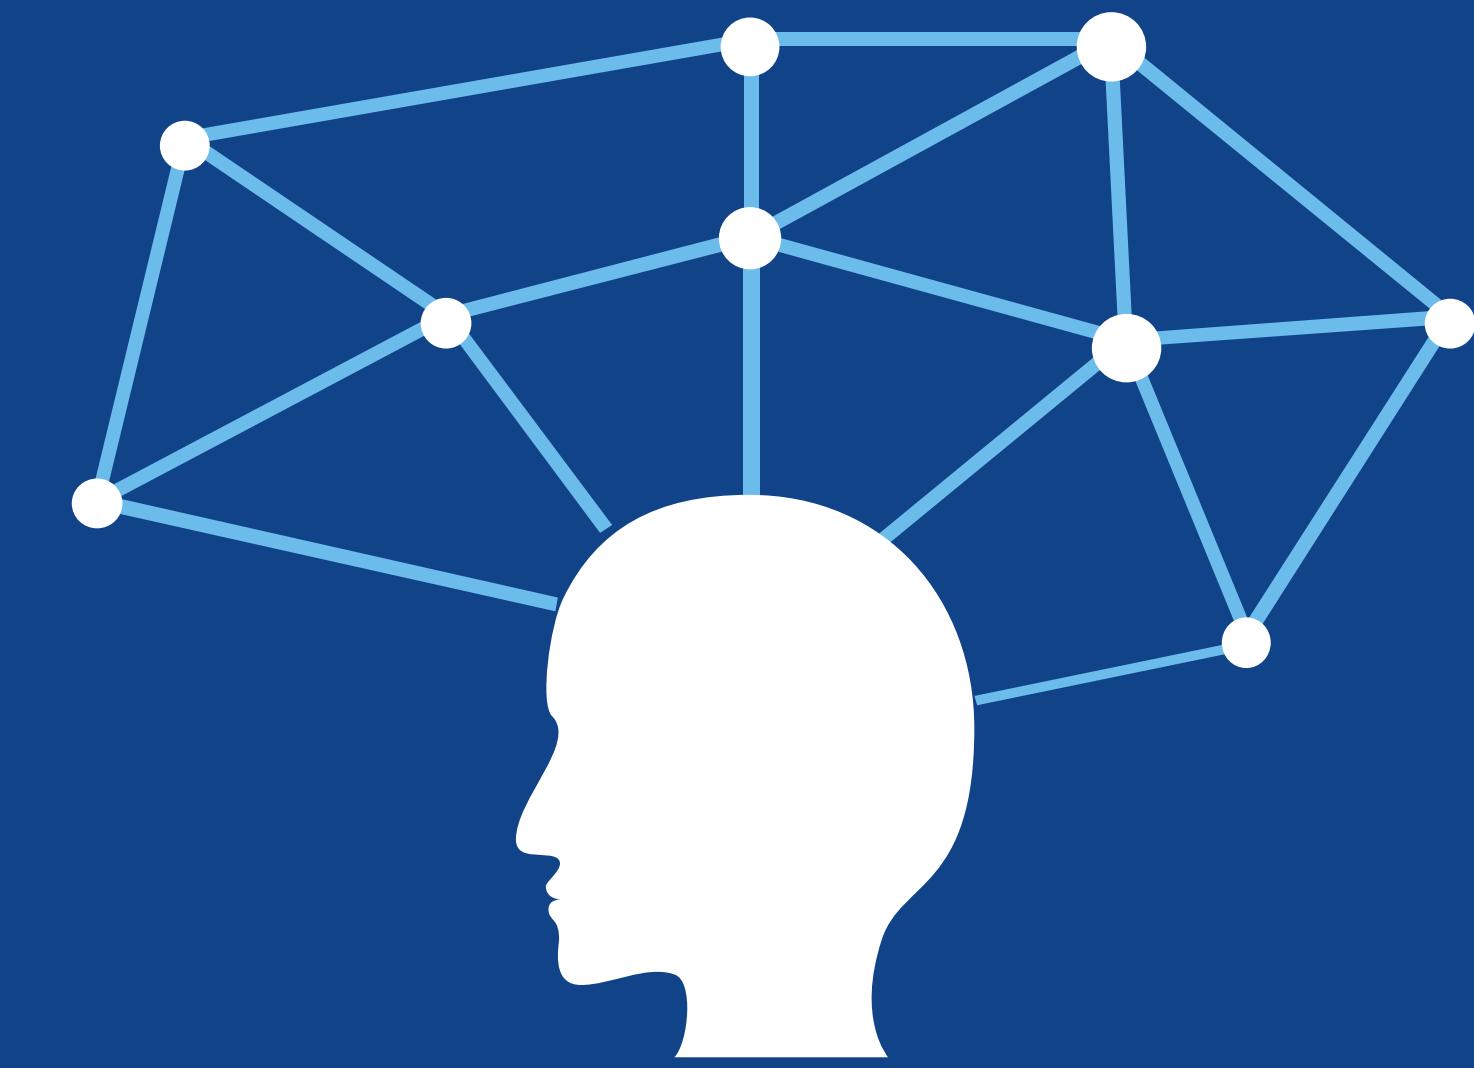
<http://karpathy.github.io/> -- Very Good Blogs on Deep Learning Architectures

<http://colah.github.io/> -- Theoretical understanding of various algorithms (RNN)

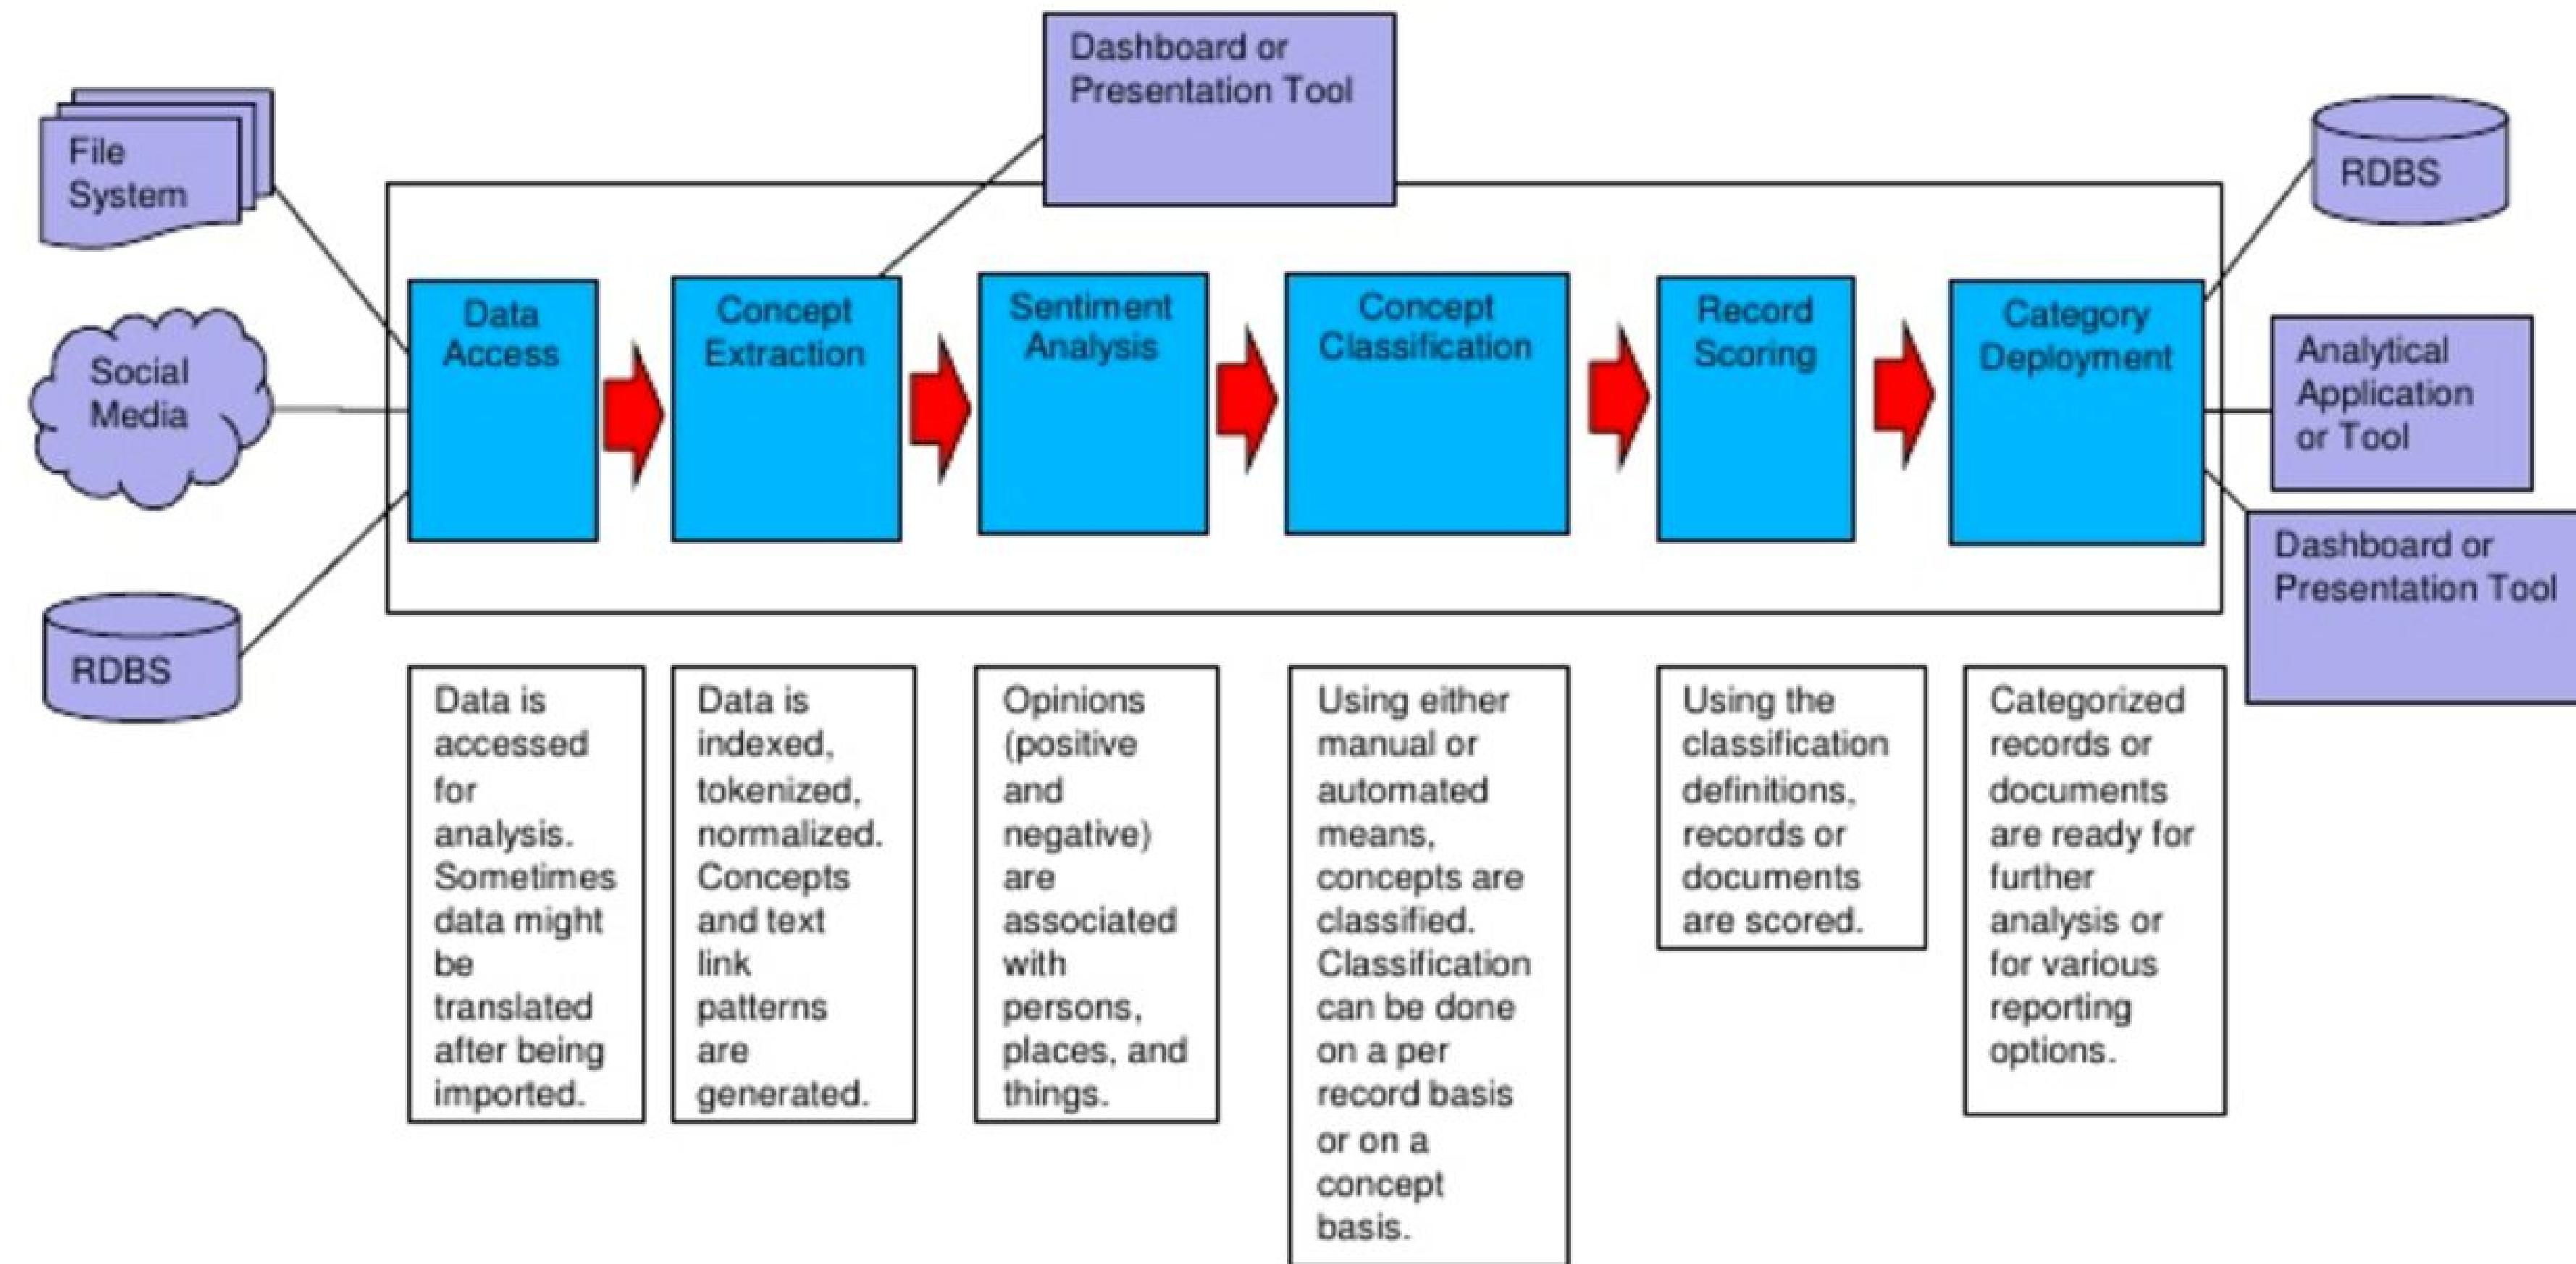
<http://deeplearning.net/reading-list/> -- Good Reading List

<https://github.com/ChristosChristofidis/awesome-deep-learning> -- Everything on Deep Learning

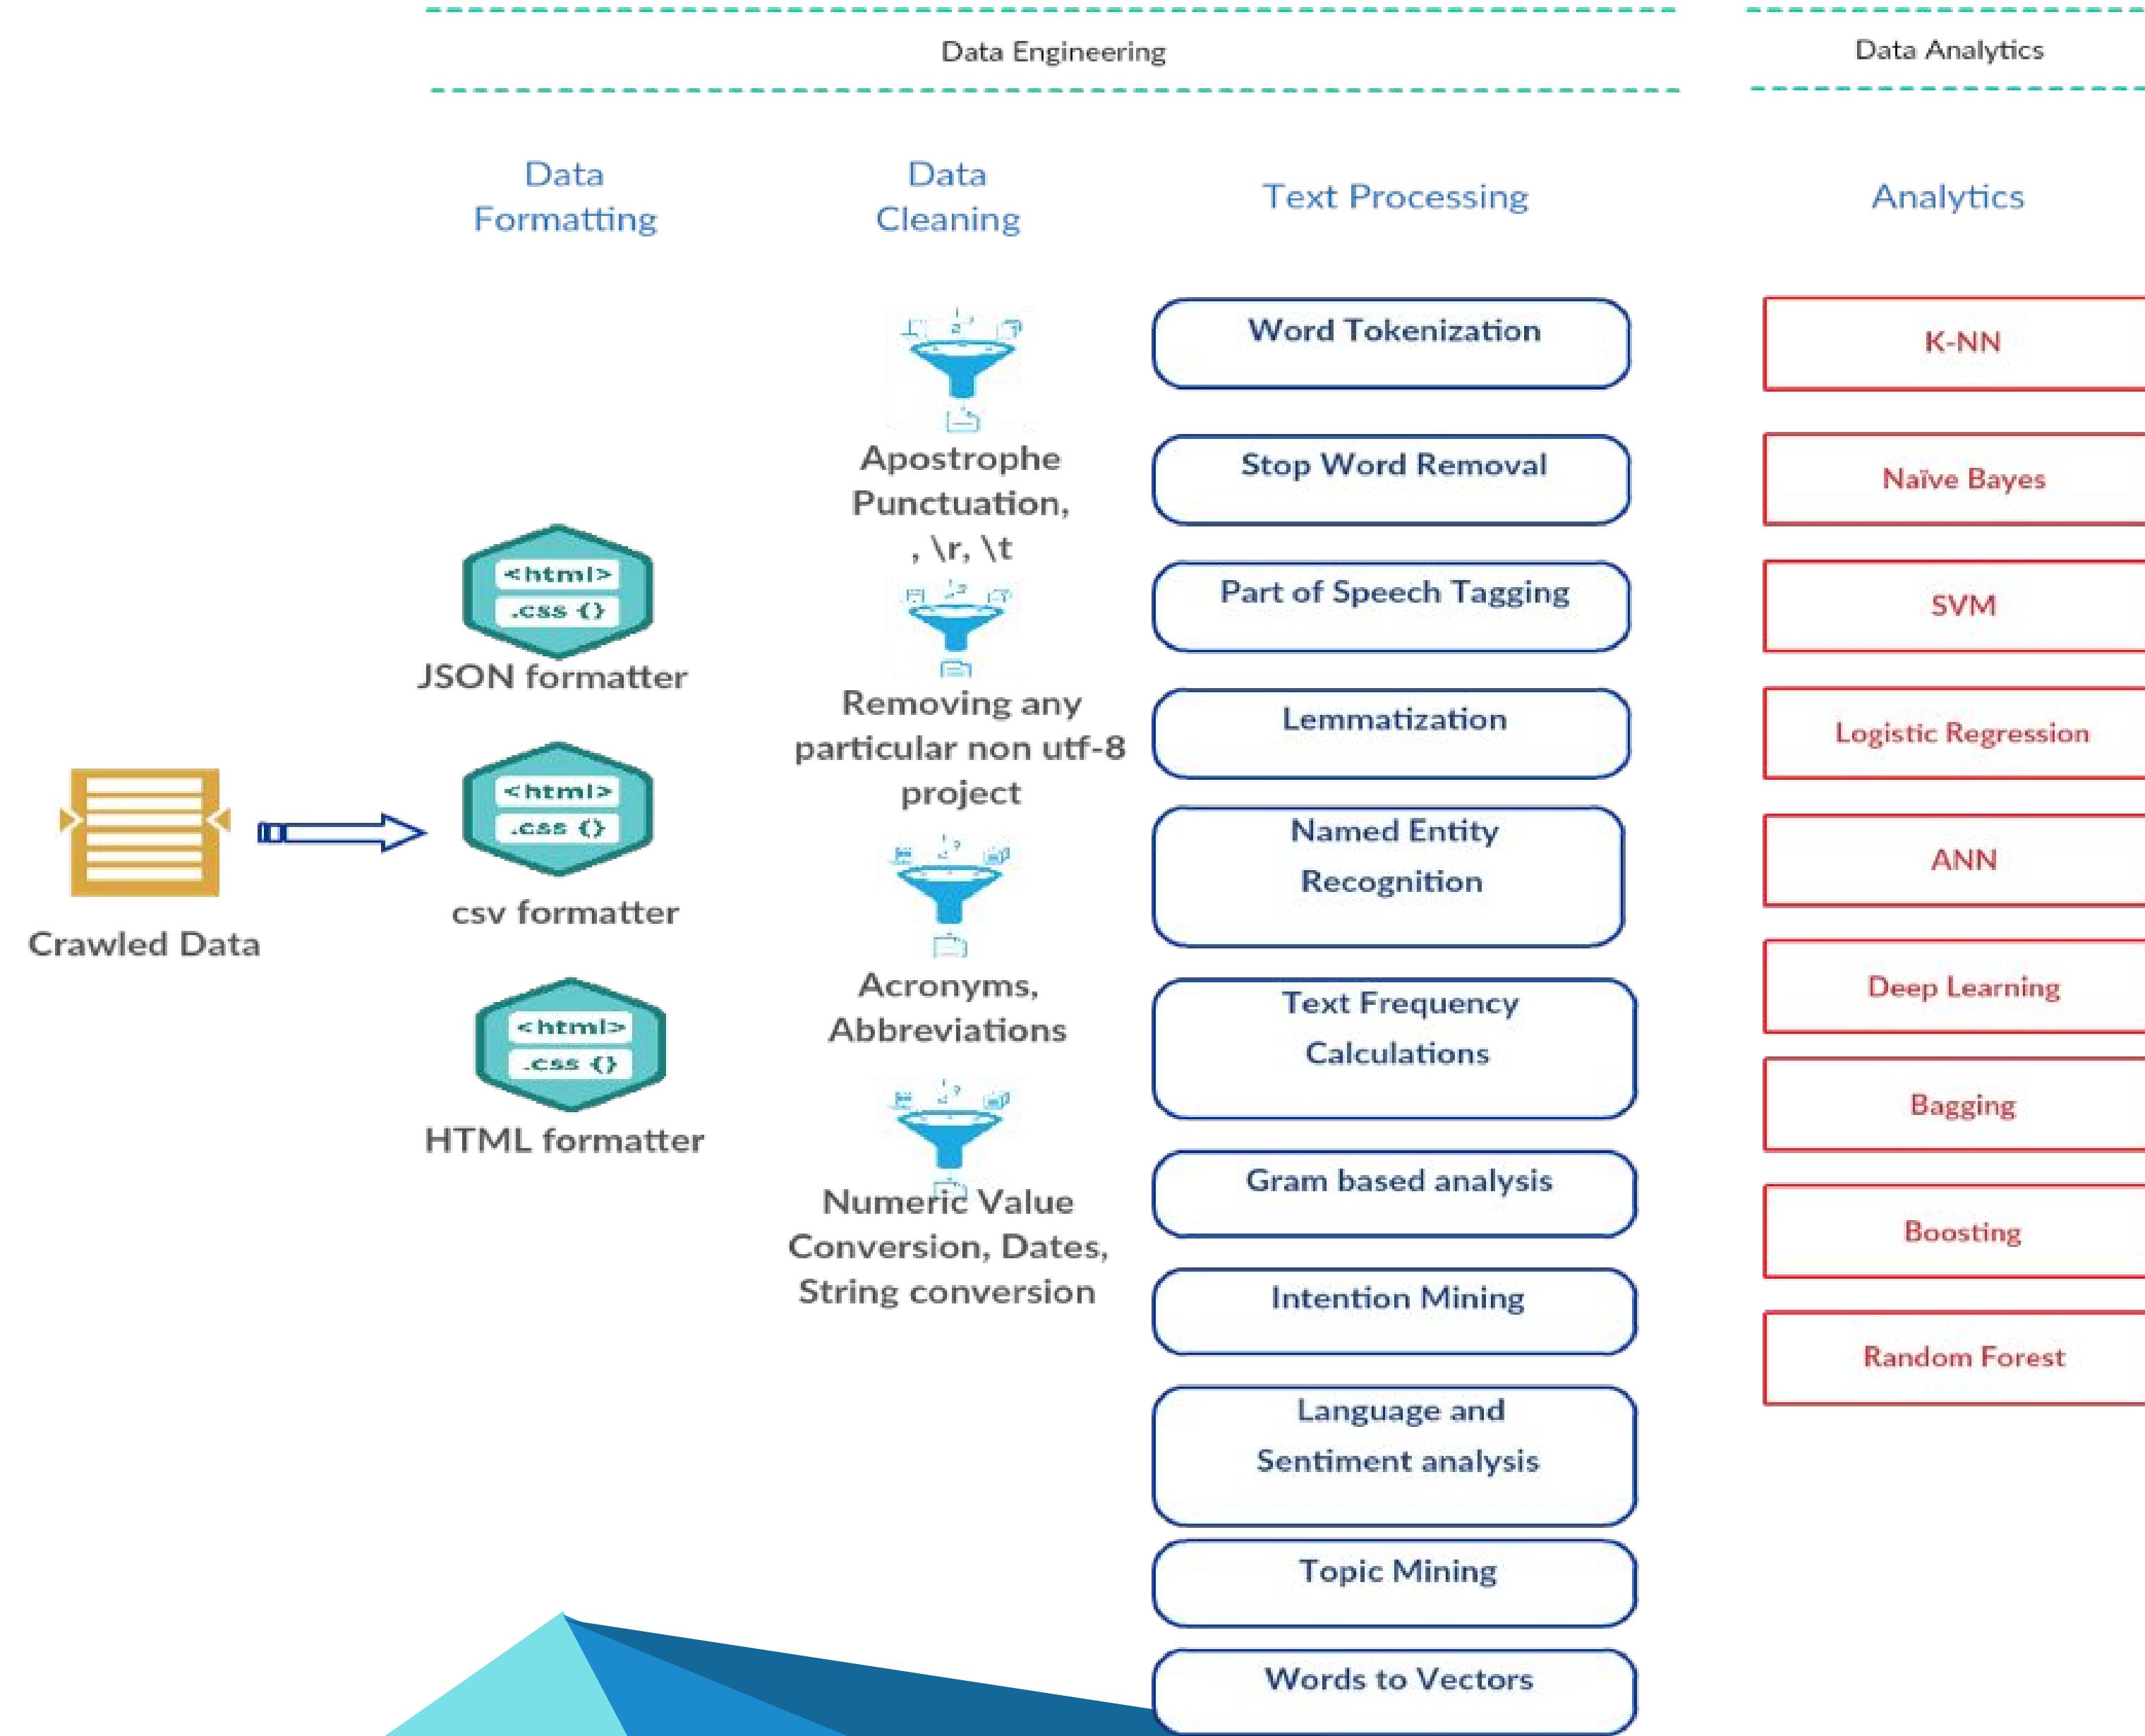
# Architectures



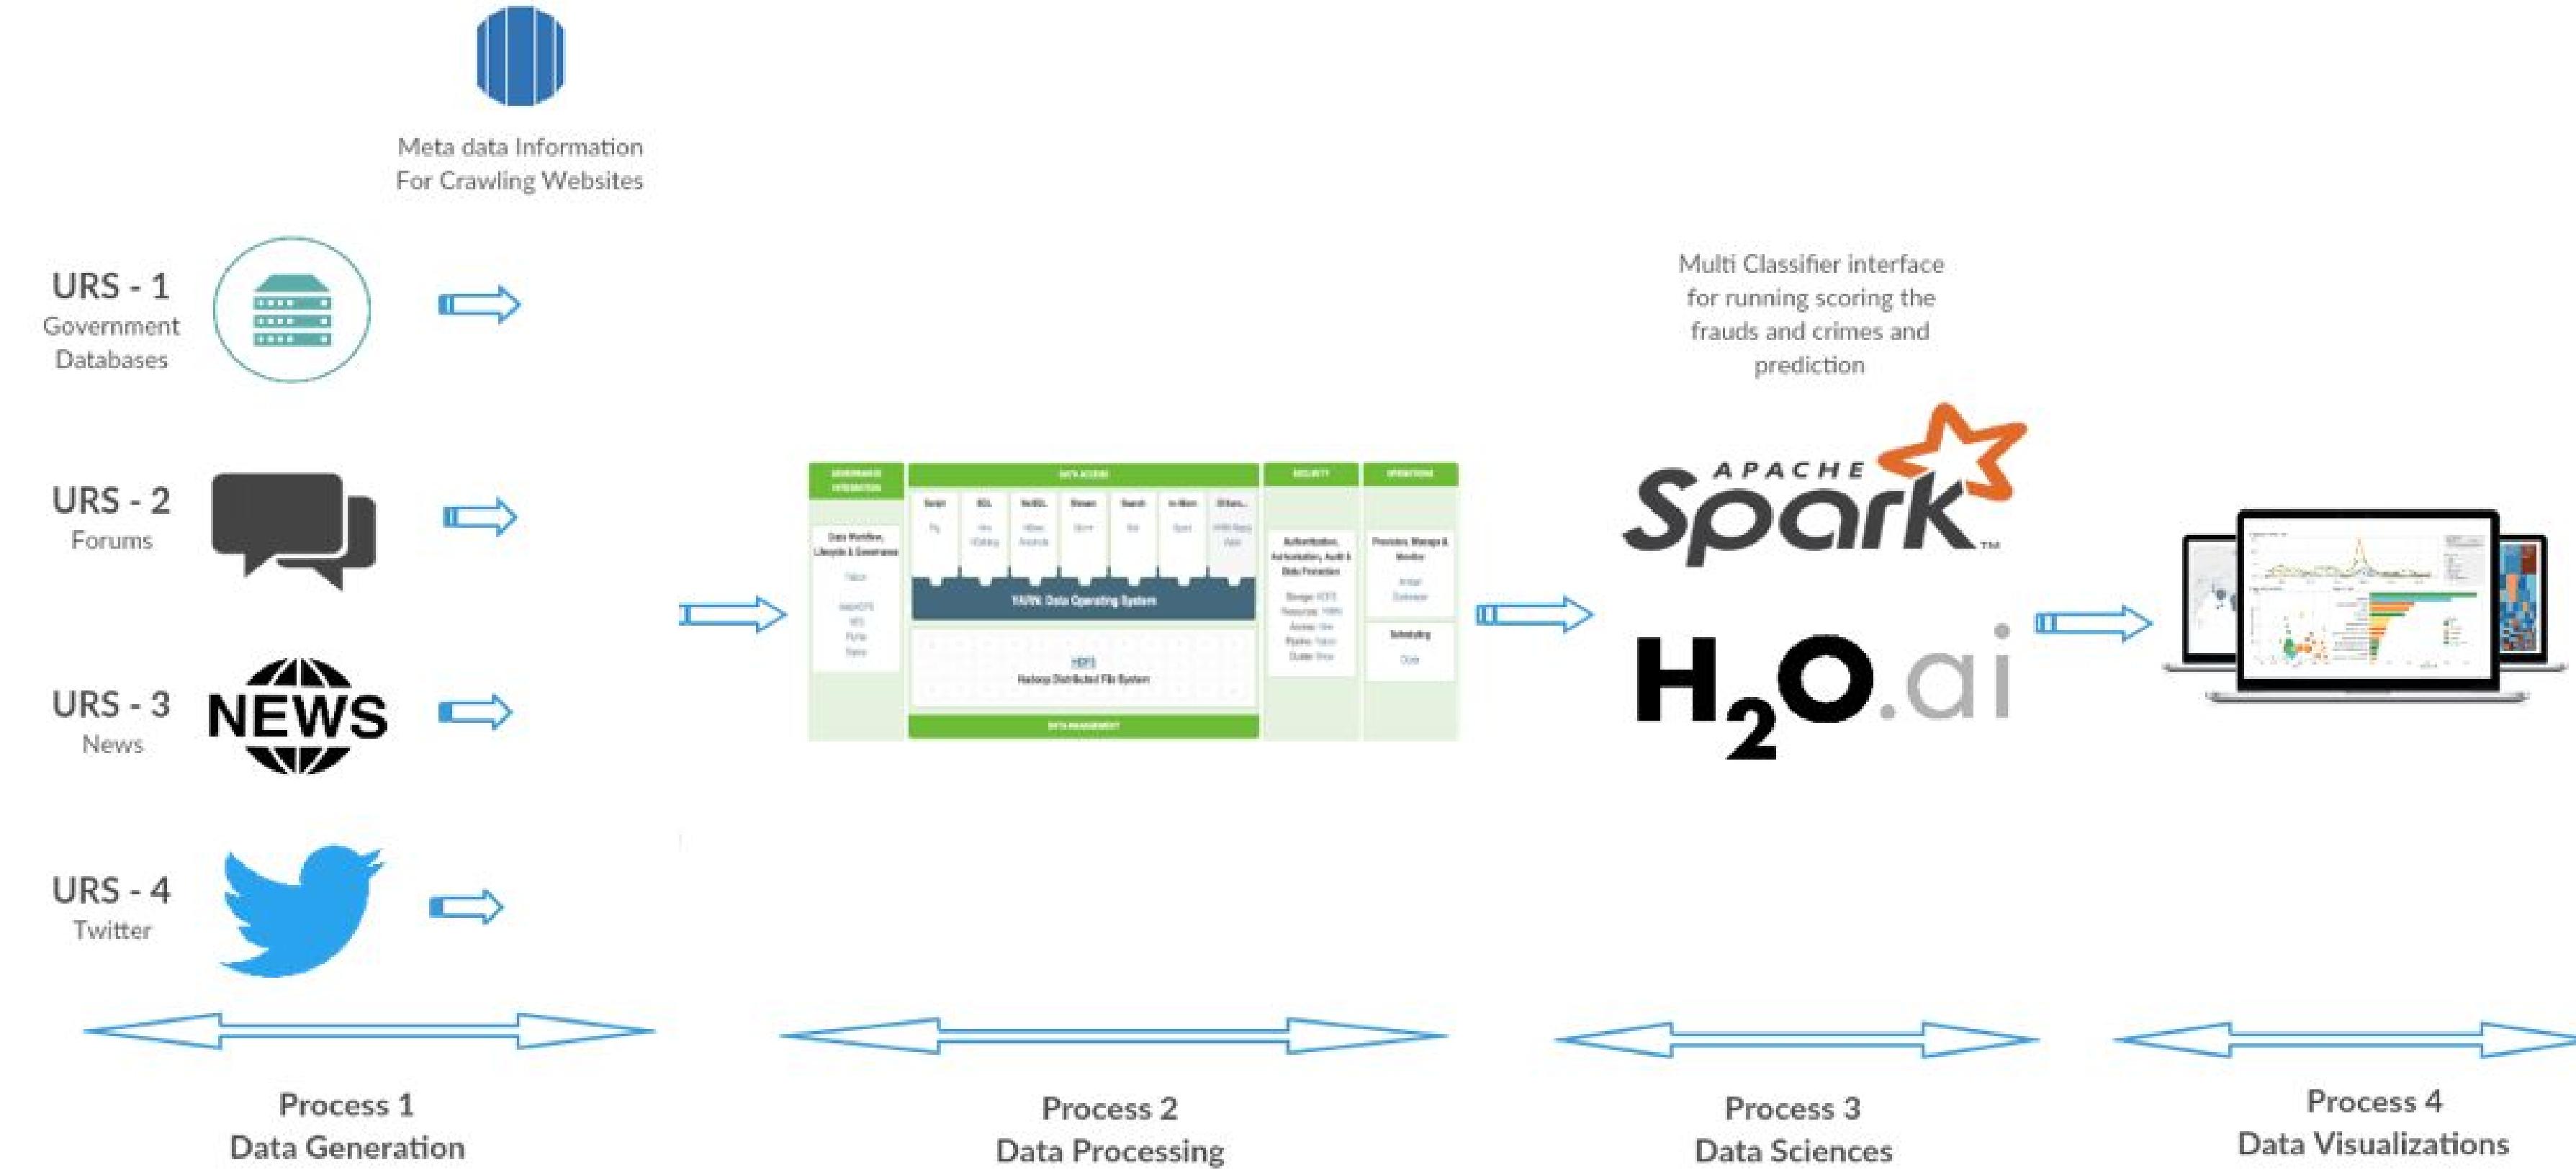
# NLP



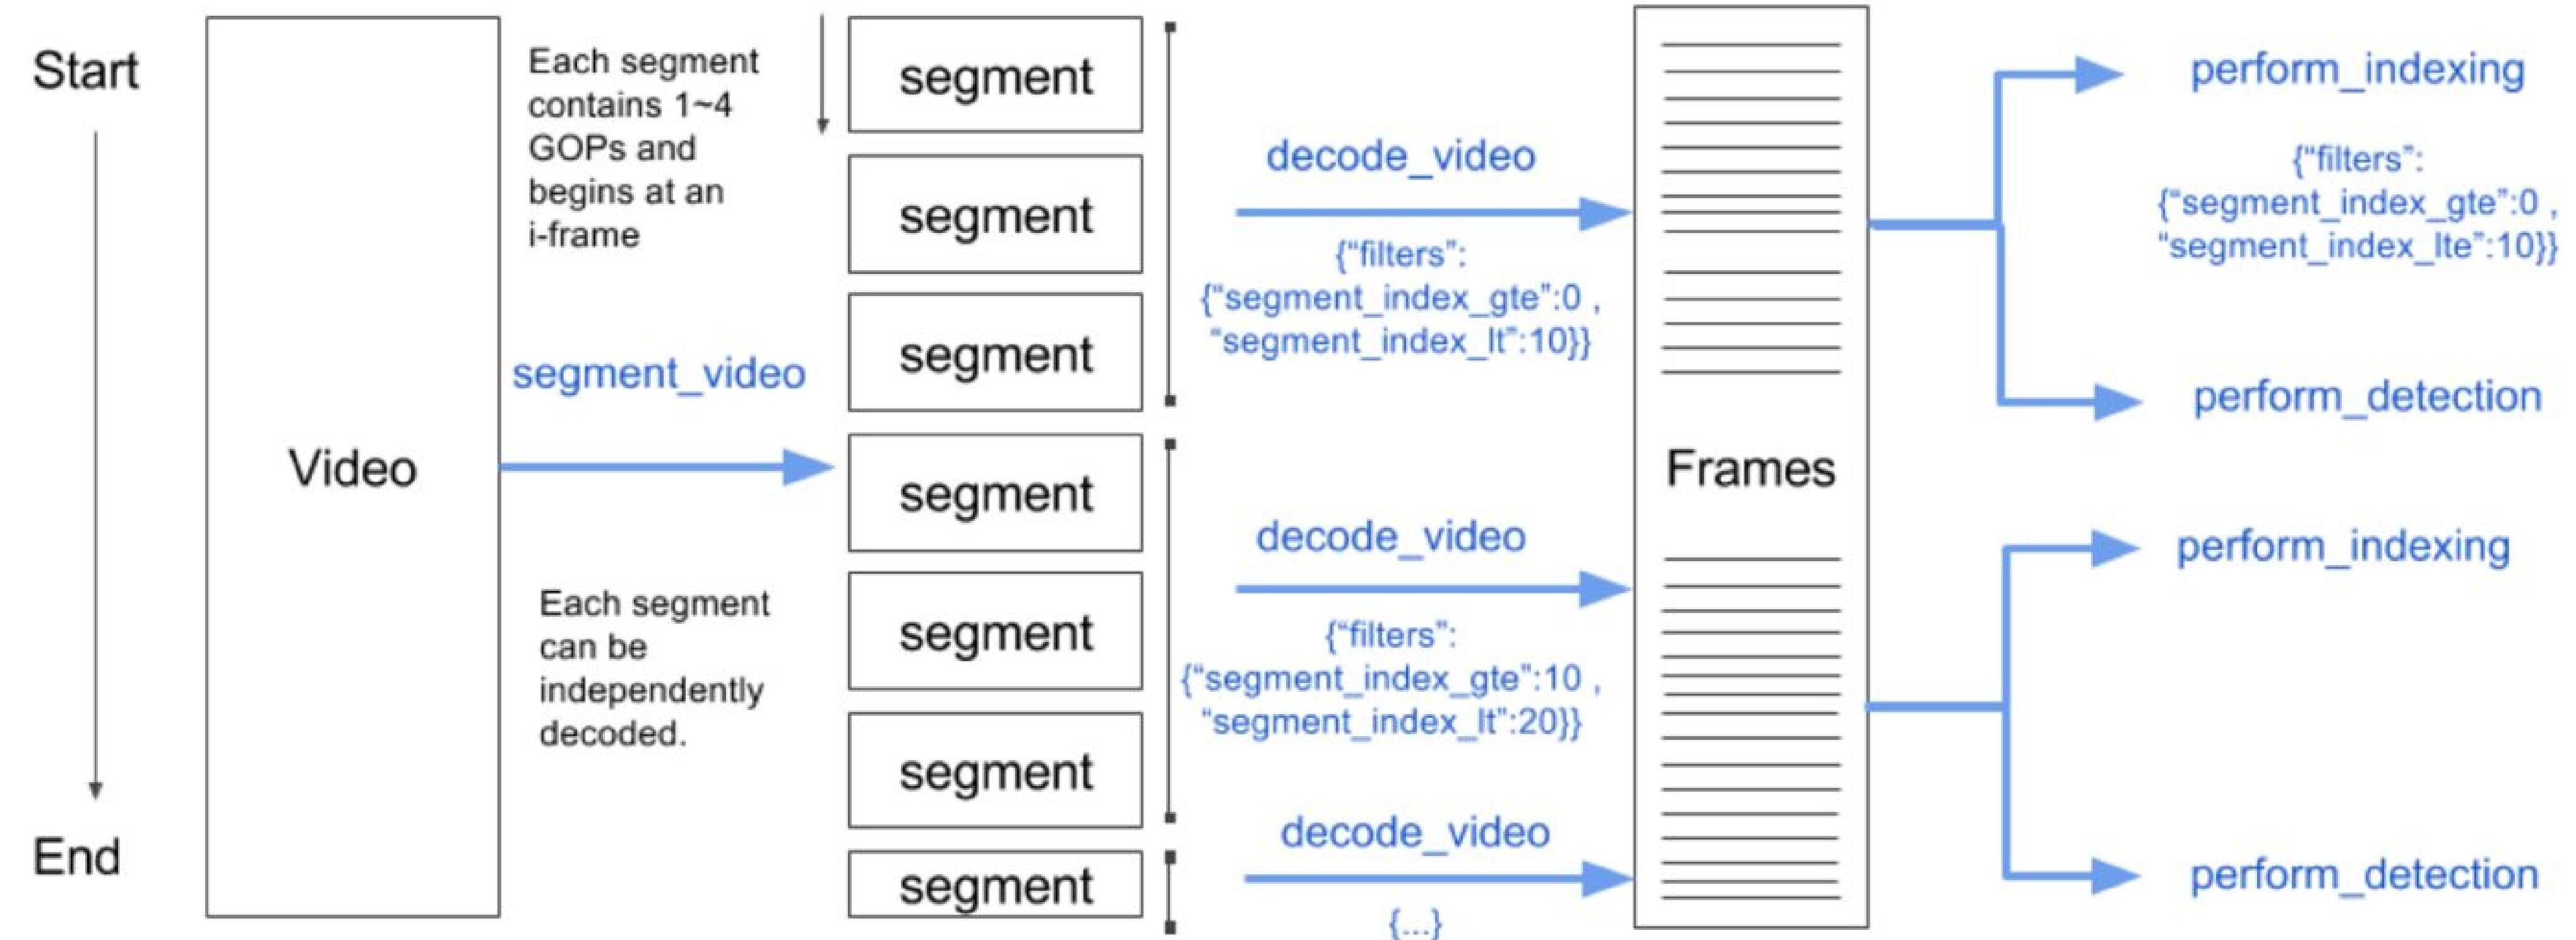
# NLP



# NLP



# Video Analytics



# Video Analytics

Data Collection



Model Development

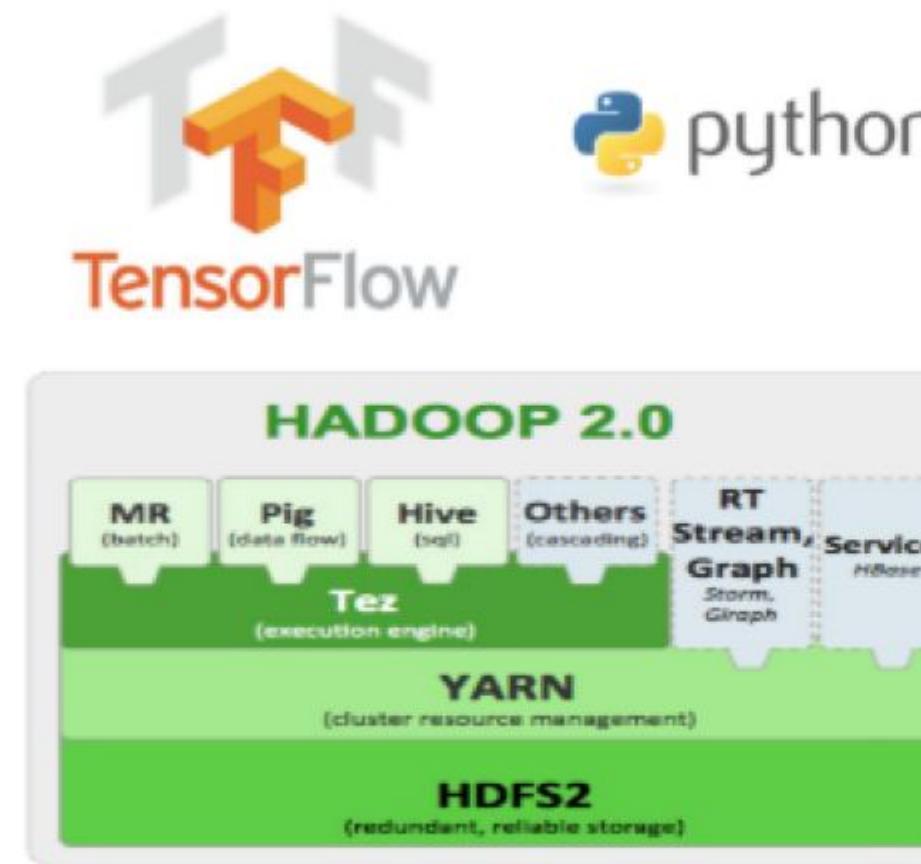
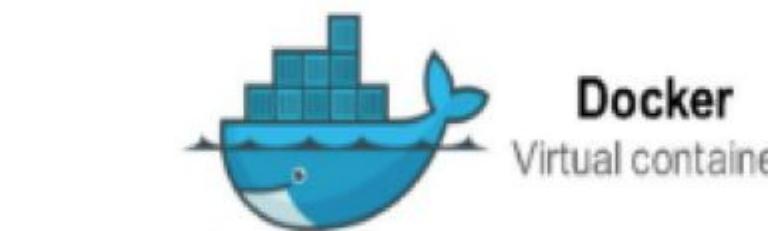


Image Detection and Video Analytics

**TensorFlow**  
Model development  
+  
**TensorBoard**  
Model analysis and tuning

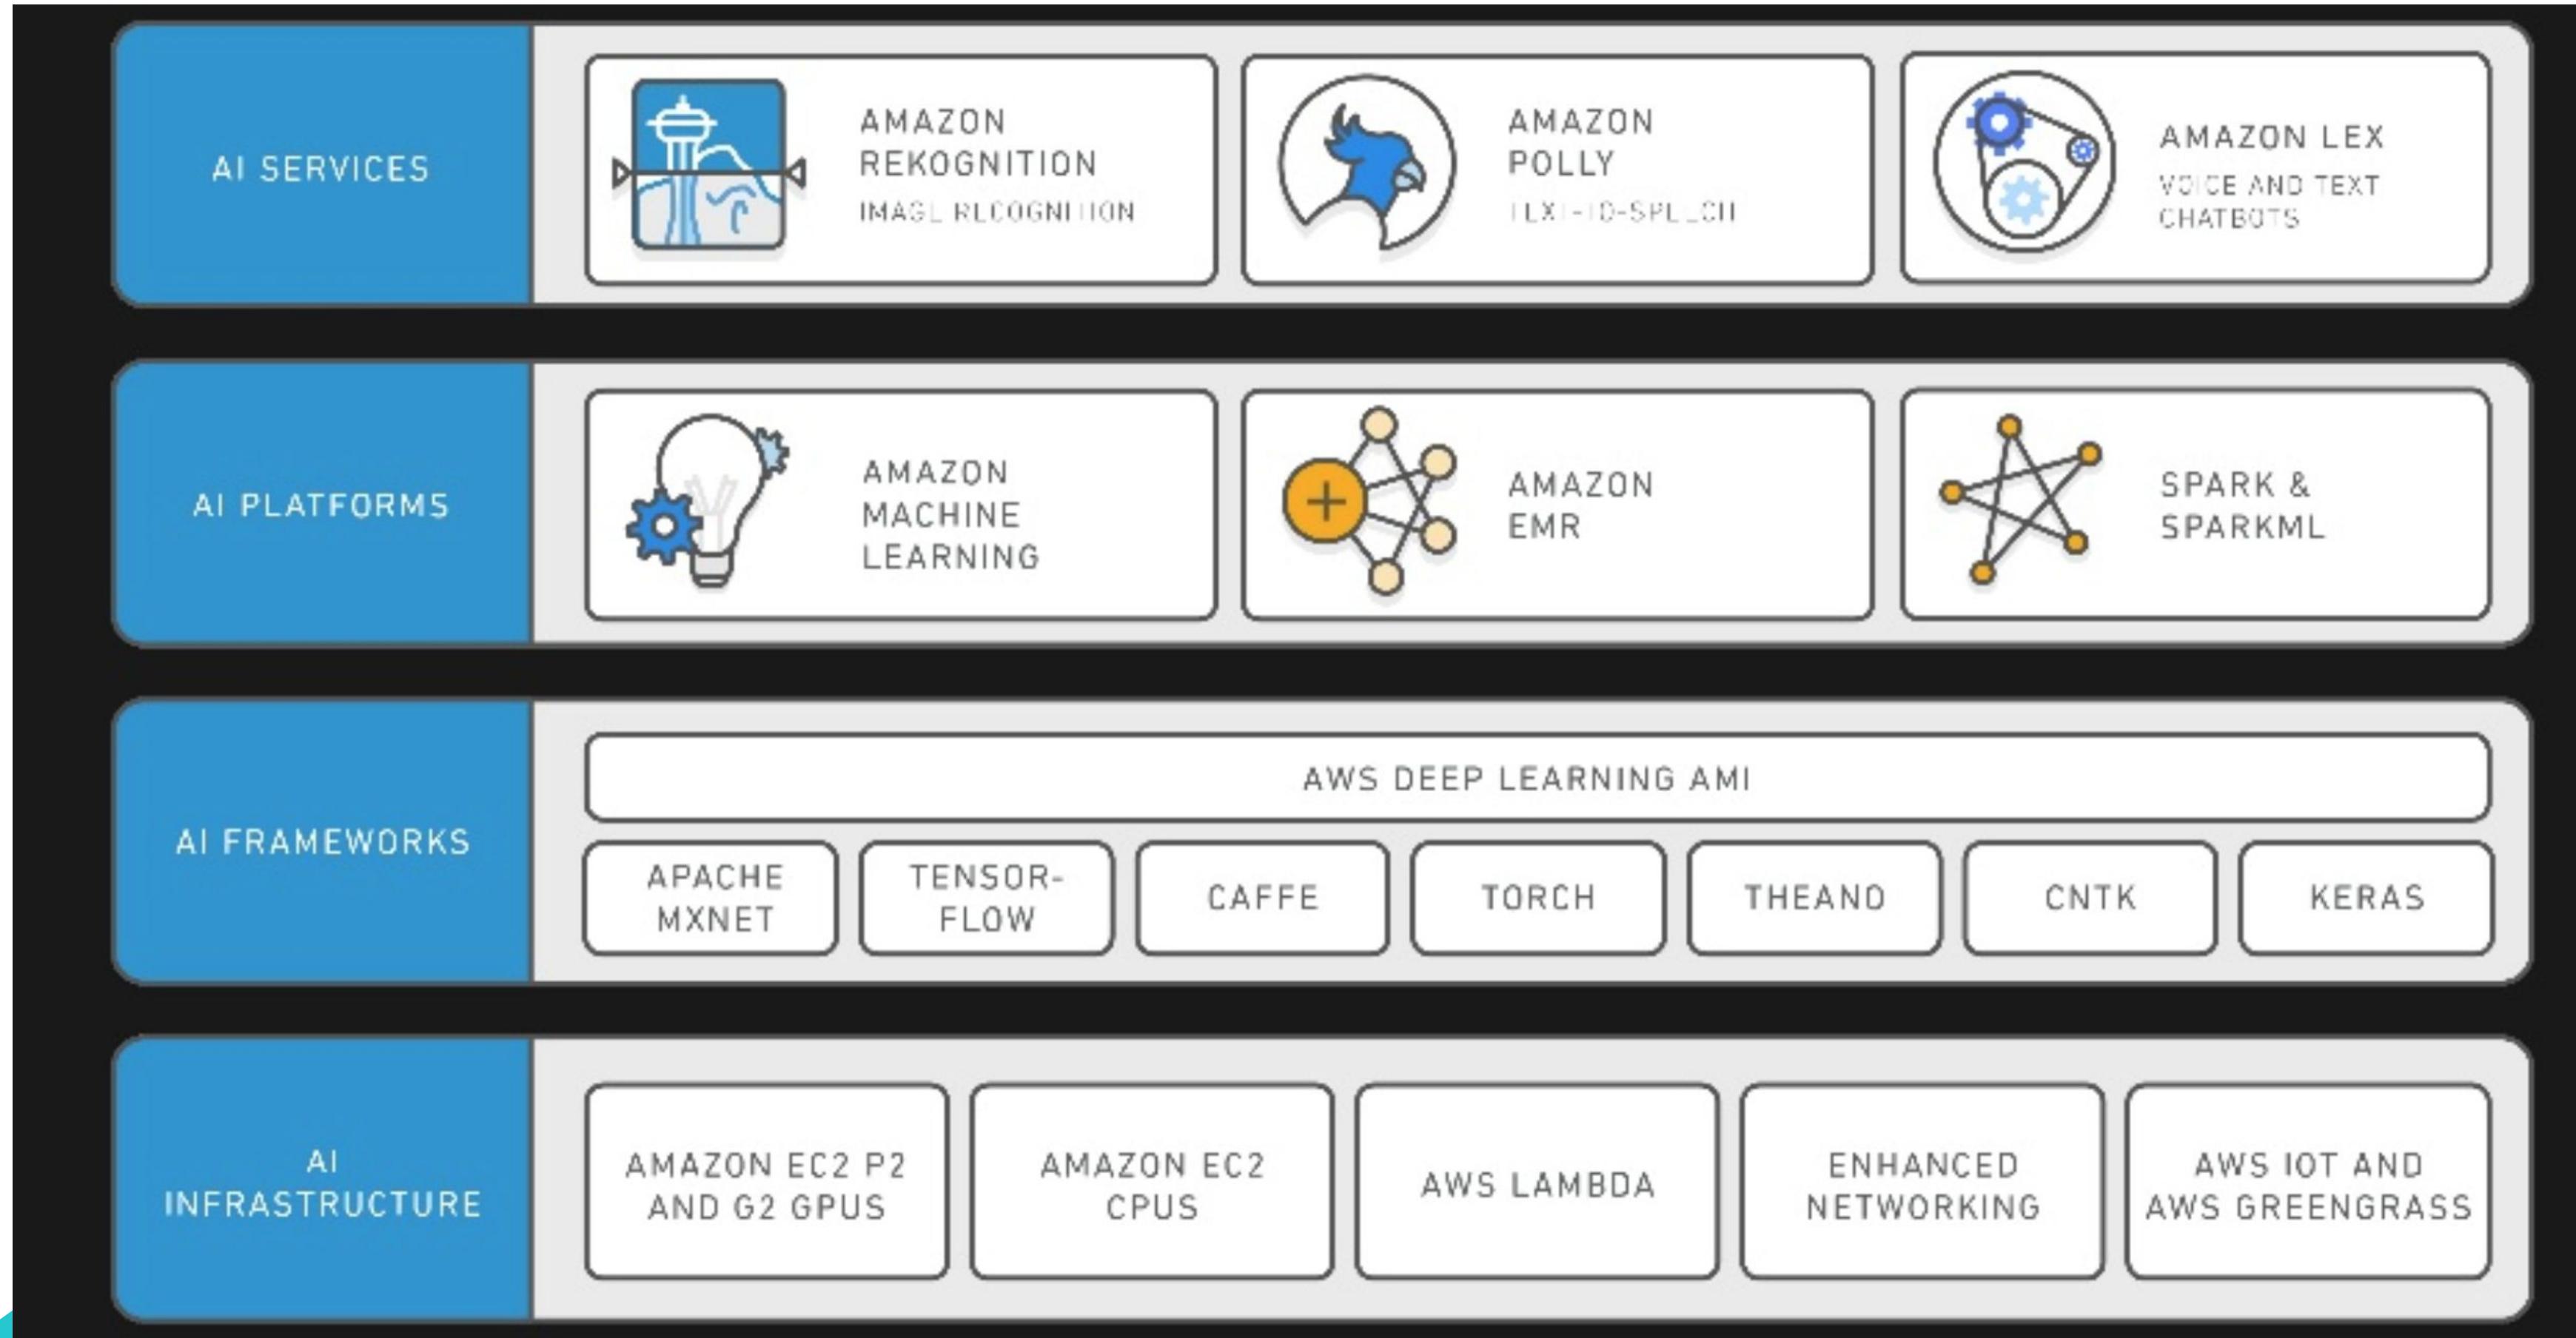


**Flask**  
Python Web App

Output



# AWS and Deep Learning

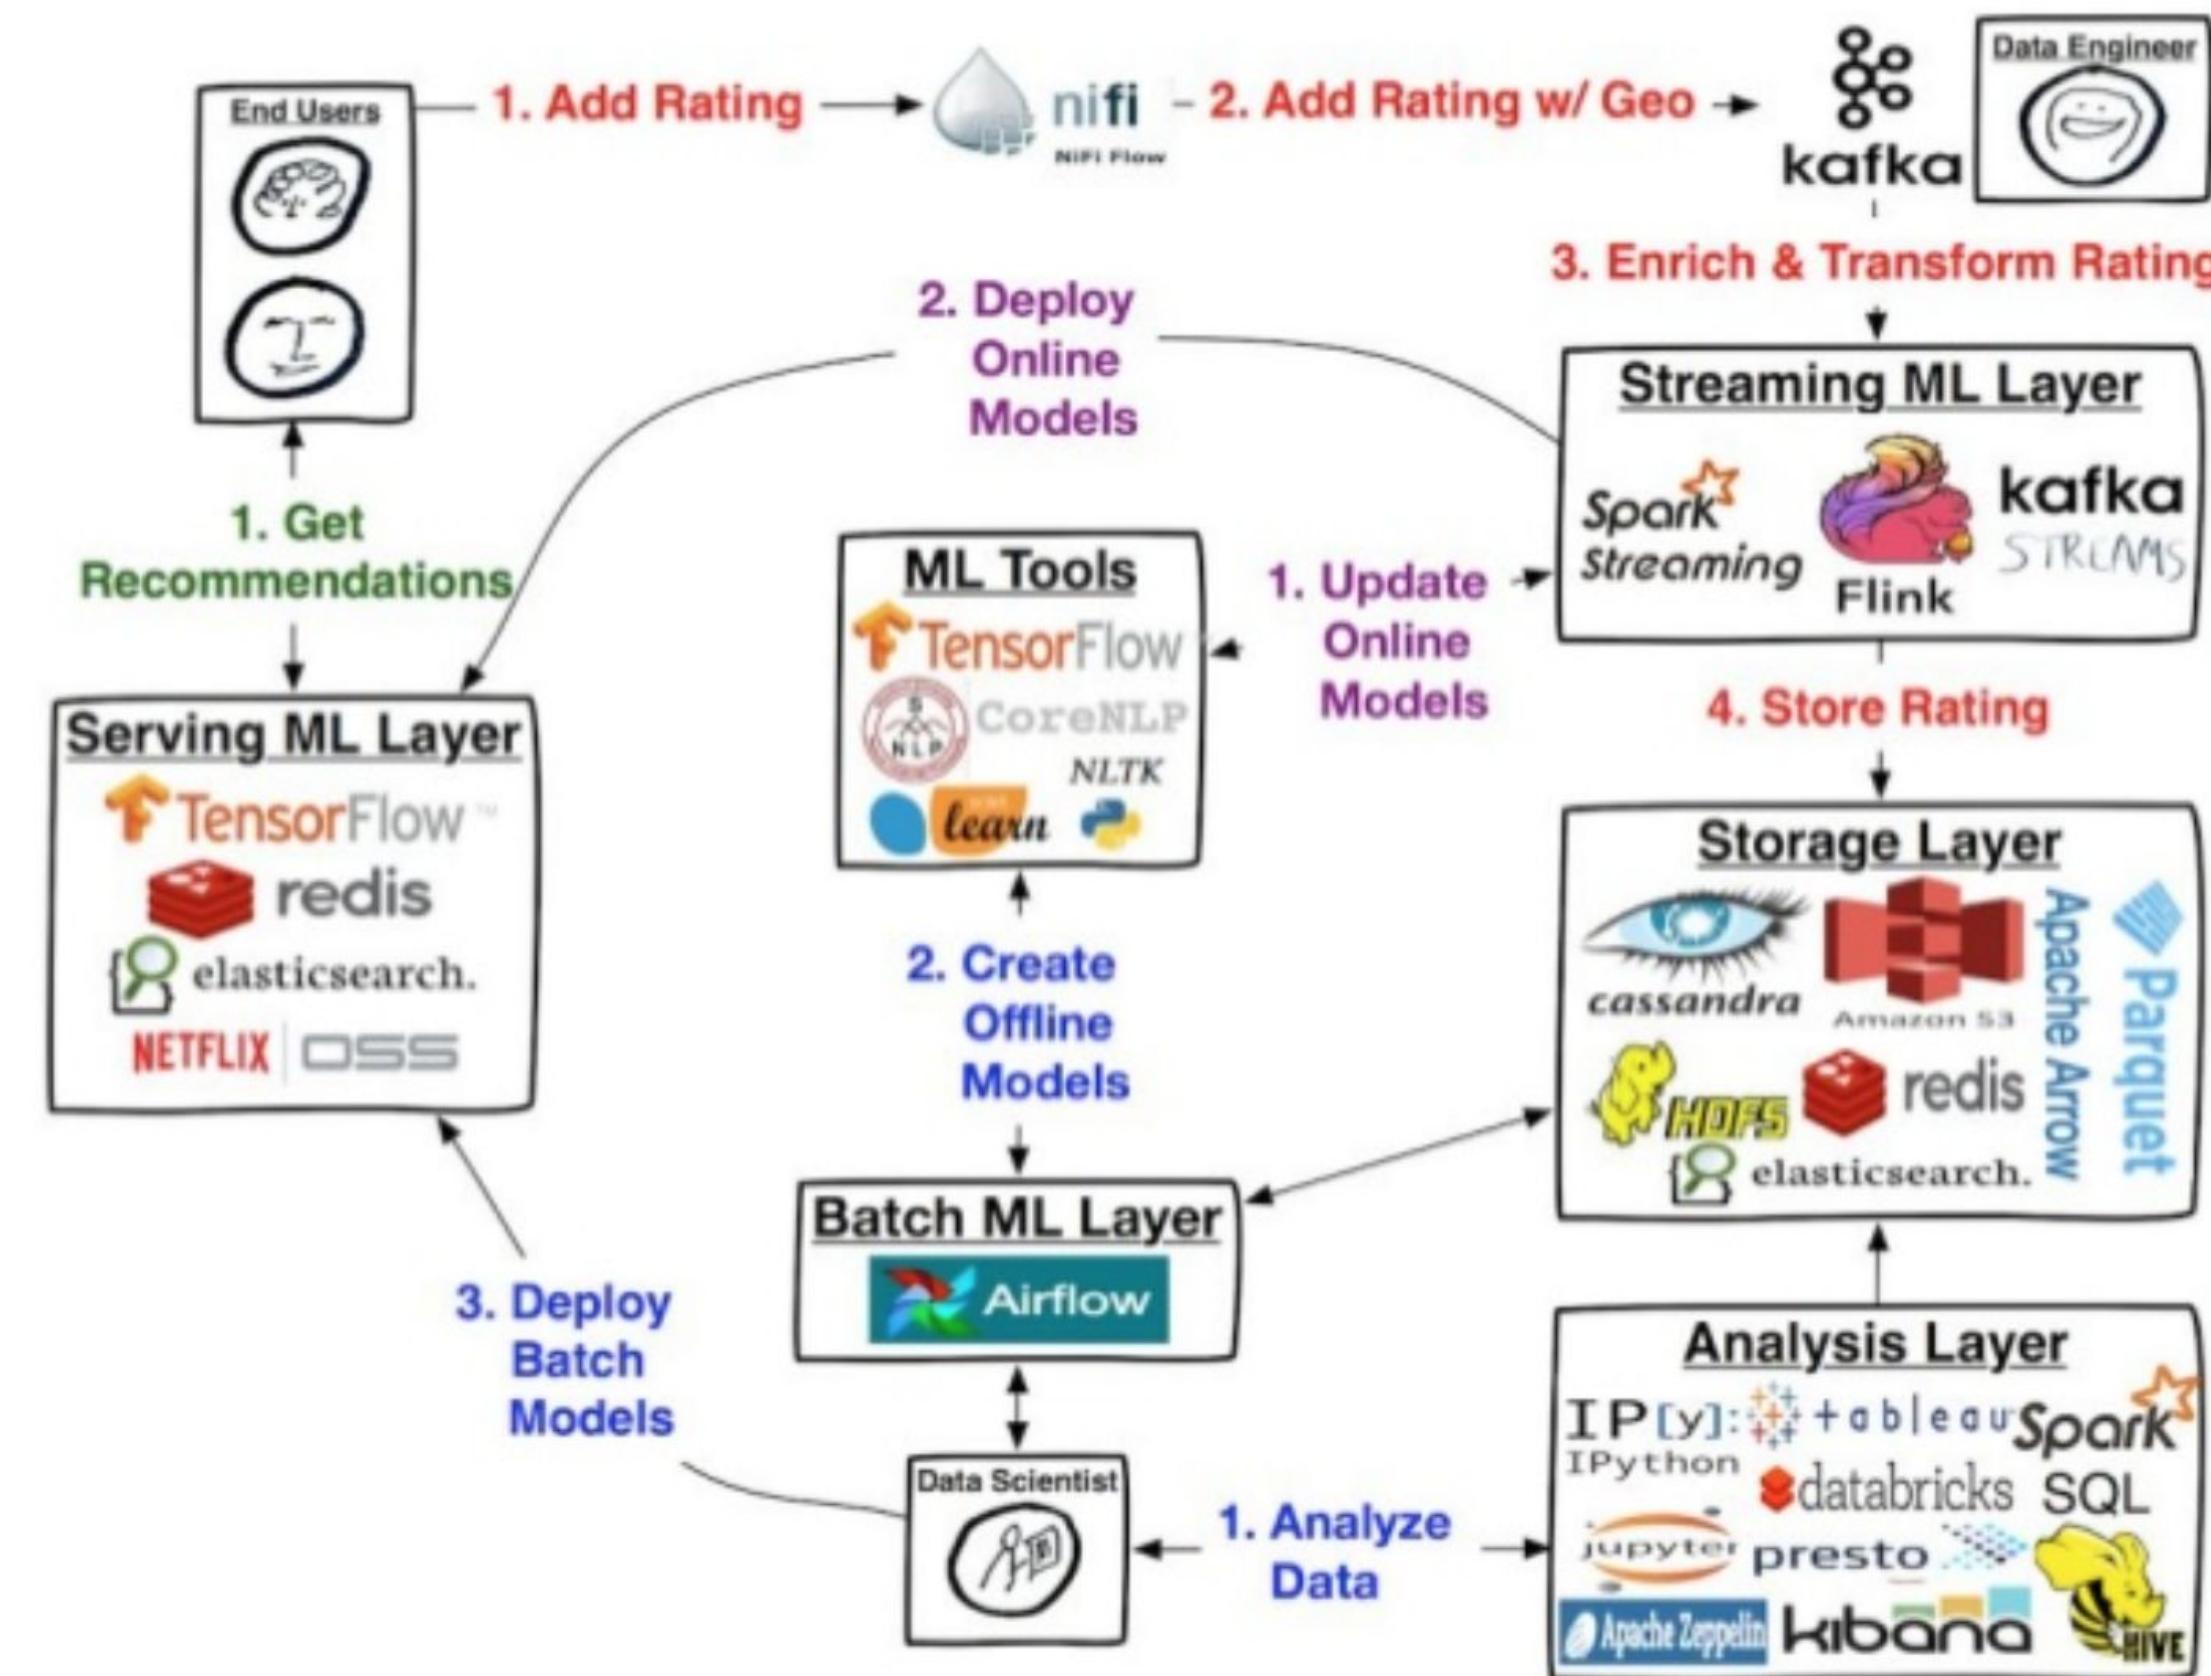


[Deep Learning AMI with Source Code \(CUDA 8, Amazon Linux\)](#)

Sold by: [Amazon Web Services](#)

The Deep Learning AMI is a base Ubuntu image provided by Amazon Web Services for use on Amazon Elastic Compute Cloud (Amazon EC2). It is designed to provide a stable, secure, and high performance execution environment for deep learning applications running on Amazon EC2. It has popular deep learning frameworks, including MXNet, Caffe, Caffe2, TensorFlow, PyTorch, Theano, CNTK, Torch and Keras as well as packages that enable easy integration with AWS. It also includes Anaconda Data Science Platform for Python2 and Python3. The Deep Learning AMI is provided at no additional charge to Amazon... [Read more](#)

# Deployment Plan



# Docker

Docker is a virtualization solution (similar to virtual machine). You can download container (or “image”) containing all the frameworks you need.

Why is it useful for DL?

- Installing all the DL frameworks takes time, so download a docker image instead.
- You are sure to have the same running environment on two different machines
- You cannot be root on the cluster.
- Don't share the code only. Share your docker image also.

# References

## Neural Networks and Deep Learning

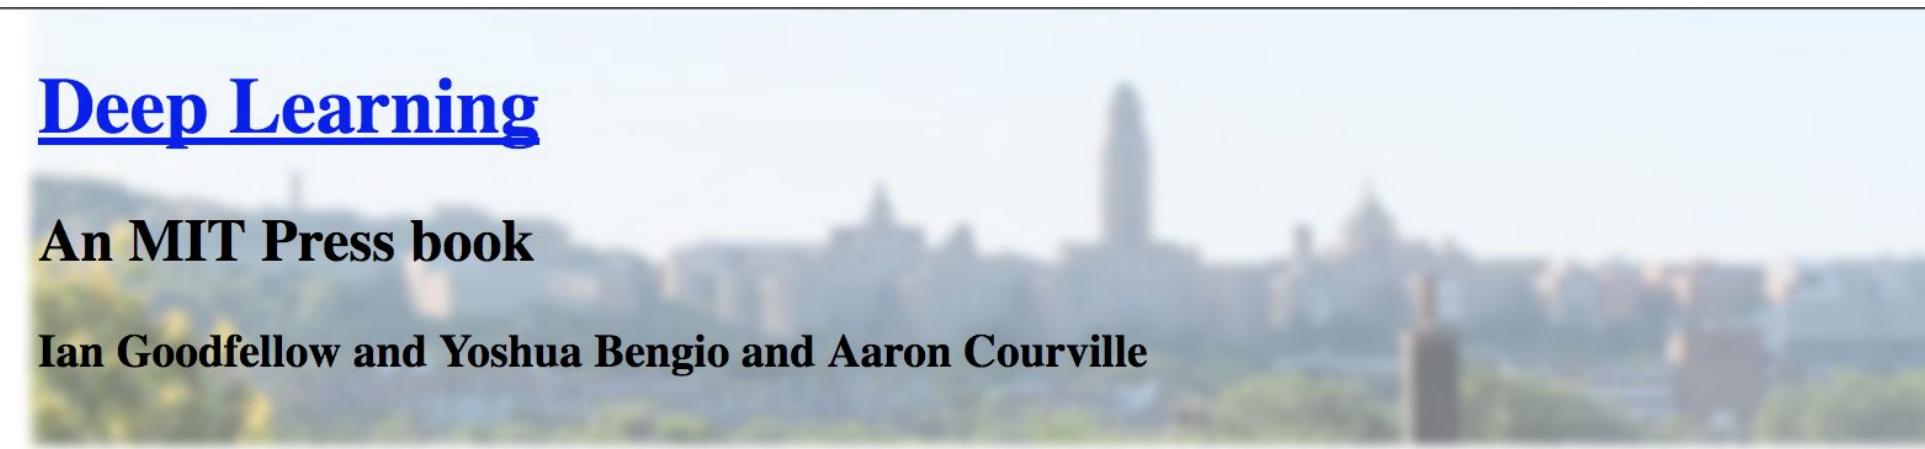
*Neural Networks and Deep Learning* is a free online book. The book will teach you about:

- Neural networks, a beautiful biologically-inspired programming paradigm which enables a computer to learn from observational data
- Deep learning, a powerful set of techniques for learning in neural networks

Neural networks and deep learning currently provide the best solutions to many problems in image recognition, speech recognition, and natural language processing. This book will teach you many of the core concepts behind neural networks and deep learning.

For more details about the approach taken in the book, [see here](#). Or you can jump directly to [Chapter 1](#) and get started.

<http://neuralnetworksanddeeplearning.com/>

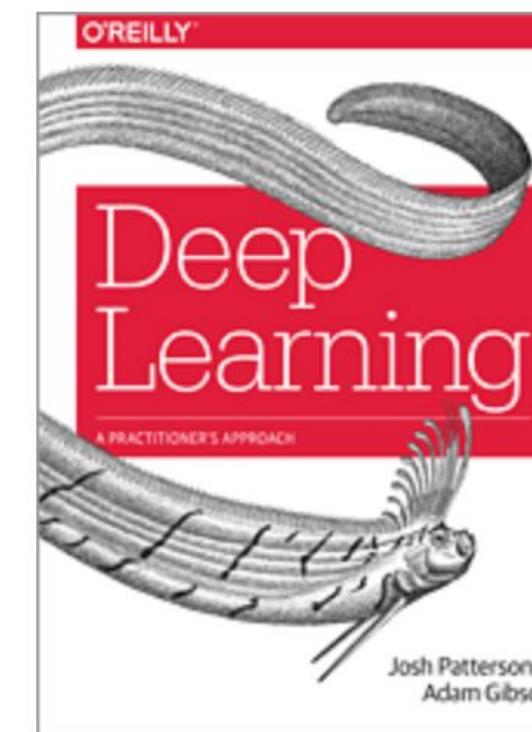


[Exercises](#) [Lectures](#) [External Links](#)

The Deep Learning textbook is a resource intended to help students and practitioners enter the field of machine learning in general and deep learning in particular. The online version of the book is now complete and will remain available online for free.

The deep learning textbook can now be pre-ordered on [Amazon](#). Pre-orders should ship on December 16, 2016.

For up to date announcements, join our [mailing list](#).



## Deep Learning A Practitioner's Approach

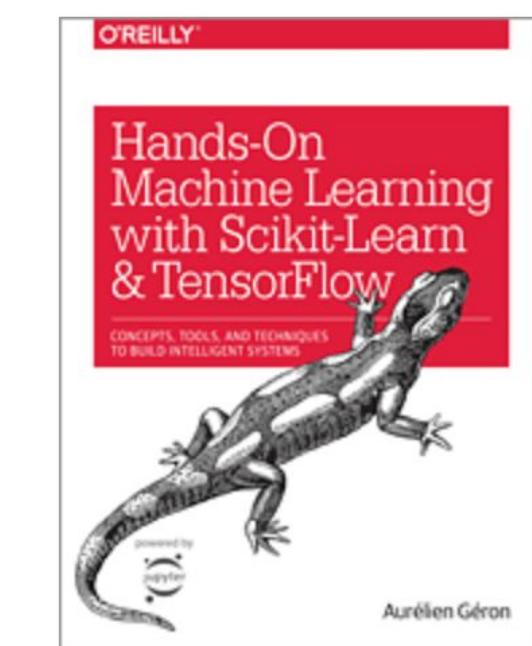
By [Adam Gibson](#), [Josh Patterson](#)

**Publisher:** [O'Reilly Media](#)

**Release Date:** August 2017

**Pages:** 532

Although interest in machine learning has reached a high point, lofty expectations often scuttle projects before they get very far. How can machine learning—especially deep neural networks—make a real difference in your organization? This hands-on guide not only provides the most practical information available on the subject, but also helps you get started building efficient deep learning networks.



## Hands-On Machine Learning with Scikit-Learn and TensorFlow

Concepts, Tools, and Techniques to Build Intelligent Systems

By [Aurélien Géron](#)

**Publisher:** [O'Reilly Media](#)

**Release Date:** March 2017

**Pages:** 570

**Graphics in this book are printed in black and white.**

Through a series of recent breakthroughs, deep learning has boosted the entire field of machine learning. Now, even programmers who know close to nothing about this technology can use simple, efficient tools to implement programs capable of learning from data. This practical book shows you how.



# Thanks!

Any questions?